



## Multiprocessor Programming Course

Dario Bahena Tapia - Ricardo Zavaleta Vazquez - Isai Barajas Cicourel

October 2015

# 1 LockFreeQueueTest

## 1.1 Particular Case (problem)

This is a particular case of the mutual exclusion problem, where the shared resource is a queue.

## 1.2 Solution

In order to guarantee the correctness of the multi-threaded access to the queue, it implements a lock free scheme on its *enq* and *deq* methods by putting waits before modifying the state of the queue. It does it incorrectly though, as we will see later.

```
1  class LockFreeQueue {
2      public int head = 0;    // next item to dequeue
3      public int tail = 0;    // next empty slot
4      Object[] items; // queue contents
5      public LockFreeQueue(int capacity) {
6          head = 0; tail = 0;
7          items = new Object[capacity];
8      }
9
10     public void enq(Object x) {
11         // spin while full
12         while (tail - head == items.length) {}; // spin
13         items[tail % items.length] = x;
14         tail++;
15     }
16
17     public Object deq() {
18         // spin while empty
19         while (tail == head) {}; // spin
20         Object x = items[head % items.length];
21         head++;
22         return x;
23     }
24 }
```

### 1.3 Experiment Description

The test program `LockFreeQueueTest` includes the following individual test cases; the parallel degree is two threads for each operation (queue or dequeue), with a `TEST_SIZE` of 512 (number of items to enqueue and dequeue) which is spread evenly among the two threads on each group:

- *testSequential*: calls *enq* method as many times as `TEST_SIZE`, and later calls *deq* method the same number of times checking that the FIFO order is preserved.
- *testParallelEnq*: enqueues in parallel (two threads) but dequeues sequentially.
- *testParallelDeq*: enqueues sequentially but dequeues in parallel (two threads)
- *testParallelBoth*: enqueues and dequeues in parallel (with a total of four threads, two for each operation).

### 1.4 Observations and Interpretations

The bottom line of this exercise is most likely, to show several types of problems that can occur if we do not use mutual exclusion; the queue implementation fails implementing it, making the test vulnerable to several cases of race conditions. Below we explain some of them.

#### 1.4.1 Non atomic dequeue: referencing invalid registers

The symptom for this problem is a `NullPointerException`:

```
There was 1 error:
1) testParallelEnq(mutex.LockFreeQueueTest) java.lang.NullPointerException
   at mutex.LockFreeQueueTest.testParallelEnq(LockFreeQueueTest.java:69)
   at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
   at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
   at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
```

where the offending line is a cast to an `Integer` from the value returned by the *deq* method; this implies that such function is returning *null*. A possible

scenario to produce this outcome is as follows. Prefixes of T1 or T2 indicate the thread running the action, and they refer to the line numbers of the code of method *deq* posted above.

- Assume that queue holds a single element which lives in *items* array at position 0; the array has *null* on position 1 (per initialization).
- T1: Executes method up to line 20.
- T2: Executes method up to line 19.
- T1: Executes line 21, getting *items[0]*.
- T2: Executes line 20, but as *head* was changed already, it gets *items[1]*.
- T1: Executes line 22 returning *items[0]*
- T2: Executes line 22 returning *items[1]*, which was *null*.

The problem with above interlacing derives from the fact that the *deq* method is not an atomic operation, hence it allowed both threads to enter into the critical section and compete for updating the shared variables.

#### 1.4.2 Non atomic dequeue: returning duplicate values

The symptom for this problem is a duplicate pop warning:

```
.parallel both
.sequential push and pop
.parallel enq
E.parallel deq
Exception in thread "Thread-7" junit.framework.AssertionFailedError:
DequeThread: duplicate pop
    at junit.framework.Assert.fail(Assert.java:47)
    at mutex.LockFreeQueueTest$DequeThread.run(LockFreeQueueTest.java:132)
```

where the offending line is an assertion that validates that nobody else has pop such value from the queue; as the threads which populate do have non overlapping ranges of values, the pop operations (*deq*) shall never return a duplicate one. But duplication is possible indeed, if we have a sequence like the one below between the two threads:

- Assume that queue holds a single element which lives in *items* array at position 0.
- T1: Executes method up to line 19.
- T2: Executes method up to line 19.
- T1: Executes line 20, getting *items[0]*.
- T2: Executes line 20, getting as well *items[0]*.
- T1: Executes line 21, setting *head* to 1.
- T2: Executes line 21, setting *head* to 2.
- T1: Executes line 22 returning *items[0]*.
- T2: Executes line 22 returning *items[0]*.

Not only we left the queue in an inconsistent state (head has incorrect value), but we also returned the same element twice, triggering then the violation on the test. The underlying problem is the same as previous case: lack of atomicity of the *deq* method.

### 1.4.3 Non atomic auto-increment: loosing values

The symptom for this problem is a never ending program, hanging on the test *testParallelBoth*. When produced several thread dumps of the Java program, we can see two hanging threads:

```
"Thread-7" #16 prio=5 os_prio=0 tid=0x00007f3140102000 nid=0x3f51
  runnable [0x00007f31226e9000]
  java.lang.Thread.State: RUNNABLE
    at mutex.LockFreeQueue.deq(LockFreeQueue.java:33)
    at mutex.LockFreeQueueTest$DeqThread.run(LockFreeQueueTest.java:141)

"main" #1 prio=5 os_prio=0 tid=0x00007f3140009800 nid=0x3f3c in Object.wait()
  [0x00007f3148382000]
  java.lang.Thread.State: WAITING (on object monitor)
    at java.lang.Object.wait(Native Method)
    - waiting on <0x00000000d6effea8> (a mutex.LockFreeQueueTest$DeqThread)
    at java.lang.Thread.join(Thread.java:1245)
    - locked <0x00000000d6effea8> (a mutex.LockFreeQueueTest$DeqThread)
    at java.lang.Thread.join(Thread.java:1319)
```

```

at mutex.LockFreeQueueTest.testParallelBoth(LockFreeQueueTest.java:123)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:497)
at junit.framework.TestCase.runTest(TestCase.java:164)
at junit.framework.TestCase.runBare(TestCase.java:130)

```

The main thread is waiting for a dequeue thread to finish; but that one is on an infinite loop at line 19 of *deq* method (line numbers per our listing in this document, not in the file). This means that we have lost some of the inserted elements in queue, and that one of the dequeue threads will never finish; as they expect each to pop a fixed amount of elements per the test.

But the amount of times we request a dequeue operation, among all threads, is the same as the number of elements we queued; how come we end up losing some of those? One possible explanation is again, the lack of atomicity but this time of the *enq* method; to be more specific, the lack of atomicity of its auto-increment operation *tail++* (line 14). Let us remember that the auto-increment operator in Java is nothing but syntactic sugar for the following sequence of operations (when applied to *tail* variable):

```

tmp = tail;
tmp = tmp + 1;
tail = tmp;

```

If two threads execute the lower level operations above, we can see how they can end up losing increments in the shared variable *tail*; the following sequence is an example of such scenario:

- T1: executes *tmp = tail*.
- T2: executes *tmp = tail*.
- T1: executes *tmp = tmp + 1*.
- T2: executes *tmp = tmp + 1*.
- T1: executes *tail = tmp*.
- T2: executes *tail = tmp*.

We can appreciate that in the above interlacing, the final value of the shared variable is  $tail + 1$ ; instead of the expected value of  $tail + 2$ . It would be enough to loose a single value this way, in order to make the enqueue threads think that they inserted the total of 512, while they really inserted 511; as each dequeue thread will try to pop 256 each, only one of them will be able to finish while the other will get blocked after having removed 255 entries. That is most likely the explanation for the hung threads we pasted above <sup>1</sup>; the solution is again to really implement mutual exclusion around the methods *deq* / *enq*; in such a way that they become atomic operations.

## 1.5 Proposed changes to fix the problems

All the three scenarios described before can be eliminated, if we make the methods *enq* and *deq* atomic; this can be easily achieved in Java by making them *synchronized*. However, by doing that, we will loose parallelism among the two groups of threads (those calling *enq* and those calling *deq*); this is because the *synchronized* keyword uses as lock the whole object, so at any moment in time, only one synchronized method can actually run within any object. In order to overcome this limitation, we can use synchronized blocks against two different lock objects (one for each operation).

Even with the changes above, we can still have issues; what if the very first thread running is one calling *deq* method? It will find the queue empty and loop forever. In order to prevent that, we should remove the waiting operation out of the queue methods, and put them in the test code itself. This is because, on the cited scenario that we try to dequeue with an empty queue, we would expect to simply try again (giving the chance to parallel enqueue threads to produce something for us to pop). The final code which incorporates these fixes is listed below:

---

<sup>1</sup>Note that it does not matter that the array *items* has populated all the correct entries, because the flow is controlled by the counters *tail* and *head*.

```

1  class LockFreeQueue {
2      private static Object enqLock = new Object();
3      private static Object deqLock = new Object();
4
5      public int head = 0;    // next item to dequeue
6      public int tail = 0;    // next empty slot
7      Object[] items; // queue contents
8      public LockFreeQueue(int capacity) {
9          head = 0; tail = 0;
10         items = new Object[capacity];
11     }
12
13     public boolean enq(Object x) {
14         synchronized(enqLock)
15         {
16             if (tail - head == items.length) {
17                 return false;
18             }
19             items[tail % items.length] = x;
20             tail++;
21             return true;
22         }
23     }
24
25     public Object deq() {
26         synchronized(deqLock)
27         {
28             if (tail == head) {
29                 return null;
30             }
31             Object x = items[head % items.length];
32             head++;
33             return x;
34         }
35     }
36 }

```



The test code was also modified, to make the enqueue and dequeue threads to iterate until they have successfully called their respective methods 256 times (total size of queue divided by number of threads). A successful call is one that does not return *false* nor *null*. The modified code was executed ten thousand times and it did not produce any of the original problems we explained. Though not a formal proof of its correctness, it is a good indication of the same (the original program produced one of the cited problems quite often, usually within 10 executions).

## 2 TreeTest

### 2.1 Particular Case (problem)

This problem belongs to chapter 12, where the idea is to present problems which look inherently serial but that surprisingly accept interesting concurrent solutions (though not always easy to explain).

The particular problem this test is exercising is that of shared counting, where we have  $N$  threads to increase a shared numeric variable up to certain value; but with the goal of producing less contention than a serial solution would generate.

### 2.2 Solution

The java program is called *Tree.java*, although it corresponds to the *CombiningTree* from the book, which is a binary tree structure where each thread is located at a leaf and at most two threads can share a leaf. The shared counter is at the root of the tree, and the rest of the nodes serve as intermediate result points. When a couple of threads collide in their attempt to increment, only one of them will serve as the representant and go up in the tree with the mission of propagating the combined increment (2) up to the shared counter at the root of the tree. In its way up, it may encounter further threads that collide again with it, making it wait or continue its journey to the root.

When a thread reaches the root it adds its accumulated result to the shared counter, and propagates down the news that the job is done to the rest of the threads that waited along the way. This solution to the count-

ing problem has worse latency than lock-based solutions ( $O(1)$  vs  $O(\log(p))$ ), where  $p$  is the number of threads or processors; but it offers a better throughput.

## 2.3 Experiment Description

The program consists of a single unit test *testGetAndIncrement*, which creates 8 threads to perform each  $2^{20}$  increments. The program uses an auxiliary array called *test* to record the individual increments done by each thread; assertions are made to ensure that none of the attempts results in a duplicated value, as well as to ensure that all threads completed their task.

## 2.4 Observations and Interpretations

The test performs without much controversy in an elapsed time between 3 and 4 seconds (laptop computer with i5 processor). Below a sample output:

```
.Parallel, 8 threads, 1048576 tries
Time: 3.641
```

To make the test a little more interesting, we created another couple of tests reusing sample template of existing one; difference was the Thread class used on each case.

The first additional test uses a thread class based on Java provider *java.util.concurrent.atomic.AtomicInteger*.

```
cnt2 = new AtomicInteger();

class MyThread2 extends Thread {
    public void run() {
        for (int j = 0; j < TRIES; j++) {
            int i = cnt2.getAndIncrement();
            if (test[i]) {
                System.out.printf("ERROR_duplicate_value%d\n", i);
            } else {
                test[i] = true;
            }
        }
    }
}
```

```

    }
}

```

while the second additional test was based on a custom class that used synchronized *getAndIncrement* method:

```

class MyThread3 extends Thread {
    public void run() {
        for (int j = 0; j < TRIES; j++) {
            int i = cnt3.getAndIncrement();
            if (test[i]) {
                System.out.printf("ERROR_duplicate_value_%d\n", i);
            } else {
                test[i] = true;
            }
        }
    }
}

class MyCounter {
    private int cnt = 0;

    public synchronized int getAndIncrement()
    {
        return cnt++;
    }
}

```

The expectation was that the test based on a synchronized method would be the slowest (*Parallel3* label on output), the one based on the *CombineTree* would become second (*Parallel1* label on output) and the fastest would be the one based on *AtomicInteger* (*Parallel2*); as the latest is likely to take advantage of hardware atomic instruction of 32bits. Surprisingly, the test based on *CombineTree* was the worst of all, by far:

```

.Parallel11, 8 threads, 1048576 tries took 3878
.Parallel12, 8 threads, 1048576 tries took 162
.Parallel13, 8 threads, 1048576 tries took 723

Time: 4.827

OK (3 tests)

```

Most likely we are not comparing apples with apples, as the theory does not match the experiment; perhaps the *CombineTest* class is not meant for real usage, so this comparison is not fair.

## 3 SynchronousDualQueueTest

### 3.1 Particular Case (problem)

The problem is to reduce the synchronization overhead of a synchronous queue.

### 3.2 Solution

The solution is given by using what is called a dual data structure; which splits the *enq* and *deq* operations in two parts. When a dequeuer tries to remove an item from the queue, it inserts a reservation object indicating that is waiting for an enqueueer. Later when an enqueueer thread realizes about the reservation, it fulfills the same by depositing an item and setting the reservation's flag. On the same way, an enqueueer thread can make another reservation when it wants to add an item and spin on its reservation's flag.

This solution has some nice properties:

- Waiting threads can spin on a locally cached flag (scalability).
- Ensures fairness in a natural way. Reservations are queued in the order they arrive.
- This data structure is linearizable, since each partial method call can be ordered when it is fulfilled.

### 3.3 Experiment Description

The test creates 16 threads, grouping them on enqueueers and dequeuers; for each group it divides the workload (512) evenly. Then each thread proceeds to either enqueue or dequeue, as many times as its share of the workload

indicated (512/8). There are some assertions to ensure that we do not repeat values (each dequeuer marks an array at the index of the value it got). among them.

### 3.4 Observations and Interpretations

The test runs normally most of the times, but in some occasions it hangs. During those cases we can see a null pointer exception, and some never ending enqueueers:

```
.parallel both
Exception in thread "Thread-9" java.lang.NullPointerException
    at queue.SynchronousDualQueueTest$DeqThread.run(SynchronousDualQueueTest.java:67)
2015-10-18 23:20:41
Full thread dump Java HotSpot(TM) 64-Bit Server VM (25.60-b23 mixed mode):

"Thread-12" #21 prio=5 os_prio=0 tid=0x00007f4314112800 nid=0x75cf runnable [0x00007f42fc980000]
    java.lang.Thread.State: RUNNABLE
        at queue.SynchronousDualQueue.enq(SynchronousDualQueue.java:49)
        at queue.SynchronousDualQueueTest$EnqThread.run(SynchronousDualQueueTest.java:60)

"Thread-8" #17 prio=5 os_prio=0 tid=0x00007f431410e800 nid=0x75cd runnable [0x00007f42fcb82000]
    java.lang.Thread.State: RUNNABLE
        at queue.SynchronousDualQueue.enq(SynchronousDualQueue.java:49)
        at queue.SynchronousDualQueueTest$EnqThread.run(SynchronousDualQueueTest.java:60)
```

The *NullPointerException* and the hanging enqueueers are related; the former appears cause the *deq* method returns null hence it fails to cast into integer for a dequeuer thread (leaving the overall situation unbalanced, as now nobody will consume some of the items being pushed into the queue).

The queue is populated with integers, so the *null* value returned must be either a defect of the tested class *SynchronousDualQueue*; or that the test case is badly designed. The easiest way to handle this is to modify the test to handle *null* values on dequeuers thread class:

```
class DeqThread extends Thread {
    public void run() {
        for (int i = 0; i < PER_THREAD; i++) {
            Object v = null;
            while ( (v = instance.deq()) == null );
            int value = (Integer) v;
            if (map[value]) {
```

```

        fail ("DeqThread: _duplicate _pop" );
    }
    map[ value ] = true;
}
}
}

```

With above modification, the test does not hang anymore.

## 4 LockFreeQueueRecycleTest

### 4.1 Particular Case (problem)

From the generic problem of improving coarse grained lock approaches, the particular approach followed on this exercise corresponds to the extreme case: lock-free data structure. The particular case is for a queue, and it has the additional bonus of recycling memory.

### 4.2 Solution

The solution is based on an 1996 ACM article from Maged M. Michel and Michael L. Scott, who based on previous publications, suggest a new way of implementing a lock-free queue with recycles. They implement the queue with a single-linked list using *tail* and *head* pointers; where *head* always points to a dummy or sentinel node which is the first in the list, and *tail* points to either the last or second to last node in the list. The algorithm uses “compare and swap” (*CAS*) with modification counters to avoid the ABA problem.

Dequeuers are allowed to free dequeued nodes by ensuring that *tail* does not point to the dequeued node nor to any of its predecessors (that is, dequeued nodes may be safely reused). To obtain consistent values of various pointers the authors relied on sequences of reads that re-check earlier values to ensure they have not changed (these reads are claimed to be simpler than snapshots).

### 4.3 Experiment Description

The test is pretty much the same described for *LockFreeQueueTest* (with the exception that it uses 8 threads instead of two).

### 4.4 Observations and Interpretations

The test does not exhibit any pitfall, which suggests that the theory works just fine on the tested machine. Sample output below:

```
.parallel enq
.parallel deq
.parallel both
.sequential push and pop
```

Time: 0.023

OK (4 tests)

## 5 BoundedQueueTest

### 5.1 Particular Case (problem)

The problem is to implement a bounded partial queue (that is, one which finite capacity which allows waits inside its methods).

### 5.2 Solution

The solution proposes to allow concurrency among enqueueers and dequeuers, by using two different locks (one for each group of threads). The code is small enough to be placed and explained in a bit more detail here, so let us give it a try:

```
1  public void enq(T x) {
2      if (x == null) throw new NullPointerException();
3      boolean mustWakeDequeuers = false;
4      enqLock.lock();
5      try {
6          while (size.get() == capacity) {
7              try {
8                  notFullCondition.await();
```

```

9         } catch (InterruptedException e) {}
10     }
11     Entry e = new Entry(x);
12     tail.next = e;
13     tail = e;
14     if (size.getAndIncrement() == 0) {
15         mustWakeDequeuers = true;
16     }
17 } finally {
18     enqLock.unlock();
19 }
20 if (mustWakeDequeuers) {
21     deqLock.lock();
22     try {
23         notEmptyCondition.signalAll();
24     } finally {
25         deqLock.unlock();
26     }
27 }
28 }

```

The main points of algorithm above are as follows (the code above was previously corrected, as it had several things inverted):

- It does not allow null values.
- There is a lock dedicated to *enq* method, which guarantees mutual exclusion to the critical section for enqueueers.
- It spins while the queue is full (the try/catch block was added just to allow compilation). This line was incorrectly asking whether queue was empty, while it needs to ask whether is still full.
- Once we know queue is room, we allocate the new node and update *tail* pointer.
- We atomically get and increment the atomic counter for the size (this was previously a decrement, which was incorrect).
- If the atomically retrieved previous value of size was zero, it means the queue was empty and there may had been waiting dequeuers; therefore we acquire the lock for *deq* method and inform all dequeuers that the queue has not at least an item.



The *deq* method is symmetric so we do not repeat same explanation (but worth to say that it had same inverted conditions and actions than *enq* method, in the code; so we needed to apply same corrective actions ... perhaps that was the purpose of the exercise?).

### 5.3 Experiment Description

The test is pretty much the same described for *LockFreeQueueTest*. (with the exception that it uses 8 threads instead of two).

### 5.4 Observations and Interpretations

The original code for the test made no sense, as conditions and actions for *enq* and *deq* methods were inverted; we did not even care testing such strange version that did not match at all the one published on the book (we also needed to fix the initialization of the *size* variable, which shall be zero instead of *capacity*). After the corrections mentioned on the Solution section above, the test ran normally. Sample output below:

```
.parallel both
.sequential push and pop
.parallel enq
.parallel deq

Time: 0.028

OK (4 tests)
```

Even if the original version of this test actually ran (we did not test it), it turns out quite counter-intuitive.

## 6 LazyListTest

### 6.1 Particular Case (problem)

The problem is that of gradually improving what coarse grain locking offers for concurrent data structures like sets (implemented with linked lists).

## 6.2 Solution

The *LazyList* solution is a refinement of the *OptimisticList* solution which does not lock while searching, but locks one it finds the interesting nodes (and then confirms that the locked nodes are correct). As one drawback of *OptimisticList* is that the most common method *contains* method locks, the next logical improvement is to make this method wait-free while keeping the *add* and *remove* methods locking (but reducing their transversings of the list from two to just one).

The refinement mentioned above is precisely that of *LazyList*, which adds a new bit to each node to indicate whether they still belong to the set or not (this prevents transversing the list to detect if the node is reachable, as the new bit introduces such invariant: if a transversing thread does not find a node or it is marked in this bit, then the corresponding item does not belong to the set. This behaviour implies that *contains* method does a single wait-free transversal of the list.

For adding an element to the list, *add* method traverses the list, locks the target's predecessor, and inserts the node. The *remove* method is lazy (hence the name of the solution), as it splits its task in two parts: first marks the node in the new bit, logically removing it; and second, update its predecessor's next field, physically removing it.

The three methods ignore the locks while transversing the list, possibly passing over both logically and physically deleted nodes. The *add* and *remove* methods still lock *pred* and *curr* nodes as with *OptimisticList* solution, but the validation reduces to check that *curr* node has not been marked; as well as validating the same for *pred* node, and that it still points to *curr* (validating a couple of nodes is much better than transversing whole list though). Finally, the introduction of logical removals implies a new contract for detecting that an item still belongs to set: it does so, if still referred by an unmarked reachable node.

## 6.3 Experiment Description

The test is pretty much the same described for *LockFreeQueueTest* (with the exception that it uses 8 threads instead of two).

## 6.4 Observations and Interpretations

The test works as expected, below sample execution:

```
.parallel deq  
.parallel both  
.sequential push and pop  
.parallel enq
```

Time: 0.03

OK (4 tests)