



Multiprocessor Programming Course

Dario Bahena Tapia - Ricardo Zavaleta Vazquez - Isai Barajas Cicourel
December 2015

Contents

1	Mutual Exclusion	2
1.1	LockFreeQueueTest	3
1.1.1	Particular Case (problem)	3
1.1.2	Solution	3
1.1.3	Experiment Description	4
1.1.4	Observations and Interpretations	4
1.1.5	Proposed changes to fix the problems	8
1.2	Peterson Algorithm	13
1.2.1	Particular Case	13
1.2.2	Solution	13
1.2.3	Experiment Description	14
1.2.4	Sample Results	14
1.2.5	Interpretation	14
1.3	Bakery Test	15
1.3.1	Particular Case	15
1.3.2	Solution	15
1.3.3	Experiment Description	17
1.3.4	Sample Results	17
1.3.5	Interpretation	17
1.3.6	Proposed solution	18
1.4	Filter Lock	19
1.4.1	Particular Case	19
1.4.2	Solution	19
1.4.3	Experiment Description	19
1.4.4	Observations and Interpretations	19
1.4.5	Proposed changes to fix the problem	20
1.4.6	Proposed solution	20

4.1	Queue Lock	76
4.1.1	Particular Case	76
4.1.2	Solution	76
4.1.3	Experiment Description	77
4.1.4	Observations and Interpretations	77
2	Foundations of Shared Memory	23
2.1	Safe Boolean MRSW Register	24
2.1.1	Particular Case	24
2.1.2	Solution	24
2.1.3	Experiment Description	24
2.1.4	Sample Results	25
2.1.5	Interpretation	26
2.2	RegMRSWRegisterTest	27
2.2.1	Particular Case (problem)	27
2.2.2	Solution	27
2.2.3	Experiment Description	28
2.2.4	Observations and Interpretations	28
2.3	Regular Boolean MRSW Register	29
2.3.1	Particular Case	29
2.3.2	Solution	29
2.3.3	Experiment Description	30
2.3.4	Observations and Interpretations	30
2.4	Sequential Snapshot	31
2.4.1	Particular Case	31
2.4.2	Solution	31
2.4.3	Experiment Description	32
2.4.4	Observations and Interpretations	32
2.4.5	Proposed changes to fix the problem	32
2.4.6	Proposed solution	33
2.5	Atomic MRSW Register	34
2.5.1	Particular Case	34
2.5.2	Solution	34
2.5.3	Experiment Description	35
2.5.4	Sample Results	35
2.5.5	Interpretation	35
2.6	AtomicMRMWRegisterTest	36
2.6.1	Particular Case (problem)	36

2.6.2	Solution	36
2.6.3	Experiment Description	37
2.6.4	Observations and Interpretations	38
2.7	SimpleSnapshotTest	39
2.7.1	Particular Case (problem)	39
2.7.2	Solution	39
2.7.3	Experiment Description	40
2.7.4	Observations and Interpretations	41
2.8	Wait-Free Snapshots	42
2.8.1	Particular Case	42
2.8.2	Solution	42
2.8.3	Experiment Description	44
2.8.4	Sample Results	45
2.8.5	Interpretation	45
3	Spin Locks and Contention	46
3.1	Problem Definition	46
3.2	Test-And-Set Locks	47
3.2.1	Particular Case	47
3.2.2	Solution	47
3.2.3	Experiment Description	48
3.2.4	Sample Results	48
3.2.5	Interpretation	49
3.3	TTASLockTest	50
3.3.1	Particular Case (problem)	50
3.3.2	Solution	50
3.3.3	Experiment Description	50
3.3.4	Observations and Interpretations	51
3.4	BackoffLockTest	52
3.4.1	Particular Case (problem)	52
3.4.2	Solution	52
3.4.3	Experiment Description	54
3.4.4	Observations and Interpretations	54
3.5	CLH Queue Lock	55
3.5.1	Particular Case	55
3.5.2	Solution	55
3.5.3	Experiment Description	55
3.5.4	Sample Results	56

3.5.5	Interpretation	56
3.6	MCSLockTest	58
3.6.1	Particular Case (problem)	58
3.6.2	Solution	58
3.6.3	Experiment Description	59
3.6.4	Observations and Interpretations	60
3.7	Array Lock	64
3.7.1	Particular Case	64
3.7.2	Solution	64
3.7.3	Experiment Description	65
3.7.4	Observations and Interpretations	65
3.7.5	Proposed changes to fix the problem	65
3.7.6	Proposed solution	66
3.8	Composite Fast Path Lock	67
3.8.1	Particular Case	67
3.8.2	Solution	67
3.8.3	Experiment Description	67
3.8.4	Observations and Interpretations	67
3.9	Time Out Lock	68
3.9.1	Particular Case	68
3.9.2	Solution	68
3.9.3	Experiment Description	68
3.9.4	Proposed changes to improve execution	69
3.9.5	Observations and Interpretations	69
3.10	CompositeLockTest	70
3.10.1	Particular Case (problem)	70
3.10.2	Solution	70
3.10.3	Experiment Description	70
3.10.4	Observations and Interpretations	71
3.11	Hierarchical Backoff Lock	72
3.11.1	Particular Case	72
3.11.2	Solution	72
3.11.3	Experiment Description	73
3.11.4	Sample Results	74
3.11.5	Interpretation	74

4	Monitors and Blocking Synchronization	75
4.1	Queue Lock	76
4.1.1	Particular Case	76
4.1.2	Solution	76
4.1.3	Experiment Description	77
4.1.4	Observations and Interpretations	77
4.2	Non-Reentrant Mutual Exclusion	78
4.2.1	Particular Case	78
4.2.2	Solution	78
4.2.3	Experiment Description	79
4.2.4	Observations and Interpretations	79
4.3	Count Down Latch	80
4.3.1	Particular Case	80
4.3.2	Solution	80
4.3.3	Experiment Description	81
4.3.4	Sample Results	83
4.3.5	Interpretation	83
5	Linked Lists: The Role of Locking	85
5.1	Coarse List	86
5.1.1	Particular Case	86
5.1.2	Solution	86
5.1.3	Experiment Description	87
5.1.4	Sample Results	88
5.1.5	Proposed solution	89
5.1.6	Interpretation	90
5.2	Fine-Grained Synchronization List	91
5.2.1	Particular Case	91
5.2.2	Solution	91
5.2.3	Experiment Description	92
5.2.4	Observations and Interpretations	92
5.3	Optimistic Synchronization	94
5.3.1	Particular Case	94
5.3.2	Solution	94
5.3.3	Experiment Description	95
5.3.4	Observations and Interpretations	95
5.3.5	Proposed changes to fix the problem	96
5.3.6	Proposed solution	97

5.4	LazyListTest	98
5.4.1	Particular Case (problem)	98
5.4.2	Solution	98
5.4.3	Experiment Description	100
5.4.4	Observations and Interpretations	100
5.5	Lock-Free List	102
5.5.1	Particular Case	102
5.5.2	Solution	102
5.5.3	Experiment Description	102
5.5.4	Sample Results	102
5.5.5	Interpretation	103
6	Concurrent Queues and the ABA Problem	104
6.8	LockFreeQueueTest Test	124
6.8.1	Particular Case	124
6.8.2	Solution	124
6.8.3	Experiment Description	125
6.8.4	Sample Results	126
6.8.5	Interpretation	126
6.8.6	Proposed solution	126
6.2	HW Queue	108
6.2.1	Particular Case	108
6.2.2	Solution	108
6.2.3	Experiment Description	108
6.2.4	Observations and Interpretations	109
6.9	Bounded Partial Queue	127
6.9.1	Particular Case	127
6.9.2	Solution	127
6.9.3	Experiment Description	129
6.9.4	Observations and Interpretations	129
6.4	Unbounded Queue	113
6.4.1	Particular Case	113
6.4.2	Solution	113
6.4.3	Experiment Description	114
6.4.4	Observations and Interpretations	114
6.5	SynchronousDualQueueTest	115
6.5.1	Particular Case (problem)	115
6.5.2	Solution	115

6.5.3	Experiment Description	115
6.5.4	Observations and Interpretations	116
6.6	SynchronousQueueTest Test	118
6.6.1	Particular Case	118
6.6.2	Solution	118
6.6.3	Experiment Description	120
6.6.4	Sample Results	120
6.6.5	Interpretation	121
6.6.6	Proposed solution	121
6.7	LockFreeQueueRecycleTest	122
6.7.1	Particular Case (problem)	122
6.7.2	Solution	122
6.7.3	Experiment Description	122
6.7.4	Observations and Interpretations	123
6.8	LockFreeQueueTest Test	124
6.8.1	Particular Case	124
6.8.2	Solution	124
6.8.3	Experiment Description	125
6.8.4	Sample Results	126
6.8.5	Interpretation	126
6.8.6	Proposed solution	126
6.9	Bounded Partial Queue	127
6.9.1	Particular Case	127
6.9.2	Solution	127
6.9.3	Experiment Description	129
6.9.4	Observations and Interpretations	129
7	Concurrent Stacks and Elimination	130
7.1	Exchanger	131
7.1.1	Particular Case	131
7.1.2	Solution	131
7.1.3	Experiment Description	134
7.1.4	Sample Results	134
7.1.5	Interpretation	134
8	Counting, Sorting and Distributed Coordination	135
8.1	TreeTest	136
8.1.1	Particular Case (problem)	136

8.1.2	Solution	136
8.1.3	Experiment Description	136
8.1.4	Observations and Interpretations	137
8.2	Balancer Test	139
8.2.1	Particular Case	139
8.2.2	Solution	139
8.2.3	Experiment Description	139
8.2.4	Sample Results	140
8.2.5	Interpretation	140
8.2.6	Proposed solution	140
8.3	Bitonic Counting Network	141
8.3.1	Particular Case	141
8.3.2	Solution	141
8.3.3	Experiment Description	142
8.3.4	Observations and Interpretations	142
9	Concurrent Hashing and Natural Parallelism	144
9.1	CoarseHashSetTest Test	145
9.1.1	Particular Case	145
9.1.2	Solution	145
9.1.3	Experiment Description	146
9.1.4	Sample Results	147
9.1.5	Interpretation	147
9.1.6	Proposed solution	147
9.2	StripedHashSetTest Test	148
9.2.1	Particular Case	148
9.2.2	Solution	148
9.2.3	Experiment Description	149
9.2.4	Sample Results	149
9.2.5	Interpretation	149
9.2.6	Proposed solution	149
9.3	CoarseCuckooHashSetTest	150
9.3.1	Particular Case (problem)	150
9.3.2	Solution	150
9.3.3	Experiment Description	151
9.3.4	Observations and Interpretations	152
9.4	LockFreeHashSetTest	153
9.4.1	Particular Case (problem)	153

9.4.2	Solution	153
9.4.3	Experiment Description	153
9.4.4	Observations and Interpretations	153
9.5	Cuckoo Hashing	155
9.5.1	Particular Case	155
9.5.2	Solution	155
9.5.3	Experiment Description	156
9.5.4	Observations and Interpretations	157
9.5.5	Proposed changes to fix the problem	157
9.5.6	Proposed solution	157
9.6	Refinable Hash Set	158
9.6.1	Particular Case	158
9.6.2	Solution	158
9.6.3	Experiment Description	159
9.6.4	Observations and Interpretations	160
9.6.5	Proposed changes to fix the problem	160
9.6.6	Proposed solution	160
9.7	RefinableCuckooHashSetTest Test	161
9.7.1	Particular Case	161
9.7.2	Solution	161
9.7.3	Experiment Description	162
9.7.4	Sample Results	163
9.7.5	Interpretation	163
9.7.6	Proposed solution	163
9.8	StripedCuckooHashSetTest	164
9.8.1	Particular Case (problem)	164
9.8.2	Solution	164
9.8.3	Experiment Description	165
9.8.4	Observations and Interpretations	165
10	Skiplists and Balanced Search	166
10.1	Double Ended Queue	167
10.1.1	Particular Case	167
10.1.2	Solution	167
10.1.3	Experiment Description	167
10.1.4	Observations and Interpretations	168
10.2	Parallel Matrix Multiplication	169
10.2.1	Particular Case	169

10.2.2	Solution	169
10.2.3	Experiment Description	170
10.2.4	Observations and Interpretations	170
10.2.5	Proposed changes to fix the problem	170
10.3	Unbounded Double Ended Queue	171
10.3.1	Particular Case	171
10.3.2	Solution	171
10.3.3	Experiment Description	172
10.3.4	Observations and Interpretations	172
10.3.5	Proposed changes to fix the problem	172
10.4	ArraySumTest	173
10.4.1	Particular Case (problem)	173
10.4.2	Solution	173
10.4.3	Experiment Description	173
10.4.4	Observations and Interpretations	173
10.5	FibTaskTest	175
10.5.1	Particular Case (problem)	175
10.5.2	Solution	175
10.5.3	Experiment Description	175
10.5.4	Observations and Interpretations	176
10.6	MMThreadTest	177
10.6.1	Particular Case (problem)	177
10.6.2	Solution	177
10.6.3	Experiment Description	178
10.6.4	Observations and Interpretations	178
11	Priority Queues	179
11.1	BitReversedCounterTest	180
11.1.1	Particular Case (problem)	180
11.1.2	Solution	180
11.1.3	Experiment Description	180
11.1.4	Observations and Interpretations	181
11.2	SimpleTreeTest Test	183
11.2.1	Particular Case	183
11.2.2	Solution	183
11.2.3	Experiment Description	184
11.2.4	Sample Results	185
11.2.5	Interpretation	185

11.2.6	Proposed solution	185
11.3	Sequential Heap	186
11.3.1	Particular Case	186
11.3.2	Solution	186
11.3.3	Experiment Description	187
11.3.4	Observations and Interpretations	187
11.4	Skiplist-Based Unbounded Priority Queue	188
11.4.1	Particular Case	188
11.4.2	Solution	188
11.4.3	Experiment Description	190
11.4.4	Observations and Interpretations	191
11.4.5	Proposed changes to fix the problem	191
11.5	FineGrainedHeapTest Test	193
11.5.1	Particular Case	193
11.5.2	Solution	193
11.5.3	Experiment Description	196
11.5.4	Sample Results	196
11.5.5	Interpretation	196
11.5.6	Proposed solution	196
11.6	SimpleLinearTest	197
11.6.1	Particular Case (problem)	197
11.6.2	Solution	197
11.6.3	Experiment Description	198
11.6.4	Observations and Interpretations	198
12	Futures, Scheduling and Work Distribution	199
12.1	ArraySumTest	200
12.1.1	Particular Case (problem)	200
12.1.2	Solution	200
12.1.3	Experiment Description	200
12.1.4	Observations and Interpretations	201
12.2	FibTaskTest	202
12.2.1	Particular Case (problem)	202
12.2.2	Solution	202
12.2.3	Experiment Description	203
12.2.4	Observations and Interpretations	203
12.3	MMThreadTest	204
12.3.1	Particular Case (problem)	204

12.3.2	Solution	204
12.3.3	Experiment Description	205
12.3.4	Observations and Interpretations	205
13	Barriers	206
13.1	DisBarrierTest	207
13.1.1	Particular Case (problem)	207
13.1.2	Solution	207
13.1.3	Experiment Description	209
13.1.4	Observations and Interpretations	209
13.2	SenseBarrierTest	210
13.2.1	Particular Case (problem)	210
13.2.2	Solution	210
13.2.3	Experiment Description	211
13.2.4	Observations and Interpretations	211
13.3	Reusable Barrier	213
13.3.1	Particular Case	213
13.3.2	Solution	213
13.3.3	Experiment Description	214
13.3.4	Observations and Interpretations	214
13.4	TreeBarrierTest	215
13.4.1	Particular Case (problem)	215
13.4.2	Solution	215
13.4.3	Experiment Description	217
13.4.4	Observations and Interpretations	217
14	Transactional Memory	218
14.1	Restricted Bloom Filter	219
14.1.1	Particular Case	219
14.1.2	Solution	219
14.1.3	Experiment Description	219
14.1.4	Observations and Interpretations	219
14.2	Lazy Skiplist	220
14.2.1	Particular Case	220
14.2.2	Solution	220
14.2.3	Experiment Description	220
14.2.4	Observations and Interpretations	220
14.3	ListTest (ofree)	222

14.3.1	Particular Case (problem)	222
14.3.2	Solution	222
14.3.3	Experiment Description	223
14.3.4	Observations and Interpretations	223

Chapter 1

Mutual Exclusion

1.1 LockFreeQueueTest

1.1.1 Particular Case (problem)

This is a particular case of the mutual exclusion problem, where the shared resource is a queue.

1.1.2 Solution

In order to guarantee the correctness of the multi-threaded access to the queue, it implements a lock free scheme on its *enq* and *deq* methods by putting waits before modifying the state of the queue. It does it incorrectly though, as we will see later.

```
1  class LockFreeQueue {
2      public int head = 0;    // next item to dequeue
3      public int tail = 0;    // next empty slot
4      Object[] items; // queue contents
5      public LockFreeQueue(int capacity) {
6          head = 0; tail = 0;
7          items = new Object[capacity];
8      }
9
10     public void enq(Object x) {
11         // spin while full
12         while (tail - head == items.length) {}; // spin
13         items[tail % items.length] = x;
14         tail++;
15     }
16
17     public Object deq() {
18         // spin while empty
19         while (tail == head) {}; // spin
20         Object x = items[head % items.length];
21         head++;
22         return x;
23     }
24 }
```


1.1.3 Experiment Description

The test program *LockFreeQueueTest* includes the following individual test cases; the parallel degree is two threads for each operation (queue or dequeue), with a `TEST_SIZE` of 512 (number of items to enqueue and dequeue) which is spread evenly among the two threads on each group:

- *testSequential*: calls *enq* method as many times as `TEST_SIZE`, and later calls *deq* method the same number of times checking that the FIFO order is preserved.
- *testParallelEnq*: enqueues in parallel (two threads) but dequeues sequentially.
- *testParallelDeq*: enqueues sequentially but dequeues in parallel (two threads)
- *testParallelBoth*: enqueues and dequeues in parallel (with a total of four threads, two for each operation).

The test program was run on two types of machines; one with two cores only and another with 24 cores.

1.1.4 Observations and Interpretations

The bottom line of this exercise is most likely, to show several types of problems that can occur if we do not use mutual exclusion; the queue implementation fails implementing it, making the test vulnerable to several cases of race conditions. Below we explain some of them.

Non atomic dequeue: referencing invalid registers

The symptom for this problem is a `NullPointerException`, and it occurred on both test machines:

```
1) testParallelEnq(mutex.LockFreeQueueTest)java.lang.NullPointerException
at mutex.LockFreeQueueTest.testParallelEnq(LockFreeQueueTest.java:67)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
```

where the offending line is a cast to an Integer from the value returned by the *deq* method; this implies that such function is returning *null*. A possible scenario to produce this outcome is as follows. Prefixes of T1 or T2 indicate the thread running the action, and they refer to the line numbers of the code of method *deq* posted above.

- Assume that queue holds a single element which lives in *items* array at position 0; the array has *null* on position 1 (per initialization).
- T1: Executes method up to line 19.
- T2: Executes method up to line 19.
- T1: Executes lines 20, getting *items[0]*.
- T1: Executes lines 21, increasing *head* to 1.
- T2: Executes line 20, but as *head* was changed already, it gets *items[1]*.
- T2: Executes lines 21, increasing *head* to 2.
- T1: Executes line 22 returning *items[0]*
- T2: Executes line 22 returning *items[1]*, which was *null*.

The problem with above interlacing derives from the fact that the *deq* method is not an atomic operation, hence it allowed both threads to enter into the critical section and compete for updating the shared variables.

Non atomic dequeue: returning duplicate values

The symptom for this problem is a duplicate pop warning, and it occurred on both test machines:

```
.parallel deq
Exception in thread "Thread-5" junit.framework.AssertionFailedError: DeqThread: duplicate pop
at junit.framework.Assert.fail(Assert.java:57)
at junit.framework.TestCase.fail(TestCase.java:227)
at mutex.LockFreeQueueTest$DeqThread.run(LockFreeQueueTest.java:129)
```

where the offending line is an assertion that validates that nobody else has pop such value from the queue; as the threads which populate do have non overlapping ranges of values, the pop operations (*deg*) shall never return a duplicate one. But duplication is possible indeed, if we have a sequence like the one below between the two threads:

- Assume that queue holds a single element which lives in *items* array at position 0.
- T1: Executes method up to line 19.
- T2: Executes method up to line 19.
- T1: Executes line 20, getting *items[0]*.
- T2: Executes line 20, getting as well *items[0]*.
- T1: Executes line 21, setting *head* to 1.
- T2: Executes line 21, setting *head* to 2.
- T1: Executes line 22 returning *items[0]*.
- T2: Executes line 22 returning *items[0]*.

Not only we left the queue in an inconsistent state (*head* has incorrect value), but we also returned the same element twice, triggering then the violation on the test. The underlying problem is the same as previous case: lack of atomicity of the *deg* method.

Non atomic auto-increment: loosing values

The symptom for this problem is a never ending program, hanging on either the test *testParallelBoth* or *testParallelEnq*; this issue occurred on both test machines (though it was easier to reproduce on the one with two cores). When produced several thread dumps of the Java program, we can see either two hanging threads (*testParallelBoth* case):

```

"Thread-7" #16 prio=5 os_prio=0 tid=0x00007f3140102000 nid=0x3f51
  runnable [0x00007f31226e9000]
  java.lang.Thread.State: RUNNABLE
    at mutex.LockFreeQueue.deq(LockFreeQueue.java:33)
    at mutex.LockFreeQueueTest$DeqThread.run(LockFreeQueueTest.java:141)

"main" #1 prio=5 os_prio=0 tid=0x00007f3140009800 nid=0x3f3c in Object.wait()
  [0x00007f3148382000]
  java.lang.Thread.State: WAITING (on object monitor)
    at java.lang.Object.wait(Native Method)
    - waiting on <0x00000000d6effea8> (a mutex.LockFreeQueueTest$DeqThread)
    at java.lang.Thread.join(Thread.java:1245)
    - locked <0x00000000d6effea8> (a mutex.LockFreeQueueTest$DeqThread)
    at java.lang.Thread.join(Thread.java:1319)
    at mutex.LockFreeQueueTest.testParallelBoth(LockFreeQueueTest.java:123)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:497)
    at junit.framework.TestCase.runTest(TestCase.java:164)
    at junit.framework.TestCase.runBare(TestCase.java:130)

```

The main thread is waiting for a dequeue thread to finish; but that one is on an infinite loop at line 19 of *deq* method (line numbers per our listing in this document, not in the file). This means that we have lost some of the inserted elements in queue, and that one of the dequeue threads will never finish; as they expect each to pop a fixed amount of elements per the test.

The other hanging scenario that comes out of (*testParallelEnq*), is a single hanging thread, which was actually the main thread doing a serial dequeue. The common factor for both scenarios was a parallel enqueue.

But the amount of times we request a dequeue operation, among all threads, is the same as the number of elements we queued; how come we end up losing some of those? One possible explanation is again, the lack of atomicity but this time of the *enq* method; to be more specific, the lack of atomicity of its auto-increment operation *tail++* (line 14). Let us remember that the auto-increment operator in Java is nothing but syntactic sugar for the following sequence of operations (when applied to *tail* variable):

```

tmp = tail;
tmp = tmp + 1;
tail = tmp;

```

If two threads execute the lower level operations above, we can see how they can end up losing increments in the shared variable *tail*; the following sequence is an example of such scenario:

- T1: executes $tmp = tail$.
- T2: executes $tmp = tail$.
- T1: executes $tmp = tmp + 1$.
- T2: executes $tmp = tmp + 1$.
- T1: executes $tail = tmp$.
- T2: executes $tail = tmp$.

We can appreciate that in the above interlacing, the final value of the shared variable is $tail + 1$; instead of the expected value of $tail + 2$. It would be enough to lose a single value this way, in order to make the enqueue threads think that they inserted the total of 512, while they really inserted 511; as each dequeue thread will try to pop 256 each, only one of them will be able to finish while the other will get blocked after having removed 255 entries. That is most likely the explanation for the hung threads we pasted above ¹; the solution is again to really implement mutual exclusion around the methods *deq* / *enq*; in such a way that they become atomic operations.

1.1.5 Proposed changes to fix the problems

In class it was mentioned that the *head* and *tail* variables were not declared as *volatile*; which in Java jargon means they are susceptible to caching on each core. This means, that a write to any of those two variables (which control the queue size) are not meant to be reflected immediately to the other cores, unless we declare the variable as *volatile*. We did several attempts to achieve lock-free implementation of this queue, by making the variables *volatile* but it was enough; then we tried by making atomic only portions of the *enq* and *deq* methods; but it did not work either.

¹Note that it does not matter that the array *items* has populated all the correct entries, because the flow is controlled by the counters *tail* and *head*.

All the three scenarios described before can be eliminated, if we make the methods *enq* and *deq* atomic; this can be easily achieved in Java by making them *synchronized*. However, by doing that, we will loose parallelism among the two groups of threads (those calling *enq* and those calling *deq*); this is because the *synchronized* keyword uses as lock the whole object, so at any moment in time, only one synchronized method can actually run within any object. In order to overcome this limitation, we can use synchronized blocks against two different lock objects (one for each operation).

Even with the changes above, we can still have issues; what if the very first thread running is one calling *deq* method? It will find the queue empty and loop forever. In order to prevent that, we should remove the waiting operation out of the queue methods, and put them in the test code itself. This is because, on the cited scenario that we try to dequeue with an empty queue, we would expect to simply try again (giving the chance to parallel enqueue threads to produce something for us to pop). The final code which incorporates these fixes is listed below:

```
class LockFreeQueue {
    private static Object enqLock = new Object();
    private static Object deqLock = new Object();

    public int head = 0;    // next item to dequeue
    public int tail = 0;    // next empty slot
    Object[] items; // queue contents
    public LockFreeQueue(int capacity) {
        head = 0; tail = 0;
        items = new Object[capacity];
    }

    public boolean enq(Object x) {
        synchronized(enqLock)
        {
            if (tail - head == items.length) {
                return false;
            }
            items[tail % items.length] = x;
            tail++;
            return true;
        }
    }
}
```

```

public Object deq() {
    synchronized(deqLock)
    {
        if (tail == head) {
            return null;
        }
        Object x = items[head % items.length];
        head++;
        return x;
    }
}

```

The test code was also modified, to make the enqueue and dequeue threads to iterate until they have successfully called their respective methods 256 times (total size of queue divided by number of threads). A successful call is one that does not return *false* nor *null*. The modified code was executed ten thousand times and it did not produce any of the original problems we explained. Though not a formal proof of its correctness, it is a good indication of the same (the original program produced one of the cited problems quite often, usually within 10 executions).

A probably more elegant solution we came out with, after reading about Java facilities for locks and conditions, was the following; it has the advantage that it does not require us to change the test case:

```

class LockFreeQueue {
    int head = 0;    // next item to dequeue
    int tail = 0;    // next empty slot
    Object[] items; // queue contents

    final Lock lock = new ReentrantLock();
    final Condition notFull = lock.newCondition();
    final Condition notEmpty = lock.newCondition();

    public LockFreeQueue(int capacity) {
        head = 0; tail = 0;
        items = new Object[capacity];
    }

    public void enq(Object x) {
        lock.lock();
        try {
            // spin while full
            while (tail - head == items.length)
                notFull.await();
            items[tail % items.length] = x;
            tail++;
            notEmpty.signalAll();
        }
        catch (InterruptedException e) {
            throw new RuntimeException(e);
        }
        finally {
            lock.unlock();
        }
    }

    public Object deq() {
        lock.lock();
        try {
            // spin while empty
            while (tail == head)
                notEmpty.await();
            Object x = items[head % items.length];
            head++;
            notFull.signalAll();
            return x;
        }
        catch (InterruptedException e) {
            throw new RuntimeException(e);
        }
    }
}

```



```

    }
    finally {
        lock.unlock();
    }
}
}

```

This second option is interesting on its usage of the *await* call to the condition objects; such call provokes that a thread which got granted a lock releases it to give a chance to other threads to come in. For example, an enqueueer may give a chance to a dequeuer to make room for it; or, a dequeuer may give an enqueueer a chance to produce something for it. Without those calls to *await*, we could end up falling into never ending loops.

Another relevant detail of this second solution is the usage of *signalAll* method on conditions; this is to awake the threads which were waiting on that particular condition. In our case, it corresponds to the enqueueers or dequeuers which called *await* on their respective *notFull* or *notEmpty* conditions. Thus, as soon as there is room an enqueueer is awakened and as soon as there is data a dequeuer is awakened. As there are tests where we can have more than one thread on each condition, we prefer to use *signalAll* method rather than *signal* (which will awake a single thread). This is to prevent the lost awakes problem that is mentioned on the book (on the chapter about monitors, if we recall correctly).

1.2 Peterson Algorithm

1.2.1 Particular Case

In this experiment, the particular case we are trying to study is again mutual exclusion where two threads try to use a shared resource.

We are trying to do so in such a way that there is no starvation.

1.2.2 Solution

During the lecture we discussed that Peterson algorithm takes the best of other two algorithms (which we called the LockOne and the LockTwo) to create a better one. The LockOne class that we discussed used an approach where mutual exclusion was attempted using a flag which indicates that a given thread has the intention of using a resource. We concluded that it satisfies the mutual exclusion but it can potentially dead lock unless executions are not interleaved.

On the other hand, the LockTwo class achieved mutual exclusion by forcing the threads let the other one execute first. However, this approach also had dead lock problems when one thread ran completely before the other one. So we observe that both algorithms complement each other.

The Peterson algorithm combines these two methods to achieve a starvation-free algorithm. It uses both the *flags* approach and the *victimization* approach.

Here we show the *lock()* and *unlock()* methods:

```
1  private boolean[] flag = new boolean[2];
2  private int victim;
3
4  public void lock() {
5      int i = ThreadID.get();
6      int j = 1-i;
7      flag[i] = true;
8      victim = i;
9      while (flag[j] && victim == i) {} // spin
10 }
11
12 public void unlock() {
13     int i = ThreadID.get();
14     flag[i] = false;
15 }
```

1.2.3 Experiment Description

Now let's discuss how the proposed test case excersices the Peterson algorithm.

The test creates 2 threads that need to be coordinate in order to increment a common counter. Both threads have to cooperate to increase this counter from 0 to 1024. Each of the threads will increase by one the counter 512 times. The expected result is that regardless of the order in which each thread executes the increment, at the end the counter must stay at 1024. If that is not the case, then it means that mutual exclusion did not work.

1.2.4 Sample Results

This experiment was successfull in all tries. This is the output of the test:

```
.parallel
Time: 0.01
OK (1 test)
```

1.2.5 Interpretation

The results in the proposed system were all OK. In all cases, the threads were able to cooperate to increase the counter to 1024.

One thing that is worth mentioning is that this algorihtm has a limitation. The limitation is that it only works for two threads. In a later experiment we will discover other algorithms that allow the coordination of more that two threads.

1.3 Bakery Test

1.3.1 Particular Case

In this experiment, we are dealing with the problem of mutual exclusion with a number of participants > 2 .

We have already seen that Peterson Algorithm works for two threads. It guarantees mutual exclusion and is starvation free. Therefore, the algorithm is deadlock free. Now let us focus in an algorithm that can be used to coordinate more threads.

1.3.2 Solution

The algorithm that we will play with is called *Bakery*. Its name comes from the fact that it is similar to the protocol used in bakeries where one enters the store and picks a number that indicates the order in which each client will be attended.

In the algorithm, the fact of entering the store is done by setting a flag. After that, the thread calculates the next number in the machine.

The algorithm then chooses the next number to be attended. When that happens, the thread is allowed to enter the critical section.

One thing that we ought to mention is that threads can have the same number assigned. If that is the case, the next in the line is decided using a lexicographical ordering of the threads ids.

Here we show the interesting methods used in this algorithm:

```
1  public void lock() {
2      int me = ThreadID.get();
3      flag[me] = true;
4      int max = Label.max(label);
5      label[me] = new Label(max + 1);
6      while (conflict(me)) {}; // spin
7  }
8
9  public void unlock() {
10     flag[ThreadID.get()] = false;
11 }
12
13 private boolean conflict(int me) {
14     for (int i = 0; i < label.length; i++) {
```

```

15         if (i != me && flag[i] && label[me].compareTo(label[i]) < 0) {
16             return true;
17         }
18     }
19     return false;
20 }
21
22 static class Label implements Comparable<Label> {
23     int counter;
24     int id;
25     Label() {
26         counter = 0;
27         id = ThreadID.get();
28     }
29     Label(int c) {
30         counter = c;
31         id = ThreadID.get();
32     }
33     static int max(Label[] labels) {
34         int c = 0;
35         for (Label label : labels) {
36             c = Math.max(c, label.counter);
37         }
38         return c;
39     }
40
41     public int compareTo(Bakery.Label other) {
42         if (this.counter < other.counter
43             || (this.counter == other.counter && this.id < other.id)) {
44             return -1;
45         } else if (this.counter > other.counter) {
46             return 1;
47         } else {
48             return 0;
49         }
50     }
51 }

```

In the *lock()* method, we observe that the thread signals its intention of accessing the critical section by setting its flag to true. After that, the thread picks its number and then it sits and waits for its turn. The thread has to wait if another thread which announced its intention to use the resource has a smaller label.

1.3.3 Experiment Description

Now let us explain how the experiments work. In this case we have 8 threads increasing a counter from 0 to 1024. Each thread will increase the counter 128 times. At the end, the counter must remain in 1024. If that is not the case, then there is a problem with the mutual exclusion algorithm.

1.3.4 Sample Results

In this exercise we observed that in the machine that we have been using, the test for the Bakery algorithm fails consistently with the following error:

```
.F
Time: 0.019
There was 1 failure:
1) testParallel(mutex.BakeryTest)junit.framework.AssertionFailedError:
expected:<1019> but was:<1024>
    at mutex.BakeryTest.testParallel(BakeryTest.java:47)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at
sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
    at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)

FAILURES!!!
Tests run: 1, Failures: 1, Errors: 0
```

1.3.5 Interpretation

Let us elaborate on the meaning of the error that we got. The assertion failure says that at the end of the algorithm, our counter didn't reach the expected 1024. This problem might indicate that our counter wasn't actually being protected by the lock because it looks like multiple threads found the counter with a value of i and increased it to $i + 1$. This is wrong.

After investigating the problem in more detail, we found out that the problem in this code has to do with the way the *label* array is accessed and modified. For example, in the lock method, multiple threads can reach the point where the max is calculated and hence these multiple threads can take the same turn. The algorithm takes care of this problem by introducing a lexicographical ordering in the comparison: if two threads have the same number, the decision is made based on the *threadID*.

Also, as explained in the textbook, we cannot be sure that we have a correct memory consistency: it is possible that the processor modified the order

in which operations in the *lock()* method are executed. It is also possible that the processor decided to execute an instruction before another.

1.3.6 Proposed solution

The following fragment of code shows one possible way of fixing the code:

```
1  public void lock() {
2      int me = ThreadID.get();
3      flag[me] = true;
4      synchronized (label){
5          int max = Label.max(label);
6          label[me] = new Label(max + 1);
7          while (conflict(me)) {}; // spin
8      }
9  }
```

Note that this solution is not quite satisfactory as we are basically using a Java mechanism to ensure synchronized access in an algorithm that promises to provide synchronized access.

In any case, let us show the output of the program after this fix:

```
.
Time: 0.015

OK (1 test)
```

After the fix the results in the proposed system were all OK. In all cases, the threads were able to cooperate to increase the counter to 1024.

One thing that is worth mentioning is that unlike the Peterson Algorithm, for example, the Bakery algorithm is able to coordinate multiple threads (at least in theory). This fact was shown in the test case where 8 threads were coordinated to increase the counter.

1.4 Filter Lock

1.4.1 Particular Case

In this experiment, the particular case we are trying to study is mutual exclusion where several threads try to use a shared resource.

We are trying to do so using waiting rooms called levels that a thread must traverse before acquiring the lock.

1.4.2 Solution

We now consider two mutual exclusion protocols that work for n threads, where n is greater than 2. The Filter lock, is a direct generalization of the Peterson lock to multiple threads.

The Filter lock, creates $n - 1$ “waiting rooms”, called levels, that a thread must traverse before acquiring the lock. The Peterson lock uses a two-element boolean flag array to indicate whether a thread is trying to enter the critical section. The Filter lock generalizes this notion with an n -element integer *level*[] array, where the value of *level*[*A*] indicates the highest level that thread *A* is trying to enter.

1.4.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to increment a common counter. All threads have to cooperate to increase this counter from 0 to 1024. Each of the threads will increase by one the counter 128 times. The expected result is that regardless of the order in which each thread executes the increment, at the end the counter must stay at 1024. If that is not the case, then it means that mutual exclusion did not work.

1.4.4 Observations and Interpretations

When executing the test on a dual core machine there was a a low chance of failure, but when executed on a Quad core and Six core machines the problems where more consistent.

1	Testcase: testParallel(mutex.FilterTest):	FAILED
2	expected:<1023> but was:<1024>	

1.4.5 Proposed changes to fix the problem

Even do the victim and level arrays are *volatile*, the atomicity is not guarantee and so to fix this I added the *synchronized* reserved word to the *sameOrHigher()* method in order to guarantee atomicity while checking the levels.

```
1 // Is there another thread at the same or higher level?
2 private synchronized boolean
3   sameOrHigher(int me, int myLevel) {
4   for (int id = 0; id < size; id++)
5     if (id != me && level[id] >= myLevel) {
6       return true;
7     }
8   return false;
9 }
```

Once the change is made the execution works fine on every equipment used.

```
1 Script Directory: .
2 Using OS: Linux
3 Linux Java 6 Runtime: ./jre/linux/bin
4 Will execute test from input: mutex.FilterTest
5 Classpath: ../lib/*../output/./build/
6 Running JUnit using Linux
7 JUnit version 4.12
8 .
9 Time: 0.021
10
11 OK (1 test)
```

1.4.6 Proposed solution

Once the comparing of the level was made atomic, the spin made for conflicting threads was properly working and able to cooperate to increase the counter to 1024.

One thing that is worth mentioning is that the *volatile* keyword seems to either have an implementation problem or not documented restriction, since it is not helping guarantee the use of the waiting rooms causing additional increases in the counter, thus ending with a different number rather than the one expected.

1.5 Queue Lock

1.5.1 Particular Case

In this experiment, the particular case we are trying to study is mutual exclusion where several threads try to use a queue.

We are trying to do so using bounded partial queue.

1.5.2 Solution

According to the theory the *enq()* and *deq()* methods operate on opposite ends of the queue, so as long as the queue is neither full nor empty, an *enq()* call and a *deq()* call should, in principle, be able to proceed without interference, guarantee mutual exclusion for the methods.

Here, we implement a bounded queue as a linked list of entries as shown by the Entry Class.

```
1  /**
2   * Individual queue item.
3   */
4  protected class Entry {
5      /**
6       * Actual value of queue item.
7       */
8      public T value;
9      /**
10     * next item in queue
11     */
12     public Entry next;
13     /**
14     * Constructor
15     * @param x Value of item.
16     */
17     public Entry(T x) {
18         value = x;
19         next = null;
20     }
21 }
```

1.5.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to *enq()* and *deq()* a range of numbers. All threads have to cooperate to add and remove elements from the queue. Each of the threads will enqueue and dequeue values into the queue, if everything works according to the test there will be mutual exclusion and the mapping of elements will corresponds to the queue elements. If that is not the case, a duplicate fail will be raised.

1.5.4 Observations and Interpretations

The tests executed as expected and no errors where found. Since the ReentrantLock is used to acquire an explicit lock, the mutual exclusion is guarantee by the Java `java.util.concurrent.locks` package.

```
1  /**
2   * Lock out other enqueueers (dequeueers)
3   */
4   ReentrantLock enqLock , deqLock ;
```

Chapter 2

Foundations of Shared Memory

2.1 Safe Boolean MRSW Register

2.1.1 Particular Case

In class we saw the difference between Safe, Regular and Atomic Registers. In this excersice we are experimenting with a Safe Boolean Register. In particular, we are trying a register that allows one single writer whose written value can be examined by multiple readers. We need to remember that Safe registers are those that :

- If a read does not overlap with a write, then the read returns the last written value
- If a read overlaps with a write, then the read can return any value in the domain of the register's values

2.1.2 Solution

One way to achive this behaviour is using a SRSW Safe register. The code that is provided with the book does this by having a boolean SRSW register per thread, in other words, it is an array of boolean values. When doing a write, the writer iterates over this array and writes the new value. The readers simply access its assigned slot in the array (based on the thread id) and return the stored value. The code for this type of register is shown next:

```
1  public Boolean read() {
2      return s_table[ThreadID.get()];
3  }
4
5  public void write(Boolean x) {
6      for (int i = 0; i < s_table.length; i++)
7          s_table[i] = x;
8  }
```

Please note how this type of register does not make use of any synchronization mechanism.

2.1.3 Experiment Description

This program provides two test cases. The first one is a sequential test and the second one is a parallel test. The former simply calls a write and then

a read from the same thread. This is not rocket science. The reader must retrieve what the writer wrote.

The second test case is slightly more interesting. The writer writes first one value and then another one. After that, 8 reader threads are started and they read the last written value. According to the rules of a safe register, all threads must read the same value because the reads and the writes did not overlap.

2.1.4 Sample Results

We found out that the tests failed occasionally. The manifestation of the failure is a

indexOutOfBoundsException as shown above.

```
[loraadm@gdlaa008 Register]$ junit register.SafeBooleanMRSWRegisterTest
.sequential read and write
.parallel read
Exception in thread "Thread-7" java.lang.ArrayIndexOutOfBoundsException: 8
    at register.SafeBooleanMRSWRegister.read(SafeBooleanMRSWRegister.java:26)
    at register.SafeBooleanMRSWRegisterTest$ReadThread.run(SafeBooleanMRSWRegisterTest.java:62)

Time: 0.008

OK (2 tests)
```

So, the way to fix this problem was to make sure that in each test case the threads ids are reset calling the static method *ThreadID.reset()*. With this hack, the tests passed.

```
1  /**
2   * Sequential calls.
3   */
4  public void testSequential() {
5      ThreadID.reset(); // We have to reset the thread ids
6      System.out.println("sequential_read_and_write");
7      instance.write(true);
8      boolean result = instance.read();
9      assertEquals(result, true);
10 }
11
12 /**
13  * Parallel reads
14  */
15 public void testParallel() throws Exception {
16     ThreadID.reset(); // We have to reset the thread ids
```

```

17     instance.write(false);
18     instance.write(true)
19     System.out.println("parallel_read")
20     for (int i = 0; i < THREADS; i++) {
21         thread[i] = new ReadThread(
22     }
23     for (int i = 0; i < THREADS; i++) {
24         thread[i].start();
25     }
26     for (int i = 0; i < THREADS; i++) {
27         thread[i].join();
28     }
29 }

```

And here is the resulting execution after this fix:

```

[oraadm@gdlaa008 Register]$ junit register.SafeBooleanMRSWRegisterTest
.sequential read and write
.parallel read

Time: 0.004

OK (2 tests)

```

2.1.5 Interpretation

In this experiment we saw a way of implementing Safe Multi-Reader Single-Writer registers. These registers were able of storing boolean values.

Unfortunately, none of the test cases exercises the case where we have concurrent writes and reads. Hence, this experiment only exercised one of the two rules of a safe register.

2.2 RegMRSWRegisterTest

2.2.1 Particular Case (problem)

The problem we face here is that of building a regular multi-valued multiple-reader single-writer register; out of simpler constructions, like the rest of chapter 4.

2.2.2 Solution

The solution, though supporting multi-values (integers), is bounded in the range of values. We use as primitive blocks regular boolean multiple-reader single-writer registers; and create an array of them (as many as values in supported range), and use it to represent numbers in unary way. That is, each position on the array that is set to true represents the number corresponding to that position (the index itself is the number).

Initially the boolean array is initialized to zero value, indicated by having true the cell at index 0. The write method of value x sets to true array position x , and updates to false the lower value positions (so it updates from right to left). In order to guarantee the regular property, the read method operates on the opposite direction: it goes from left to right, starting at zero position, and returns the first position whose value is true (an invariant of the array should be that there is at least one cell in with true value). The code from the book is presented below, for further reference (the range of values picked by the book's authors was that of byte Java type):

```
1 public class RegMRSWRegister implements Register<Byte> {
2     private static int RANGE = Byte.MAX_VALUE - Byte.MIN_VALUE + 1;
3     // regular boolean mult-reader single-writer register
4     boolean[] r_bit = new boolean[RANGE];
5     public RegMRSWRegister(int capacity) {
6         r_bit[0] = true; // least value
7     }
8     public void write(Byte x) {
9         r_bit[x] = true;
10        for (int i = x - 1; i >= 0; i--)
11            r_bit[i] = false;
12    }
13    public Byte read() {
```



```

14     for (int i = 0; i < RANGE; i++)
15         if (r.bit[i]) {
16             return (byte)i;
17         }
18     return -1; // never reached
19 }
20 }

```

2.2.3 Experiment Description

The test program *RegMRSWRegisterTest* includes the following individual test cases (which is pretty much the same thing as test *AtomicMRMWRegisterTest*):

- *testSequential*: calls *write* method first with a value of 11, then calls *read* method expecting read 11. A single thread is used (main thread).
- *testParallel*: calls twice method *write*, putting first 11 then 22 value. Then proceeds to create 8 reader threads, which expect all to read 22.

2.2.4 Observations and Interpretations

Just like *AtomicMRMWRegisterTest*, this does not exhibit any issue on two nor in 24 core machines. Again, part of that reason is that the test is not really interlacing writes with reads. A sample successful execution is shown below:

```

$ junit register.RegMRSWRegisterTest
.sequential read and write
.parallel read

Time: 0.004

OK (2 tests)

```

2.3 Regular Boolean MRSW Register

2.3.1 Particular Case

A register, is an object that encapsulates a value that can be observed by a *read()* method and modified by a *write()* method

For Boolean registers, the only difference between safe and regular arises when the newly written value x is the same as the old. A regular register can only return x , while a safe register may return either Boolean value.

2.3.2 Solution

A register that implements the *RegisterBoolean* interface is called a Boolean register (we sometimes use 1 and 0 as synonyms for true and false). This register uses the *ThreadLocal* Java `java.lang` package, which provides thread-local variables. These variables differ from their normal counterparts in that each thread that accesses one (via its *get* or *set* method) has its own, independently initialized copy of the variable. *ThreadLocal* instances are typically private static fields in classes that wish to associate state with a thread

```
1 public class RegBooleanMRSWRegister
2 implements Register<Boolean> {
3     ThreadLocal<Boolean> last;
4     private boolean s_value;
5     RegBooleanMRSWRegister(int capacity) {
6         this.last = new ThreadLocal<Boolean>() {
7             protected Boolean initialValue() { return false; };
8         };
9     }
10    public void write(Boolean x) {
11        if (x != last.get()) { // if new value different ...
12            last.set(x);      // remember new value
13            s_value = x;       // update register
14        }
15    }
16    public Boolean read() {
17        return s_value;
18    }
19 }
```

2.3.3 Experiment Description

The test creates 8 threads that reads the value of a register. All threads have to be able to read the current value of the register despite the previous value is contrary to the current one. The expected result is must be a 1 or true value.

2.3.4 Observations and Interpretations

The tests executed as expected and no errors where found.

2.4 Sequential Snapshot

2.4.1 Particular Case

In order to read multiple register values atomically we use an atomic snapshot. An atomic snapshot constructs an instantaneous view of an array of atomic registers. We construct a wait-free snapshot, meaning that a thread can take an instantaneous snapshot of memory without delaying any other thread.

2.4.2 Solution

Each thread repeatedly calls *collect()*, and returns as soon as it detects a clean double collect (one in which both sets of timestamps were identical). This construction always returns correct values. The *update()* calls are wait-free, but *scan()* is not because any call can be repeatedly interrupted by *update()*, and may run forever without completing. It is however obstruction-free, since a *scan()* completes if it runs by itself for long enough.

```
1 public class SeqSnapshot<T> implements Snapshot<T> {
2     private T[] a_value;
3     public SeqSnapshot(int capacity, T init) {
4         a_value = (T[])new Object[capacity];
5         for (int i = 0; i < a_value.length; i++) {
6             a_value[i] = init;
7         }
8     }
9     public synchronized void update(T v) {
10        a_value[ThreadID.get()] = v;
11    }
12    public synchronized T[] scan() {
13        T[] result = (T[]) new Object[a_value.length];
14        for (int i = 0; i < a_value.length; i++)
15            result[i] = a_value[i];
16        return result;
17    }
18 }
```

2.4.3 Experiment Description

The test creates 2 threads that write the register twice with the the values *FIRST* and *SECOND* and later parallel reads the values.

2.4.4 Observations and Interpretations

When executing the test the test fails on the parallel part due to a index out of bound exception.

```
1 Testcase: testSequential(register.SeqSnapshotTest):      Caused
  an ERROR
2 2
3 java.lang.ArrayIndexOutOfBoundsException: 2
4     at register.SeqSnapshot.update(SeqSnapshot.java:26)
5     at register.SeqSnapshotTest.testSequential(
      SeqSnapshotTest.java:40)
```

2.4.5 Proposed changes to fix the problem

By resetting the thread id at the beginning of the test the id is guarantee to be the experted while doing the parallel test.

```
1  class MyThread extends Thread {
2      public void run() {
3          ThreadID.reset();
4          instance.update(FIRST);
5          instance.update(SECOND);
6          Object[] a = instance.scan();
7          for (Object x : a) {
8              Integer i = (Integer) x;
9              int me = ThreadID.get();
10             for (int j = 0; j < THREADS; j++) {
11                 results[me][j] = (Integer)a[j];
12             }
13         }
14     }
15 }
```

Once the change is made the execution works fine on every equipment used.

```
1 Tests run: 2, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
  0.097 sec
2
```

```
3 |----- Standard Output -----|
4 | sequential
5 | parallel
6 |-----|
7 | test-single:
8 | BUILD SUCCESSFUL (total time: 0 seconds)
```

2.4.6 Proposed solution

Once the ThreadID is reset before changing from the sequential to the parallel test the execution never fails, since the out of bound happens due to the id being incremental.

2.5 Atomic MRSW Register

2.5.1 Particular Case

The characteristics of this type of register are:

- It should never happen that $R^i \rightarrow W^i$
- It should never happen that for any j , $W^i \rightarrow W^j \rightarrow R^j$
- If $R^i \rightarrow R^j$ then $i \leq j$

The book proposes to do this by using multiple SRSW registers.

2.5.2 Solution

The algorithm then requires a 2 dimensional table. When a writer decides to update the register, it has to update the values in cells $A[i][i]$, where i is the thread id. Apart from writing the value, it has to also update the cell with a timestamp.

In the other hand, when a reader wants to read from the register, it checks the timestamp of the cell $A[i][i]$. After that, it has to check other cells in the same column (ie, cells $A[x][i]$) to see if there has been an update in between. That is done by comparing the timestamps. If there is a newer timestamp, the reader has then to update all cells in its row (ie $A[i][x]$). This is the way to indicate subsequent readers which version of the value has the previous reader retrieved.

It is easy to see that the first property is satisfied because there is no way to read a future value. In order to read a value, this value has to be written before.

The second property is also satisfied because when a reader reads the value, it reads the one with the most recent timestamp by scanning the timestamps in its corresponding column. This reader also helps subsequent readers by updating all cells in its row. Forcing property 3 to be satisfied as well.

The *read()* and *write()* methods are shown next:

```
1  public synchronized T read() {
2      int me = ThreadID.get();
3      StampedValue<T> value = a_table[me][me];
4      for (int i = 0; i < a_table.length; i++) {
5          value = StampedValue.max(value, a_table[i][me]);
```

```

6      }
7      for (int i = 0; i < a_table.length; i++) {
8          a_table[me][i] = value;
9      }
10     return value.value;
11 }
12
13 public synchronized void write(T v) {
14     int me = ThreadID.get();
15     long stamp = lastStamp.get() + 1;
16     lastStamp.set(stamp); // remember for next time
17     StampedValue<T> value = new StampedValue<T>(stamp, v);
18     for (int i = 0; i < a_table.length; i++) {
19         a_table[me][i] = value;
20     }
21 }

```

2.5.3 Experiment Description

The two test cases presented demonstrates the same as previous examples so we will not describe the experiment further.

2.5.4 Sample Results

For this test, the results were always successful. Here is the sample output:

```

.sequential read and write
.parallel read

Time: 0.004

OK (2 tests)

```

2.5.5 Interpretation

In this experiment we observed a way of implementing MRSW registers. We saw that, at least in the processor we used, the algorithm seems to work correctly.

It would have been great, though, to also have test cases that demonstrate that this particular register satisfies property 3 which is the one that makes it different from the Safe version.

2.6 AtomicMRMWRegisterTest

2.6.1 Particular Case (problem)

The problem we are trying to solve here, is that of implementing an atomic multi-valued ¹ multiple-reader multiple-writer register. Let us briefly recall that atomic registers are the most powerful ones, when compared with the other two categories: safe and regular. Informally, atomic registers behave like one would expect when reads overlap writes: the reads always “see” the last written value; while regular registers allows to see either previous or latest value. Finally “safe” registers allow reads to see any value (not safe at all then!).

2.6.2 Solution

The solution from the book uses as primitive blocks the atomic multi-valued multi-reader single-writer registers; we build an array of them whose size equals the number of threads we want to support (assuming number of writers is same as readers). When a given thread *A* wants to write a value, it needs to read all the values on the array and writes to its own cell an stamped value that is bigger than any one observed. When same thread wants to read a value, it reads again whole array and returns the one with biggest stamped value (resolving ties with thread ids). This is essentially the same as Lamport’s Bakery algorithm from mutual exclusion chapter; the code is small enough to fit here, so there it goes for reference:

¹As opposed to boolean values, which have only a couple of values; while multi-valued can range over a big set, like integers.

```

1 public class AtomicMRMWRegister<T> implements Register<T>{
2     // array of multi-reader single-writer registers
3     private StampedValue<T>[] a_table;
4     public AtomicMRMWRegister(int capacity, T init) {
5         a_table = (StampedValue<T>[]) new StampedValue[capacity];
6         StampedValue<T> value = new StampedValue<T>(init);
7         for (int j = 0; j < a_table.length; j++) {
8             a_table[j] = value;
9         }
10    }
11    public void write(T value) {
12        int me = ThreadID.get();
13        StampedValue<T> max = StampedValue.MIN_VALUE;
14        for (int i = 0; i < a_table.length; i++) {
15            max = StampedValue.max(max, a_table[i]);
16        }
17        a_table[me] = new StampedValue<T>(max.stamp + 1, value);
18    }
19    public T read() {
20        StampedValue<T> max = StampedValue.MIN_VALUE;
21        for (int i = 0; i < a_table.length; i++) {
22            max = StampedValue.max(max, a_table[i]);
23        }
24        return max.value;
25    }
26 }

```

2.6.3 Experiment Description

The test program AtomicMRMWRegisterTest includes the following individual test cases:

- *testSequential*: calls *write* method first with a value of 11, then calls *read* method expecting read 11. A single thread is used (main thread).
- *testParallel*: calls twice method *write*, putting first 11 then 22 value. Then proceeds to create 8 reader threads, which expect all to read 22. This test is not really exercising multi-write capability of the register. sequentially.

2.6.4 Observations and Interpretations

The test performs well on 2 and 24 core machines, without any race conditions nor abnormal behaviors. This partly because, the test is not really exercising the multi-write capability, nor the interlacing of writes and reads. We believe the authors just copied paste a previous test (probably one that just cared about single reader and multiple writers), and forgot to update. Due time constraints we did not modify the test to exercise those missing features; a sample execution is shown below:

```
$ junit register.AtomicMRMWRegisterTest
.sequential read and write
.parallel read

Time: 0.004

OK (2 tests)
```

2.7 SimpleSnapshotTest

2.7.1 Particular Case (problem)

The problem we want to solve is that of atomic snapshots: we want to read atomically a set of registers. For this problem we assume the registers are themselves atomic, and that each one supports multiple-readers but just one writer. The problem is better defined in terms of the Java interface we want to implement:

```
public interface Snapshot<T> {  
    public void update(T v);  
    public T[] scan();  
}
```

While the method *update* allows each thread to modify its own register (single-writer), the *scan* method allows any thread to read all the registers as a single atomic operation (multiple-readers). Ideally, we would like to have a wait-free implementation of these two methods.

2.7.2 Solution

The solution presented in the book is a natural evolution to the sequential implementation (which uses *synchronized* methods for both *update* and *scan* methods). The solution is called *SimpleSnapshot* and it uses stamped values for the *update* method; there is no need to seek for the maximum, just to increase the current stamp per thread (each thread only writes to its own register). This makes the *update* method wait-free, which is the ideal (it was not very hard to achieve indeed, given the single-writer restriction).

For the *scan* method we call an auxiliary function *collect*, which represents a non atomic read of all the registers (which are copied into a new array and returned); if two consecutive *collect* calls return same values, it means that between those two calls there were no writes. When we reach that condition, we return the array of values; otherwise we repeat the iteration. Please note that *scan* method is not wait-free (we do not know how many times we will need to iterate), but at least is obstruction-free (we are not blocking other

threads from trying their own *scan* calls).

The main parts of the code are presented below for reference:

```
1 public class SimpleSnapshot<T> implements Snapshot<T> {
2     // array of atomic MRSW registers
3     private StampedValue<T>[] a_table;
4     ...
5     public void update(T value) {
6         int me = ThreadID.get();
7         StampedValue<T> oldValue = a_table[me];
8         StampedValue<T> newValue =
9             new StampedValue<T>((oldValue.stamp)+1, value);
10        a_table[me] = newValue;
11    }
12    private StampedValue<T>[] collect() {
13        StampedValue<T>[] copy = (StampedValue<T>[]) new StampedValue[a_table.length];
14        for (int j = 0; j < a_table.length; j++)
15            copy[j] = a_table[j];
16        return copy;
17    }
18    public T[] scan() {
19        StampedValue<T>[] oldCopy, newCopy;
20        oldCopy = collect();
21        collect: while (true) {
22            newCopy = collect();
23            if (! Arrays.equals(oldCopy, newCopy)) {
24                oldCopy = newCopy;
25                continue collect;
26            }
27            // clean collect
28            T[] result = (T[]) new Object[a_table.length];
29            for (int j = 0; j < a_table.length; j++)
30                result[j] = newCopy[j].value;
31            return result;
32        }
33    }
34 }
```

2.7.3 Experiment Description

The test program *SimpleSnapshotTest* consists of two individual test cases:

- *testSequential*: with a single thread we update its value first (11), and then obtain an snapshot (*scan*); expectation is to have a single value in array (the one we wrote).
- *testParallel*: we create a couple of threads, and each one writes its own register twice (first putting 11, then 22). Then each thread proceeds to call *scan* method and saves its own returned values into a global results table. At the end of test, main thread check that both threads got the same results (which should be a two element array, both with value 22).

2.7.4 Observations and Interpretations

The test runs fine, both in 2 and 24 core machines; sample output below:

```
$ junit register.SimpleSnapshotTest
.sequential
.parallel

Time: 0.002

OK (2 tests)
```

2.8 Wait-Free Snapshots

2.8.1 Particular Case

In this experiment we are dealing with the problem of creating an algorithm to create a Snapshot of a set of Register in such a way that the algorithm guarantees a Wait-Free operation.

2.8.2 Solution

The purpose of a Wait-Free snapshot is to overcome the problems of the Simple snapshot presented before. Let us remember that such snapshot executes sucesive `collect()` operations. Once it achieves a *clean double collect*, it returns the snapshot. Otherwise, it uses the following idea:

When a double collect fails, it is because a update interfered. That means that the updater could take a snapshot right before its update and other threads could use it as their snapshot too.

Since this algorithm is a bit more complex, let us explain in more detail how the implementation works.

```
1  public WFSnapshot(int capacity, T init) {
2      a_table = (StampedSnap<T>[]) new StampedSnap[capacity];
3      for (int i = 0; i < a_table.length; i++) {
4          a_table[i] = new StampedSnap<T>(init);
5      }
6  }
```

Firstly, the algorithm uses an array to store the information. The size of the array is, in this case, n , where n is the number of threads. The way in which this array is constructed is as an array of *StampedSnaps*. A *StampedSnap* is internally an array whose elements are of type T . A *StampedSnap* is a subclass of *StampedValue* which has 3 fields:

- *stamp*, which is a counter
- *owner*, which contains the thread id
- *value*, the actual value in the register

Now let us look at an auxiliary operation which is *collect()*. This operation returns a copy of the array. Observe that the method creates a new array of StampedSnap values and then iterates the existing array to copy each element.

```

1  private StampedSnap<T>[] collect() {
2      StampedSnap<T>[] copy = (StampedSnap<T>[]) new StampedSnap[a_table.length];
3      for (int j = 0; j < a_table.length; j++)
4          copy[j] = a_table[j];
5      return copy;
6  }
```

The next method to understand is *scan()*. The method first creates a backup or an old copy of the array. Then it takes a second collect. It then compares both copies. If both copies are equal, then it means that we have a double clean collect. And the result of the scan operation is any of the copies just as in the Sequential Snapshot case.

On the other hand, if the copies are different, then there are two cases. For this, we have a counter that indicates how many times, the copies have been compared and found different. If it is the first time, then we take another collect. If in the second collect, the copies are equal, then we fall in the case described in the previous paragraph. If the copies are different again, then it means that we can take the old copy as a result of the scan.

The justification of this is that if a thread A sees a thread B move twice while it is performing repeated collects, then B executed a complete *update()* call within the interval of A's scan, and so it is correct for A to use B's snapshot.

```

1  public T[] scan() {
2      StampedSnap<T>[] oldCopy;
3      StampedSnap<T>[] newCopy;
4      boolean[] moved = new boolean[a_table.length];
5      oldCopy = collect();
6      collect: while (true) {
7          newCopy = collect();
8          for (int j = 0; j < a_table.length; j++) {
9              // did any thread move?
10             if (oldCopy[j].stamp != newCopy[j].stamp) {
11                 if (moved[j]) { // second move
12                     return oldCopy[j].snap;

```



```

13         } else {
14             moved[j] = true;
15             oldCopy = newCopy;
16             continue collect;
17         }
18     }
19 }
20 // clean collect
21 T[] result = (T[]) new Object[a_table.length];
22 for (int j = 0; j < a_table.length; j++)
23     result[j] = newCopy[j].value;
24 return result;
25 }
26 }

```

Finally, we have the *update()* operation which simply updates the register with the new stamp.

```

1 public void update(T value) {
2     int me = ThreadID.get();
3     T[] snap = this.scan();
4     StampedSnap<T> oldValue = a_table[me];
5     StampedSnap<T> newValue =
6         new StampedSnap<T>(oldValue.stamp+1, value, snap);
7     a_table[me] = newValue;
8 }

```

2.8.3 Experiment Description

In this experiment, the test looks a bit different. Let us explain how. In this case we have two experiments:

1. testSequential. We spawn a new thread that updates the register with a value of *FIRST* and right after that, it scans the value of the register. The expected result is that this read retrieves the value of *FIRST*
2. testParallel. We spawn *THREADS* number of threads (in our case it is 2). Each of them will first update its register with a value of *FIRST* and then with a value of *SECOND*. Each thread then updates a position in a matrix of size *THREADS* \times *THREADS* with the result

of doing a *scan()* operation. At the end we compare consecutive rows in the matrix. The comparison is done entry by entry and we should see that for some entries the first entry is greater and for some others it is smaller than the second. What we must see is that all entries are equal and we allow the first to sometimes be either greater or smaller than the second one.

2.8.4 Sample Results

For this test, we saw that in every try, both test cases passed. Here is the output:

```
[oraadm@gdlaa008 Register]$ junit register.WFSnapshotTest
.sequential
.parallel

Time: 0.003

OK (2 tests)
```

2.8.5 Interpretation

In this experiment we were able to observe how a Wait-Free Snapshot can be constructed based on the idea of *clean double collects* and on the idea of using the updater's collect when the clean double collect method fails.

Chapter 3

Spin Locks and Contention

3.1 Problem Definition

The concept of Spin Lock comes from the following problem: *In the case where a thread tries to acquire a lock on a resource and it fails, what should be done?*

There are at least two ways to react to this problem:

1. Keep trying getting the lock. In this case the lock is called a Spin Lock. The thread is then said to be in a busy wait.
2. Suspend the execution of the current thread and let the others do some work. This action is called *blocking*

In this chapter we will study different ways in which spin locks can be implemented.

3.2 Test-And-Set Locks

3.2.1 Particular Case

The particular case we are trying to solve is that of mutual exclusion with two participants. The problem with previous approaches for solving this problem (e.g. Peterson Algorithm) is that in the real world the assumption of having a “sequential consistent memory” and a “program order guarantee among reads-writes by a given thread” are not true.

There are at least two reasons for that. The first one is that compilers, in order to improve performance of the program, can in some cases reorder certain operations. As the book clearly mentions, most programming languages preserve execution order that affect individual variables, but not across variables. This means that reads and/or writes of different variables can end up in a different order than the one appearing in the source code once the program is compiled or executed. This can clearly affect our Peterson Algorithm.

The second reason that has to do with the contradiction of sequential consistent memory is due to the hardware design which in most multiprocessor systems allow writes to be written to store buffers. These buffers are written to memory only when needed. Hence, race conditions arise, where one variable might be read before another because the latter was in a store buffer.

The generic way to solve this problem is by using instructions called *memory barriers* or *memory fences*. These instructions are able to prevent reordering operations due to write buffering. The idea is that these instructions can be called in specific parts of the code to guarantee our proposed order of execution.

The problem in this section is *how can we implement an actually working mutual exclusion algorithm with consensus number of two using memory barrier?*

3.2.2 Solution

One solution for this problem is using a synchronization instruction that includes memory barriers. For example, we can use the *testAndSet()* operation which has a consensus number of two and which atomically stores *true* in a variable and returns the previously stored value.

To design an algorithm using this instruction the idea is the following. A value of false is used to indicate that a lock is free. Then, we store a value of true to indicate that we are using the lock. Once we are done, we can set the value back to false to allow other threads take the lock.

In the java code, the *testAndSet()* is done with the *getAndSet()* method. And setting the value to false is achieved with the method *set()*. The code for this is show next:

```
1 public class TASLock implements Lock {
2     AtomicBoolean state = new AtomicBoolean( false );
3     public void lock() {
4         while ( state.getAndSet( true ) ) {} // spin
5     }
6     public void unlock() {
7         state.set( false );
8     }
9     ...
10 }
```

Observe how calling the *lock* consists simply of setting the value of the *state* variable to true by using a memory barrier operation. The same technique is used for the *unlock()* method.

3.2.3 Experiment Description

To demonstrate that this implementation works, the test that was provided does the following:

1. Initiate a shared counter with a value of 0
2. Start 8 threads
3. Each thread has to increment the counter by one 1024 times
4. At the end, the counter must hold a value of $8 * 1024$

3.2.4 Sample Results

For this test, we saw that in every try, it always passed. Here is the output:

```
[oraadm@gdlaa008 Spin]$ junit spin.TASLockTest
.parallel
```

```
Time: 0.026
```

```
OK (1 test)
```

3.2.5 Interpretation

It was shown that the algorithm seems to work and allows synchronization of two threads (consensus number was two) to access a shared resource. The way of achieving this is through the use of the memory barrier functions *getAndSet()* and *set()* of the *AtomicBoolean* class.

3.3 TTASLockTest

3.3.1 Particular Case (problem)

Here we are trying to address the shortcomings of the *TASLock* solution, which we know per the book causes a lot of bus traffic.

3.3.2 Solution

The solution is to loop on a locally cached variable, so that we do not generate traffic on the bus. The code is simple enough and is presented below:

```
1 public class TTASLock implements Lock {
2     AtomicBoolean state = new AtomicBoolean(false);
3     public void lock() {
4         while (true) {
5             while (state.get()) {}; // spin
6             if (!state.getAndSet(true))
7                 return;
8         }
9     }
10    public void unlock() {
11        state.set(false);
12    }
13    ...
14 }
```

Only the first time the *get* in the loop is called it generates bus traffic, but after that and until the moment where the state is changed by the owner of the lock, each thread uses its core cache. The drawback of this solution though, is that traffic comes once the owner releases the lock; in that moment the caches of all the cores waiting are invalidated, generating bus traffic. Furthermore, once unleashed all the threads/cores will call the synchronization primitive *getAndSet* (which again will generate traffic on the bus). After some fierce competition, once we have a winner, the system will reach again an stable configuration (looser threads spinning on local variables).

3.3.3 Experiment Description

The test is exactly the same as that of *BackoffLockTest*: 8 threads try to increment 1024 times a shared counter, and they do that concurrently. At

the end the counter shall have $8 * 1024$ as final value.

3.3.4 Observations and Interpretations

The test runs fine on 2 and 24 cores, below a sample successful execution:

```
.parallel
Time: 0.024
OK (1 test)
```

We did not require to set to volatile any of the variables, as they were already atomic (*AtomicBoolean* implies same guarantees than *volatile*, and even more).

3.4 BackoffLockTest

3.4.1 Particular Case (problem)

This is part of the chapter 7, where we are dealing with the generic issue of building real-life lock implementations that solve the mutual exclusion problem. The particular subproblem is doing that with spinning locks (those which repeatedly try to acquire the lock); and the very concrete subproblem here is how to reduce the contention out of the *TTASLock*.

3.4.2 Solution

The solution is based on a key observation, for which we need to cite the *TTASLock* code first:

```
1 public class TTASLock implements Lock {
2     AtomicBoolean state = new AtomicBoolean(false);
3     public void lock() {
4         while (true) {
5             while (state.get()) {}; // spin
6             if (!state.getAndSet(true))
7                 return;
8         }
9     }
10    public void unlock() {
11        state.set(false);
12    }
13    ...
14 }
```

If between the calls to *get* and *getAndSet*, another thread took the lock; then it must be that there is high contention for such resource (many threads trying to acquire it). Insisting on acquiring the lock on high contention scenarios is a bad idea; we will add a lot of traffic to the bus used by cores to communicate (as primitive *getAndSet* forces reads to be consistent) and the chances of getting it are scarce (too much competition).

The solution then, is to wait some time before retrying; how much? Well, the more we fail to acquire the lock the more we wait next time. The concrete

solution for that, on the context of this test, is to use what is called “adaptive exponential backoff”. Each time we fail to acquire the lock, a random number is generated within a given range and the thread sleeps that time ¹. Every time we do this waiting, the range is enlarged twice (meaning that on average the random number is twice large); this is done up to a maximum established (1024 in this case). Code is presented below for reference (note that we also have a lower bound for the minimum amount of time to wait; being 32 milliseconds here):

```

1 public class BackoffLock implements Lock {
2     private Backoff backoff;
3     private AtomicBoolean state = new AtomicBoolean(false);
4     private static final int MIN_DELAY = 32;
5     private static final int MAX_DELAY = 1024;
6     public void lock() {
7         Backoff backoff = new Backoff(MIN_DELAY, MAX_DELAY);
8         while (true) {
9             while (state.get()) {}; // spin
10            if (!state.getAndSet(true)) { // try to acquire lock
11                return;
12            } else { // backoff on failure
13                try {
14                    backoff.backoff();
15                } catch (InterruptedException ex) {
16                }
17            }
18        }
19    }
20    public void unlock() {
21        state.set(false);
22    }
23    ...
24 }
25
26 public class Backoff {
27     ...
28
29     /**
30      * Backoff for random duration.
31      * @throws java.lang.InterruptedException

```

¹This sample uses milliseconds, a common unit when dealing with multi-thread programming

```

32     */
33     public void backoff() throws InterruptedException {
34         int delay = random.nextInt(limit);
35         if (limit < maxDelay) { // double limit if less than max
36             limit = 2 * limit;
37         }
38         Thread.sleep(delay);
39     }
40 }

```

By the way, the code for *Backoff* has an error (not shown above to save space in this report); the argument *max* passed to the constructor is not really used (there seems to be a typo there, where we used *min* parameter instead).

3.4.3 Experiment Description

The test consists in a single case *testParallel*, which creates 8 threads and puts each one to increase 1024 times a shared counter. The access to the counter is protected by the *BackoffLock* solution we described already. The assertion of the test is that the shared counter ends up with a final value, that reflects the fact that all the thread contributions are present (nobody lost increments, and we do not have more than expected).

3.4.4 Observations and Interpretations

With or without the fix mentioned about the *max* argument constructor in *Backoff* class, the test works fine on 2 and 24 cores. Below a sample execution:

```

.parallel
Time: 0.075
OK (1 test)

```

3.5 CLH Queue Lock

3.5.1 Particular Case

Using queue locks solve two problems that appear in simpler locking protocols. Firstly, by using a queue, cache-coherence traffic is reduced because each thread spin on a different memory location; and secondly, by using a queue, threads can know whether their turn has come by inspecting the status of their predecessor in the queue.

The particular case that we want to study in this experiment is *How can we implemente a space efficient Queue Lock?*

3.5.2 Solution

One solution for this problem is given by the CLHLock protocol. The algorithm that implement this protocol need two fields local to the thread: A pointer to the current node and a pointer to the previous node. The queue is constructed by having the pointer *previous* pointing to the previous node's *current* field. This field contains a boolean value which is false when the thread releases the lock. When this is the case, the next thread is allowed to acquire the lock.

3.5.3 Experiment Description

To demonstrate that this implementation works, the test that was provided does the following:

1. Initiate a shared counter with a value of 0
2. Start 8 threads
3. Each thread has to increment the counter by one 1024 times
4. At the end, the counter must hold a value of $8 * 1024$

These are the details of the system we used to run the experiments:

- Processor: Intel Core i5 @2.5 GHz. 2 Cores.
- L2 Cache per Core: 256 KB

- L3 Cache: 3 MB
- System Memory: 16 GB

3.5.4 Sample Results

For this test, we saw that in every try, it always passed.

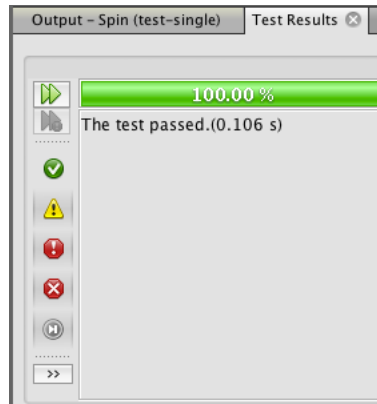


Figure 3.1: Successful execution of the tests for CLH Queue Lock test

```
compile-test-single:
  Testsuite: spin.CLHLockTest
  Tests run: 1, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 0.106 sec

test-single:
  BUILD SUCCESSFUL (total time: 0 seconds)
```

Figure 3.2: Successful execution of the tests for CLH Queue Lock test

3.5.5 Interpretation

It was shown that the algorithm seems to work and allows synchronization of Multiple threads trying to access a shared variable.

One important thing about this algorithm is that threads spin checking for different memory addresses which avoids invalidation of cached copies. Another important advantage is that it requires a limited number of memory and also provides first-in-first-out fairness.

The case where this algorithm is not good is when, in a NUMA architecture, the *previous* pointer points to a remote location. In that case, the performance of the algorithm degrades.

3.6 MCSLockTest

3.6.1 Particular Case (problem)

The problem is still how to build scalable spinning locks, and in particular, how to overcome the limitations from the backoff approaches. The queue-based approach introduced both *ALock* and *CLHLock*, the first is space inefficient and the second overcomes such limitation at the expense of not being suitable for cache-less NUMA architecture.

3.6.2 Solution

The solution or rather alternative to *CLHLock* approach is called *MCSLock*; and that is the algorithm tested in this section. Just like the *CLHLock* approach, the lock is modeled with a linked list of *QNode* objects, where each one represents either a lock holder or a thread waiting to acquire the lock; the difference though, is that the list is explicit, not virtual². On the *MCSLock* approach, instead of embedding the list in thread-local variables, it is placed in the globally accessible *QNode* objects.

The *MCSLock* code is presented below for further reference:

```
1 public class MCSLock implements Lock {
2     AtomicReference<QNode> queue;
3     ThreadLocal<QNode> myNode;
4     public MCSLock() {
5         queue = new AtomicReference<QNode>(null);
6         // initialize thread-local variable
7         myNode = new ThreadLocal<QNode>() {
8             protected QNode initialValue() {
9                 return new QNode();
10            }
11        };
12    }
13    public void lock() {
14        QNode qnode = myNode.get();
15        QNode pred = queue.getAndSet(qnode);
16        if (pred != null) {
17            qnode.locked = true;
```

²The author introduces the term “virtual”, when describing *CLHLock* because the list is implicit: each thread refers to its predecessor through a thread-local *pred* variable.

```

18         pred.next = qnode;
19         while (qnode.locked) {} // spin
20     }
21 }
22 public void unlock() {
23     QNode qnode = myNode.get();
24     if (qnode.next == null) {
25         if (queue.compareAndSet(qnode, null))
26             return;
27         while (qnode.next == null) {} // spin
28     }
29     qnode.next.locked = false;
30     qnode.next = null;
31 }
32 ...
33 static class QNode { // Queue node inner class
34     boolean locked = false;
35     QNode next = null;
36 }
37 }

```

This *MCSLock* solution is not a perfect replacement for *CLHLock*: on one side it has same advantages than *CLHLock*, in particular, the property that each lock release invalidates only the successor's cache entry, with the plus that it is better suited to cache-less NUMA architectures because each thread controls the location on which it spins. The space complexity is the same as *CLHLock*, $O(L + n)$, as the nodes can be recycled.

On the other side, the disadvantages of *MCSLock* algorithm is that releasing a lock requires spinning; and that it requires more reads, writes, and `compareAndSet()` calls than the *CLHLock* algorithm.

3.6.3 Experiment Description

The test is exactly the same as that of *BackoffLockTest*: 8 threads try to increment 1024 times a shared counter, and they do that concurrently. At the end the counter shall have $8 * 1024$ as final value.

3.6.4 Observations and Interpretations

The test works fine in two-cores, but in 24 cores it hangs time to time; below a sample thread dump captured when hanging occurred (after trying the test a bit more than 300 times):

Full thread dump OpenJDK 64-Bit Server VM (23.7-b01 mixed mode):

```
"Thread-7" prio=10 tid=0x00007f3c4c0e0000 nid=0x6aaa runnable [0x00007f3bf6e4c000]
  java.lang.Thread.State: RUNNABLE
    at spin.MCSLock.lock(MCSLock.java:41)
    at spin.MCSLockTest$MyThread.run(MCSLockTest.java:52)

"Thread-6" prio=10 tid=0x00007f3c4c0de000 nid=0x6aa9 runnable [0x00007f3bf6f4d000]
  java.lang.Thread.State: RUNNABLE
    at spin.MCSLock.lock(MCSLock.java:41)
    at spin.MCSLockTest$MyThread.run(MCSLockTest.java:52)

"Thread-5" prio=10 tid=0x00007f3c4c0dc000 nid=0x6aa8 runnable [0x00007f3bf704e000]
  java.lang.Thread.State: RUNNABLE
    at spin.MCSLock.lock(MCSLock.java:41)
    at spin.MCSLockTest$MyThread.run(MCSLockTest.java:52)

"Thread-4" prio=10 tid=0x00007f3c4c0da000 nid=0x6aa7 runnable [0x00007f3bf714f000]
  java.lang.Thread.State: RUNNABLE
    at spin.MCSLock.lock(MCSLock.java:41)
    at spin.MCSLockTest$MyThread.run(MCSLockTest.java:52)

"Thread-3" prio=10 tid=0x00007f3c4c0d7000 nid=0x6aa6 runnable [0x00007f3bf7250000]
  java.lang.Thread.State: RUNNABLE
    at spin.MCSLock.lock(MCSLock.java:41)
    at spin.MCSLockTest$MyThread.run(MCSLockTest.java:52)

"Thread-2" prio=10 tid=0x00007f3c4c0d5800 nid=0x6aa5 runnable [0x00007f3bf7351000]
  java.lang.Thread.State: RUNNABLE
    at spin.MCSLock.lock(MCSLock.java:41)
    at spin.MCSLockTest$MyThread.run(MCSLockTest.java:52)

"Thread-1" prio=10 tid=0x00007f3c4c0d4000 nid=0x6aa4 runnable [0x00007f3bf7452000]
  java.lang.Thread.State: RUNNABLE
    at spin.MCSLock.lock(MCSLock.java:41)
    at spin.MCSLockTest$MyThread.run(MCSLockTest.java:52)

"Thread-0" prio=10 tid=0x00007f3c4c0d8800 nid=0x6aa3 runnable [0x00007f3bf7553000]
  java.lang.Thread.State: RUNNABLE
    at spin.MCSLock.lock(MCSLock.java:41)
    at spin.MCSLockTest$MyThread.run(MCSLockTest.java:52)
...
"main" prio=10 tid=0x0000000000c8a800 nid=0x6a79 in Object.wait() [0x00007f3c5ce0d000]
  java.lang.Thread.State: WAITING (on object monitor)
    at java.lang.Object.wait(Native Method)
      - waiting on <0x0000000058386adc8> (a spin.MCSLockTest$MyThread)
    at java.lang.Thread.join(Thread.java:1260)
      - locked <0x0000000058386adc8> (a spin.MCSLockTest$MyThread)
    at java.lang.Thread.join(Thread.java:1334)
    at spin.MCSLockTest.testParallel(MCSLockTest.java:43)
```

```

at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:606)
at junit.framework.TestCase.runTest(TestCase.java:176)
at junit.framework.TestCase.runBare(TestCase.java:141)
at junit.framework.TestResult$1.protect(TestResult.java:122)
at junit.framework.TestResult.runProtected(TestResult.java:142)
at junit.framework.TestResult.run(TestResult.java:125)
at junit.framework.TestCase.run(TestCase.java:129)
at junit.framework.TestSuite.runTest(TestSuite.java:252)
at junit.framework.TestSuite.run(TestSuite.java:247)
at junit.textui.TestRunner.doRun(TestRunner.java:116)
at junit.textui.TestRunner.start(TestRunner.java:183)
at junit.textui.TestRunner.main(TestRunner.java:137)

```

From the above stacks, we can observe that the main thread hangs because none of the launched 8 threads gets out of the spinning while of *lock* method (line 41):

```

1      while (qnode.locked) {}          // spin

```

The reason of the above must be that the *locked* flag is not declared as *volatile*; without such keyword, the JVM does not guarantee any synchronization time limits between the *RAM* and the local caches of each core. Thus, the hanging shall occur when the write made by last thread to its local flag (*locked=false*, did not get propagated to any other core for a while; we do now know when they would get propagated, actually.

This is an interesting aspect of the *AtomicReference* class; while it does cover the “reference” to a *QNode*, it does not appear to affect the attributes of the object. Thus, changes to the *queue* reference will be atomic indeed; but they would point to attributes whose writes are not.

One may be tempted to think that the *volatile* keyword is only applicable to the *locked* boolean flag; but even with that chance the test may still hang. Below sample thread dump on the 24 cores machine, after retrying test 1118 times:

Full thread dump OpenJDK 64-Bit Server VM (23.7-b01 mixed mode):

```

"Thread-7" prio=10 tid=0x00007f55b00ef800 nid=0x9e07 runnable [0x00007f555c6e8000]
  java.lang.Thread.State: RUNNABLE
    at spin.MCSLock.unlock(MCSLock.java:49)
    at spin.MCSLockTest$MyThread.run(MCSLockTest.java:56)

"Thread-6" prio=10 tid=0x00007f55b00ed800 nid=0x9e06 runnable [0x00007f555c7e9000]
  java.lang.Thread.State: RUNNABLE
    at spin.MCSLock.lock(MCSLock.java:41)
    at spin.MCSLockTest$MyThread.run(MCSLockTest.java:52)

...

"main" prio=10 tid=0x00000000018eb800 nid=0x9dd6 in Object.wait() [0x00007f55c2736000]
  java.lang.Thread.State: WAITING (on object monitor)
    at java.lang.Object.wait(Native Method)
    - waiting on <0x0000000058386b910> (a spin.MCSLockTest$MyThread)
    at java.lang.Thread.join(Thread.java:1260)
    - locked <0x0000000058386b910> (a spin.MCSLockTest$MyThread)
    at java.lang.Thread.join(Thread.java:1334)
    at spin.MCSLockTest.testParallel(MCSLockTest.java:43)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:606)
    at junit.framework.TestCase.runTest(TestCase.java:176)
    at junit.framework.TestCase.runBare(TestCase.java:141)
    at junit.framework.TestResult$1.protect(TestResult.java:122)
    at junit.framework.TestResult.runProtected(TestResult.java:142)
    at junit.framework.TestResult.run(TestResult.java:125)
    at junit.framework.TestCase.run(TestCase.java:129)
    at junit.framework.TestSuite.runTest(TestSuite.java:252)
    at junit.framework.TestSuite.run(TestSuite.java:247)
    at junit.textui.TestRunner.doRun(TestRunner.java:116)
    at junit.textui.TestRunner.start(TestRunner.java:183)
    at junit.textui.TestRunner.main(TestRunner.java:137)

```

In the above dump there are only hanging threads two; “Thread-6” is waiting on the already seen spinning loop of *lock* method, which should no longer be attributable to the lack of *volatile* attribute (we just fixed that). This spinning loop of the “Thread-6” hanging thread then, must be valid: the *locked* flag of the other thread has not changed.

The new part is the other hanging thread “Thread-7”, which is waiting on the spinning loop of the *unlock* method:

```

1      while (qnode.next == null) {} // spin

```

We immediately notice same problem as with *locked* flag; the *next* pointer also needs immediate propagation to main memory, otherwise the threads

will not detect then is being set to null. Let us also note that the hanging due *next* variable, triggers the hanging on the spinning loop for *lock* method; the *locked* flag is volatile, so a write to it would propagate right away and invalidate other's caches; but since the thread doing the *unlock* never exits the loop, it can not perform such write. Putting as volatile both attributes of class *Qnode* solves this issue; we ran the test 10,000 times without a hang.

An interesting observation made while discussing this exercise on the distribution list, was that the solution imposes a serialization, after all; but that is correct. Let us remember that the queue is just the internal implementation artifact of the high level lock construction. And for this particular test, we are trying to solve the mutual-exclusion problem for a counter; thus serialization is what we want, indeed.

Even if serialization is required, there are several ways of achieving it; some are more efficient at hardware level than others. The family of solutions that *MCSLock* belongs to, tries to minimize traffic on the core buses (as we only invalidate the cache of the core that owes the modified variables *locked* / *next*).

3.7 Array Lock

3.7.1 Particular Case

The ALock, is a simple array-based queue lock. In a queue, each thread can learn if its turn has arrived by checking whether its predecessor has finished.

3.7.2 Solution

The threads share an *AtomicInteger* tail field, initially zero. To acquire the lock, each thread atomically increments tail. Call the resulting value the threads slot. The slot is used as an index into a Boolean flag array. If *flag[j]* is true, then the thread with slot *j* has permission to acquire the lock.

```
1 public class ALock implements Lock {
2     // thread-local variable
3     ThreadLocal<Integer> mySlotIndex = new ThreadLocal<Integer> ()
4     {
5         protected Integer initialValue() {
6             return 0;
7         }
8     };
9     volatile AtomicInteger tail;
10    volatile boolean[] flag;
11    int size;
12    /**
13     * Constructor
14     * @param capacity max number of array slots
15     */
16    public ALock(int capacity) {
17        size = capacity;
18        tail = new AtomicInteger(0);
19        flag = new boolean[capacity];
20        flag[0] = true;
21    }
22    public void lock() {
23        int slot = tail.getAndIncrement() % size;
24        mySlotIndex.set(slot);
25        while (! flag[mySlotIndex.get()]) {} // spin
26    }
27    public void unlock() {
28        flag[mySlotIndex.get()] = false;
29        flag[(mySlotIndex.get() + 1) % size] = true;
30    }
31 }
```

```

30 // any class implementing Lock must provide these methods
31 public Condition newCondition() {
32     throw new java.lang.UnsupportedOperationException();
33 }
34 public boolean tryLock(long time,
35     TimeUnit unit)
36     throws InterruptedException {
37     throw new java.lang.UnsupportedOperationException();
38 }
39 public boolean tryLock() {
40     throw new java.lang.UnsupportedOperationException();
41 }
42 public void lockInterruptibly() throws InterruptedException {
43     throw new java.lang.UnsupportedOperationException();
44 }
45 }

```

3.7.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to increment a common counter. All threads have to cooperate to increase this counter from 0 to 1024. Each of the threads will increase by one the counter 128 times. The expected result is that regardless of the order in which each thread executes the increment, at the end the counter must stay at 1024. If that is not the case, then it means that mutual exclusion did not work.

3.7.4 Observations and Interpretations

When executing the test values where lost from the queue.

```

1 1009
2 1010
3 1011
4 1013
5 1014
6 1015

```

3.7.5 Proposed changes to fix the problem

By adding the volatile keyword in the ALock class the execution was fixed.

```
1  volatile AtomicInteger tail;  
2  volatile boolean[] flag;  
3  int size;
```

Once the change is made the execution works fine on every equipment used.

3.7.6 Proposed solution

Apparently the `AtomicInteger` keyword is not guarantee that the values are propagated through the memory hierarchy and by making the value atomic the execution works.

3.8 Composite Fast Path Lock

3.8.1 Particular Case

In this experiment, we extend the *CompositeLock* algorithm to encompass a fast path in which a solitary thread acquires an idle lock without acquiring a node and splicing it into the queue.

3.8.2 Solution

We use a *FASTPATH* flag to indicate that a thread has acquired the lock through the fast path. Because we need to manipulate this flag together with the tail fields reference, we take a high-order bit from the tail fields integer stamp. The private *fastPathLock()* method checks whether the tail fields stamp has a clear *FASTPATH* flag and a null reference. If so, it tries to acquire the lock simply by applying *compareAndSet()* to set the FAST-PATH flag to true, ensuring that the reference remains null. An uncontended lock acquisition thus requires a single atomic operation. The *fastPathLock()* method returns true if it succeeds, and false otherwise.

3.8.3 Experiment Description

The test creates 2 threads that need to be coordinate in order to increase a counter. All threads have to cooperate in order to increase the counter into the desire value, this lock implements a explicit lock that the test calls. If that is not the case, an assertion fail will be raised.

3.8.4 Observations and Interpretations

The tests executed as expected and no errors where found. Since the *CompositeFastPathLock* is used to acquire an explicit lock, the mutual exclusion is guarantee.

```
1 Tests run: 1, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   0.895 sec
2
3 test-single:
4 BUILD SUCCESSFUL (total time: 3 seconds)
```


3.9 Time Out Lock

3.9.1 Particular Case

In this experiment we consider the Java Lock interface which includes a *tryLock()* method that allows the caller to specify a *timeout* or maximum duration the caller is willing to wait to acquire the lock. The TOLock implements this mechanism in order to wait for a resource.

3.9.2 Solution

The *tryLock()* method creates a new QNode with a null pred field and appends it to the list as in the CLHLock class. If the lock was free, the thread enters the critical section. Otherwise, it spins waiting for its predecessors QNodes pred field to change. If the predecessor thread times out, it sets the pred field to its own predecessor, and the thread spins instead on the new predecessor.

```
1      while (System.nanoTime() - startTime < patience) {
2          QNode predPred = pred.pred;
3          if (predPred == AVAILABLE) {
4              return true;
5          } else if (predPred != null) { // skip predecessors
6              pred = predPred;
7          }
8      }
9      // timed out; reclaim or abandon own node
10     if (!tail.compareAndSet(qnode, pred))
11         qnode.pred = pred;
```

3.9.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to increase a counter. All threads have to cooperate in order to increase the counter into the desire value, this lock implements a explicit lock that the test calls, but with the addition of a time out. If that is not the case, an assertion fail will be raised.

3.9.4 Proposed changes to improve execution

By making volatile the QNode's in the TOLock class the execution finishes faster.

```
1  volatile static QNode AVAILABLE = new QNode();
2  volatile AtomicReference<QNode> tail;
3  ThreadLocal<QNode> myNode;
```

3.9.5 Observations and Interpretations

The tests executed as expected and no errors where found. But the time out generates some waiting depending on the number of cores the machine has.

```
1  Tests run: 2, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   22.551 sec
2
3  _____ Standard Output _____
4  locking
5  Created Threads
6  Started Threads
7  Finished Threads
8  timeout
9  timeouts: 8191
10 _____
11 test-single:
12 BUILD SUCCESSFUL (total time: 24 seconds)
```

3.10 CompositeLockTest

3.10.1 Particular Case (problem)

The particular problem we are trying to solve, is that of looking for a good balance between the trade-offs imposed by queue locks vs backoff locks. On one hand, queue locks provide FIFO fairness, fast lock release, and low contention, but require nontrivial protocols for recycling abandoned nodes; on the other hand, backoff locks support trivial timeout protocols, but are not scalable, and may have slow lock release if timeout parameters are not well-tuned.

3.10.2 Solution

The solution proposed is to combine both type of locks, and is based on the observation that only the threads at the front of the queue need to do lock handoffs; while the rest may be happy doing backoffs. This solution occupies several pages at the book, and putting all the code and the details here may be cumbersome; let us just put the main high-level details.

The *CompositeLock* solution uses a short array of auxiliary lock nodes (the array is static, it can not be resized). Once a thread needs to acquire “the” lock it picks a random slot in the array; which would have a node. If the auxiliary lock protecting that node is used, the thread uses an adaptive backoff approach and retries (after sleeping the required time). Once the thread could acquire the lock over the picked node, it places that node into a *TOLock*-style queue. Then, the thread spins on the preceding node, and when that node’s owner indicates it has finished, the thread can finally acquire “the” lock (and enter its critical section). When the thread leaves due completion or time-out, it releases ownership of the node, and another thread doing backoff can acquire it. There are far more details regarding the procedure to recycle the freed nodes of the array, while multiple threads attempt to acquire control over them; we omit those details here.

3.10.3 Experiment Description

There is a single test case *testParallel*, which creates a couple of threads and asks each to increment 2048 times a shared counter. The counter increment

is the critical section, and is protected with a *CompositeLock*. At the end the shared counter shall have a value of 4096.

3.10.4 Observations and Interpretations

The test runs fine on our two and 24 core machines, sample execution below:

```
.  
Time: 0.066  
OK (1 test)
```

3.11 Hierarchical Backoff Lock

3.11.1 Particular Case

In the following exercises we will be dealing with the idea of Hierarchical locks. The motivation of this type of locks are the architecture of modern cache-coherent computers where processors are organized as clusters. Communication between processors is then categorized in intra-cluster communication and inter-cluster communication. The former is faster and cheaper than the latter.

Given the following background, the idea is to build locking algorithms that work efficiently in these architectures. The problem is then: *how can we create locking algorithm for this type of systems?*

3.11.2 Solution

The first algorithm that we are going to study is the Hierarchical Backoff Lock. This algorithm uses the *test-and-test-and-set* lock that we presented before. If a thread from a cluster releases a lock, then it is more likely that another thread from the same cluster acquires it. This strategy reduces the overall time needed to switch the lock ownership. Let us take a quick look at the source code in java:

```
1  public void lock() {
2      int myCluster = ThreadID.getCluster();
3      Backoff localBackoff =
4          new Backoff(LOCAL_MIN_DELAY, LOCAL_MAX_DELAY);
5      Backoff remoteBackoff =
6          new Backoff(REMOTE_MIN_DELAY, REMOTE_MAX_DELAY);
7      while (true) {
8          if (state.compareAndSet(FREE, myCluster)) {
9              return;
10         }
11         int lockState = state.get();
12         try {
13             if (lockState == myCluster) {
14                 localBackoff.backoff();
15             } else {
16                 remoteBackoff.backoff();
17             }
18         } catch (InterruptedException ex) {
```

```

19     }
20   }
21 }
22
23 public void unlock() {
24     state.set(FREE);
25 }

```

For locking, we create two backoff objects which are the ones that have a method to perform the backoff for a random period. Notice that we distinguish between two cases: a local backoff and a remote backoff. The delays are as follows:

```

1  private static final int LOCAL_MIN_DELAY = 8;
2  private static final int LOCAL_MAX_DELAY = 256;
3  private static final int REMOTE_MIN_DELAY = 256;
4  private static final int REMOTE_MAX_DELAY = 1024;

```

A remote delay is longer than a local delay. Then we perform a *compare-AndSet()* operation. If the state of the lock is free, then we set its value to the value of our cluster. That indicates that we are holding the lock now.

If the state was not free, then we perform a backoff. If the lock is being hold by another thread in the cluster, then we do a short backoff. Otherwise, we do a long backoff.

The unlock method simply sets the state to free.

3.11.3 Experiment Description

To demonstrate that this implementation works, the test that was provided does the following:

1. Initiate a shared counter with a value of 0
2. Start 8 threads
3. Each thread has to increment the counter by one 1024 times
4. At the end, the counter must hold a value of $8 * 1024$

3.11.4 Sample Results

For this test, we saw that in every try, it always passed. Here is the output:

```
[oraadm@gdlaa008 Spin]$ junit spin.HBOLockTest
.  
Time: 0.262  
OK (1 test)
```

3.11.5 Interpretation

In this experiment we discovered how to implement a locking algorithm in a cache-coherent architecture by performing backoffs of different delays depending on whether the thread holding the lock is in the same or in a different cluster of CPUs.

One problem that we can foresee though, is that this algorithm could cause starvation of remote threads because the algorithms give more priority to local threads. This is not a desirable property. In the next algorithm we shall see a way of overcoming this problem

Chapter 4

Monitors and Blocking Synchronization

4.1 Queue Lock

4.1.1 Particular Case

In this experiment, the particular case we are trying to study is mutual exclusion where several threads try to use a queue.

We are trying to do so using bounded partial queue.

4.1.2 Solution

According to the theory the *enq()* and *deq()* methods operate on opposite ends of the queue, so as long as the queue is neither full nor empty, an *enq()* call and a *deq()* call should, in principle, be able to proceed without interference, guarantee mutual exclusion for the methods.

Here, we implement a bounded queue as a linked list of entries as shown by the Entry Class.

```
1  /**
2   * Individual queue item.
3   */
4  protected class Entry {
5      /**
6       * Actual value of queue item.
7       */
8      public T value;
9      /**
10     * next item in queue
11     */
12     public Entry next;
13     /**
14     * Constructor
15     * @param x Value of item.
16     */
17     public Entry(T x) {
18         value = x;
19         next = null;
20     }
21 }
```

4.1.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to *enq()* and *deq()* a range of numbers. All threads have to cooperate to add and remove elements from the queue. Each of the threads will enqueue and dequeue values into the queue, if everything works according to the test there will be mutual exclusion and the mapping of elements will corresponds to the queue elements. If that is not the case, a duplicate fail will be raised.

4.1.4 Observations and Interpretations

The tests executed as expected and no errors where found. Since the ReentrantLock is used to acquire an explicit lock, the mutual exclusion is guarantee by the Java `java.util.concurrent.locks` package.

```
1  /**
2   * Lock out other enqueueers (dequeueers)
3   */
4   ReentrantLock enqLock , deqLock ;
```

4.2 Non-Reentrant Mutual Exclusion

4.2.1 Particular Case

In this experiment we use a re-entrant mutual exclusion Lock with the same basic behavior and semantics as the implicit monitor lock accessed using synchronized methods and statements, but with extended capabilities.

4.2.2 Solution

The Mutex lock extends a *AbstractQueuedSynchronizer*, we have a *tryAcquire()* that only acquire the lock when the state is FREE, an *tryRelease()* that releases the lock by setting the state to FREE and a method to verify the current state of the lock. Due to the AbstractQueued we use a framework for implementing blocking locks and related synchronizers (semaphores, events, etc) that rely on first-in-first-out (FIFO) wait queues. This extended class is designed to be a useful basis for most kinds of synchronizers that rely on a single atomic int value to represent state, thus the object does all the work and we forward the conditions.

```
1  private final LockSynch sync = new LockSynch();
2
3  public void lock() {
4      sync.acquire(0);
5  }
6  public boolean tryLock() {
7      return sync.tryAcquire(0);
8  }
9  public void unlock() {
10     sync.release(0);
11 }
12 public Condition newCondition() {
13     return sync.newCondition();
14 }
15 public boolean isLocked() {
16     return sync.isHeldExclusively();
17 }
18 public boolean hasQueuedThreads() {
19     return sync.hasQueuedThreads();
20 }
21 public void lockInterruptibly() throws InterruptedException {
22     sync.acquireInterruptibly(0);
23 }
```

```

24 | public boolean tryLock(long timeout, TimeUnit unit) throws
    |     InterruptedException {
25 |     return sync.tryAcquireNanos(1, unit.toNanos(timeout));
26 | }

```

4.2.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to increase a counter. All threads have to cooperate in order to increase the counter into the desire value, this lock implements a explicit lock that the test calls, but with the addition of a re-entrant lock sync, which in addition to implementing the Lock interface, this class defines methods *isLocked()* and *hasQueueThreads()*, as well as some associated protected access methods that may be used for instrumentation and monitoring. If that is not the case, an assertion fail will be raised.

4.2.4 Observations and Interpretations

The tests executed as expected and no errors where found. Since we extend the AbstractQueuedSynchronizer, we are mainly using the Java `java.util.concurrent.locks` package.

```

1 | Tests run: 1, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
  |     0.382 sec
2 |
3 | _____ Standard Output _____
4 | parallel
5 | _____
6 | test-single:
7 | BUILD SUCCESSFUL (total time: 1 second)

```

4.3 Count Down Latch

4.3.1 Particular Case

The purpose of this experiment to our eyes is to demonstrate how monitor locks and conditions are used. Let us remember both concepts before going further.

The concept of monitor consists on encapsulating three components: data, methods and synchronization mechanism. This allows to have our data structures (or classes) take care of synchronization instead of overwhelming the user of the API with this task.

Now, the concept of *conditions* comes into play because with them we can signal threads about an event. For example, instead of having threads spinning waiting for a value to become false, we can instead signal threads and let them know that they can now re-check for their locking condition.

4.3.2 Solution

The java interface for *Lock* suggests two methods to implement the protocol drafted above. First, there is a method *await()* which basically asks the thread wait till it is signaled about an event. The accompanying method *signalAll* achieves this latter task.

Let us explain how the implementation of the methods *countDown()* and *await()* work:

```
1      public void await() throws InterruptedException {
2          lock.lock();
3          try {
4              while (counter > 0)
5                  condition.await();
6          } finally {
7              lock.unlock();
8          }
9      }
10
11     public void countDown() {
12         lock.lock();
13         try {
14             counter--;
15             if (counter == 0) {
```

```

16         condition.signalAll();
17     }
18 } finally {
19     lock.unlock();
20 }
21 }

```

Basically, when we create an object of type *CountDownLatch*, we initialize a counter with a value. The goal of the object is to count down to zero. To achieve this, we need to execute the method *countDown()* a number of times.

For example, one thread can create a *CountDownLatch* object with a counter of 10 and then it sits to wait, ie. calls the *await()* function. Other threads then call the method *countDown* to decrement the counter. When the counter reaches zero, then the first thread gets signaled.

4.3.3 Experiment Description

This experiment is a bit different from others showed before. The idea in this experiment is to show how the methods *await()* and *signalAll()* are combined in a working program.

So, the idea of the test is the following. Initially we have 8 threads. We start them as usual but we do not allow them to proceed normally. Instead, we will ask them to wait for a signal. Internally, what we do is decrement a shared variable. When this variable becomes 0, the signal is triggered. Initially this signal is 1, meaning that we only need to decrement the variable once.

After that, each thread will start and wait for the next signal. The variable that controls this other signal is initially set to 8. Each thread will be in charge of decrementing this variable by one. Once it becomes 0, the second signal will be triggered announcing that our test must finish.

Let us quickly review the test case:

At the beginning we see the two *CountDownLatch* objects. One is used to start and one is used to finish. The one for start is initialized with 1 because only one thread will trigger it. The one for start is initialized with 8 because we will create 8 threads and each of them will call *countDown()* once.

```

1    ...
2    static final int THREADS = 8; // number threads to test

```

```

3  static final int START = 1492;
4  static final int STOP  = 1776;
5  CountdownLatch startSignal = new CountdownLatch(1);
6  CountdownLatch doneSignal = new CountdownLatch(THREADS);
7  ...
8  public void testCountDownLatch() {
9      System.out.format(" Testing_%d_threads\n", THREADS);
10     CountdownLatch startSignal = new CountdownLatch(1);
11     CountdownLatch doneSignal  = new CountdownLatch(THREADS);
12     Thread[] thread = new Thread[THREADS];
13     int[] log = new int[THREADS];
14     // create threads
15     for (int j = 0; j < THREADS; j++)
16         thread[j] = new TestThread(j, THREADS, startSignal, doneSignal, log);
17     // start threads
18     for (int j = 0; j < THREADS; j++)
19         thread[j].start();
20     // do something
21     for (int i = 0; i < THREADS; i++)
22         log[i] = START;
23     // give threads permission to start
24     startSignal.countDown();
25     // wait for threads to complete
26     try {
27         doneSignal.await();
28     } catch (InterruptedException e) {
29         fail("interrupted_exception");
30     }
31     for (int i = 0; i < THREADS; i++)
32         if (log[i] != STOP)
33             fail("found_bad_value:" + log[i]);
34 }

```

So, at start the main thread creates 8 threads and make them run. When these threads start, the are put to wait immediately. Then the main thread signals them to actually start and it puts himself to wait.

```

1  public class TestThread extends Thread {
2      int threads;
3      CountdownLatch startSignal;
4      CountdownLatch doneSignal;
5      int[] log;
6      int index;
7      public TestThread(int index, int threads,

```

```

8         CountdownLatch startSignal, CountdownLatch doneSignal, int[] log) {
9         this.threads = threads;
10        this.log = log;
11        this.index = index;
12        this.startSignal = startSignal;
13        this.doneSignal = doneSignal;
14    }
15    public void run() {
16        ThreadID.set(index);
17        try {
18            startSignal.await(); // wait for permission to start
19        } catch (InterruptedException ex) {
20            ex.printStackTrace();
21        }
22        if (log[index] != START) {
23            System.out.format("%d\tError expected %d found %d at %d\n", START,
24 log[index], index);
25        }
26        log[index] = STOP;
27        doneSignal.countDown(); // notify driver we are done
28    }
29 }

```

Then each of the 8 threads will run and do their stuff. At the end, each of the threads calls *countDown()* for the doneSignal. When the counter for this latch reaches 0, the program ends.

4.3.4 Sample Results

For this test, we saw that in every try, it always passed. Here is the output:

```

[oraadm@gdlaa008 Monitor]$ junit monitor.CountDownLatchTest
.Testing 8 threads

Time: 0.01

OK (1 test)

```

4.3.5 Interpretation

This experiment showed us the way in which Monitors are used. We saw that we did not need to continuously poll for a flag to acquire a lock. Instead, the monitor sent the threads a signal to indicate them that a particular condition

has been met. In our particular test, these signals indicated: (1) that the threads can start; and (2) that the threads must stop.

Chapter 5

Linked Lists: The Role of Locking

5.1 Coarse List

5.1.1 Particular Case

In this exercise we are trying to solve the problem of accessing a List data structure by multiple threads simultaneously. While allowing multiple threads accessing the data structure, we still want to maintain its consistency.

As mentioned in the book, the invariants that we want to satisfy are:

1. Sentinels are neither added or removed (ie. the head and the tail)
2. Nodes are sorted by key
3. Keys are unique

5.1.2 Solution

Our first approach to solving this problem is then through a coarse grained list. This means that we use a lock in such a way that any access to the list is sequential because this lock must be hold and there is only one lock.

Let's take a quick look at the add and remove items to grasp the idea behind this algorithm. Notice how the whole logic of the *add()* and *remove()* operations are guarded by the calls to *lock()* and *unlock()*.

```
1  /**
2   * Add an element.
3   * @param item element to add
4   * @return true iff element was not there already
5   */
6
7  public boolean add(T item) {
8      Node pred, curr;
9      int key = item.hashCode();
10     lock.lock();
11     try {
12         pred = head;
13         curr = pred.next;
14         while (curr.key < key) {
15             pred = curr;
16             curr = curr.next;
17         }
18         if (key == curr.key) {
```

```

19         return false;
20     } else {
21         Node node = new Node(item);
22         node.next = curr;
23         pred.next = node;
24         return true;
25     }
26 } finally {
27     lock.unlock();
28 }
29 }
30
31 /**
32  * Remove an element.
33  * @param item element to remove
34  * @return true iff element was present
35  */
36 public boolean remove(T item) {
37     Node pred, curr;
38     int key = item.hashCode();
39     lock.lock();
40     try {
41         pred = this.head;
42         curr = pred.next;
43         while (curr.key < key) {
44             pred = curr;
45             curr = curr.next;
46         }
47         if (key == curr.key) { // present
48             pred.next = curr.next;
49             return true;
50         } else {
51             return false; // not present
52         }
53     } finally { // always unlock
54         lock.unlock();
55     }
56 }

```

5.1.3 Experiment Description

We have four test cases:

1. *testSequential()*. The main thread inserts multiple elements to the list, then checks if all elements are there and finally removes all elements one by one.
2. *testParallelAdd()*. The main thread spawns multiple threads that add elements to the list concurrently. When they are done, the main thread checks if all elements are in there and finally removes all elements.
3. *testParallelRemove()*. The main thread inserts multiple elements to the list, then checks if all elements are in there and finally it spawns multiple threads that remove the elements from the list concurrently.
4. *testParallelBoth()*. The main thread creates multiple threads. Some of them insert elements to the list and some of them remove elements.

5.1.4 Sample Results

For this test, we saw that it failed consistently with the following error:

```
[oraadm@gdlaa008 Lists]$ junit lists.CoarseListTest
.sequential add, contains, and remove
.parallel both
Exception in thread "Thread-3" Exception in thread "Thread-5" Exception in
thread "Thread-9" Exception in thread "Thread-15"
junit.framework.AssertionFailedError: RemoveThread: duplicate remove: 32
    at junit.framework.Assert.fail(Assert.java:57)
    at junit.framework.TestCase.fail(TestCase.java:227)
    at lists.CoarseListTest$RemoveThread.run(CoarseListTest.java:145)
junit.framework.AssertionFailedError: RemoveThread: duplicate remove: 121
    at junit.framework.Assert.fail(Assert.java:57)
    at junit.framework.TestCase.fail(TestCase.java:227)
    at lists.CoarseListTest$RemoveThread.run(CoarseListTest.java:145)
junit.framework.AssertionFailedError: RemoveThread: duplicate remove: 17
    at junit.framework.Assert.fail(Assert.java:57)
    at junit.framework.TestCase.fail(TestCase.java:227)
    at lists.CoarseListTest$RemoveThread.run(CoarseListTest.java:145)
junit.framework.AssertionFailedError: RemoveThread: duplicate remove: 66
    at junit.framework.Assert.fail(Assert.java:57)
    at junit.framework.TestCase.fail(TestCase.java:227)
    at lists.CoarseListTest$RemoveThread.run(CoarseListTest.java:145)
.parallel add
.parallel remove

Time: 0.017

OK (4 tests)
```

So, the test that fails consistently is *parallel both*. It complains saying that an specific element was removed multiple times by one or more threads.

5.1.5 Proposed solution

This failure turns out to be a test issue. Each remover thread is in charge of removing a specific set of elements, say from 25 to 50. In the same way, an inserter thread is in charge of adding a specific set of elements. Since the threads for a given interval are spawned concurrently, it might be the case that a remover thread is trying to remove an element that hasn't been inserted yet.

One way of solving this problem is by asking the remover thread to retry till it is successful. (Obviously this fix should be improved further since it might lead to a hang in user's application – Imagine the case where an specific element is never inserted. However, for this specific case, the fix is good enough.)

This is the code of the remover thread before the fix:

```
1  class removethread extends thread {
2      int value;
3      removethread(int i) {
4          value = i;
5      }
6      public void run() {
7          for (int i = 0; i < per_thread; i++) {
8              if (!instance.remove(value + i)) {
9                  fail("removethread: duplicate remove: " + (value + i));
10             }
11         }
12     }
13 }
```

And here is the code with the fix:

```
1  class removethread extends thread {
2      int value;
3      removethread(int i) {
4          value = i;
5      }
6      public void run() {
7          for (int i = 0; i < per_thread; i++) {
8              while (!instance.remove(value + i)) {}
9          }
10     }
11 }
```

11 }

And here is the test output after the fix:

```
[oraadm@gdlaa008 Lists]$ junit lists.CoarseListTest
.sequential add, contains, and remove
.parallel both
.parallel add
.parallel remove

Time: 0.016

OK (4 tests)
```

5.1.6 Interpretation

In this experiment we observed how to implement a simple list that allows concurrent access by making all the access sequential. We saw that even-though the algorithm is correct, a correct usage of the API is also required.

5.2 Fine-Grained Synchronization List

5.2.1 Particular Case

The Fine-Grained synchronization improves concurrency by locking individual nodes, rather than locking the list as a whole. Instead of placing a lock on the entire list, we add a Lock to each node, along with *lock()* and *unlock()* methods.

5.2.2 Solution

Just as in the coarse-grained list, *remove()* makes currA unreachable by setting predA next field to currA successor. To be safe, *remove()* must lock both predA and currA. To guarantee progress, it is important that all methods acquire locks in the same order, starting at the head and following next references toward the tail. The *FineList* class *add()* method uses hand-over-hand locking to traverse the list. The finally blocks release locks before returning.

```
1  public boolean add(T item) {
2      int key = item.hashCode();
3      head.lock();
4      Node pred = head;
5      try {
6          Node curr = pred.next;
7          curr.lock();
8          try {
9              while (curr.key < key) {
10                 pred.unlock();
11                 pred = curr;
12                 curr = curr.next;
13                 curr.lock();
14             }
15             if (curr.key == key) {
16                 return false;
17             }
18             Node newNode = new Node(item);
19             newNode.next = curr;
20             pred.next = newNode;
21             return true;
22         } finally {
23             curr.unlock();
24         }
25     }
```



```

25     } finally {
26         pred.unlock();
27     }
28 }

```

```

1  public boolean remove(T item) {
2      Node pred = null, curr = null;
3      int key = item.hashCode();
4      head.lock();
5      try {
6          pred = head;
7          curr = pred.next;
8          curr.lock();
9          try {
10             while (curr.key < key) {
11                 pred.unlock();
12                 pred = curr;
13                 curr = curr.next;
14                 curr.lock();
15             }
16             if (curr.key == key) {
17                 pred.next = curr.next;
18                 return true;
19             }
20             return false;
21         } finally {
22             curr.unlock();
23         }
24     } finally {
25         pred.unlock();
26     }
27 }

```

5.2.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to add and remove values from a list. All threads have to cooperate in order to parallel add and remove values from the list while avoiding duplicate remove of the values. If that is not the case, a fail will be raised.

5.2.4 Observations and Interpretations

The tests executed as expected and no errors where found.

```
1 Tests run: 4, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   0,183 sec
2
3 ----- Standard Output -----
4 sequential add, contains, and remove
5 parallel add
6 parallel remove
7 parallel both
8 -----
9 test-single:
10 BUILD SUCCESSFUL (total time: 2 seconds)
```

5.3 Optimistic Synchronization

5.3.1 Particular Case

In this experiment, the particular case we are trying to study an Optimistic Synchronization which to reduce synchronization costs we take a chance by, search without acquiring locks, lock the nodes found, and then confirm that the locked nodes are correct.

5.3.2 Solution

The *OptimisticList* *add()* method traverses the list ignoring locks, acquires locks, and validates before adding the new node, we must validate and guarantee freedom from interference by checking that *predA* points to *currA* and is reachable from *head*.

```
1  public boolean add(T item) {
2      int key = item.hashCode();
3      while (true) {
4          Entry pred = this.head;
5          Entry curr = pred.next;
6          while (curr.key <= key) {
7              pred = curr; curr = curr.next;
8          }
9          pred.lock(); curr.lock();
10         try {
11             if (validate(pred, curr)) {
12                 if (curr.key == key) { // present
13                     return false;
14                 } else { // not present
15                     Entry entry = new Entry(item);
16                     entry.next = curr;
17                     pred.next = entry;
18                     return true;
19                 }
20             }
21         } finally { // always unlock
22             pred.unlock(); curr.unlock();
23         }
24     }
25 }
```

The *remove()* method traverses ignoring locks, acquires locks, and validates before removing the node.

```

1  public boolean remove(T item) {
2      int key = item.hashCode();
3      while (true) {
4          Entry pred = this.head;
5          Entry curr = pred.next;
6          while (curr.key < key) {
7              pred = curr; curr = curr.next;
8          }
9          pred.lock(); curr.lock();
10         try {
11             if (validate(pred, curr)) {
12                 if (curr.key == key) { // present in list
13                     pred.next = curr.next;
14                     return true;
15                 } else { // not present in list
16                     return false;
17                 }
18             }
19         } finally { // always unlock
20             pred.unlock(); curr.unlock();
21         }
22     }
23 }

```

5.3.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to add and remove values from a list. All threads have to cooperate in order to parallel add and remove values from the list while avoiding duplicate remove of the values. If that is not the case, a fail will be raised.

5.3.4 Observations and Interpretations

When executing the test duplicate removes where found.

```

1  junit.framework.AssertionFailedError: RemoveThread: duplicate
   remove:

```

5.3.5 Proposed changes to fix the problem

The test approach is naive, since it assumes that data would exist always prior deletion, thus not handling the case when there is no data to remove. By doing a retry until there is data to remove the problem seems to be fix.

```
1  public boolean remove(T item) {
2      int key = item.hashCode();
3      boolean bRetry = true;
4      while (true) {
5          Entry pred = this.head;
6          Entry curr = pred.next;
7          while (curr.key < key) {
8              pred = curr; curr = curr.next;
9          }
10         pred.lock(); curr.lock();
11         try {
12             while (bRetry) {
13                 if (validate(pred, curr)) {
14                     bRetry = false;
15                     if (curr.key == key) { // present in list
16                         pred.next = curr.next;
17                         return true;
18                     } else { // not present in list
19                         return false;
20                     }
21                 }
22             }
23         } finally { // always unlock
24             pred.unlock(); curr.unlock();
25         }
26     }
27 }
28 }
```

Once the change is made the execution works fine on every equipment used.

```
1  Tests run: 4, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   0,071 sec
2
3  _____ Standard Output _____
4  sequential add, contains, and remove
5  parallel add
6  parallel remove
7  parallel both
8  _____
```

```
9 | test-single:
10 | BUILD SUCCESSFUL (total time: 3 seconds)
```

5.3.6 Proposed solution

This error may imply a bad scheduling of the list. Apparently old JVM depending on the OS, executed using a logic that avoided this from happening by working in a round robin fashion.

5.4 LazyListTest

5.4.1 Particular Case (problem)

The problem is that of gradually improving what coarse grain locking offers for concurrent data structures like sets (implemented with linked lists).

5.4.2 Solution

The *LazyList* solution is a refinement of the *OptimisticList* solution which does not lock while searching, but locks one it finds the interesting nodes (and then confirms that the locked nodes are correct). As one drawback of *OptimisticList* is that the most common method *contains* method locks, the next logical improvement is to make this method wait-free while keeping the *add* and *remove* methods locking (but reducing their transversings of the list from two to just one).

The refinement mentioned above is precisely that of *LazyList*, which adds a new bit to each node to indicate whether they still belong to the set or not (this prevents transversing the list to detect if the node is reachable, as the new bit introduces such invariant: if a transversing thread does not find a node or it is marked in this bit, then the corresponding item does not belong to the set. This behavior implies that *contains* method does a single wait-free transversal of the list.

For adding an element to the list, *add* method traverses the list, locks the target's predecessor and successor nodes, and inserts the new node in between. The *remove* method is lazy (hence the name of the solution), as it splits its task in two parts: first marks the node in the new bit, logically removing it; and second, update its predecessor's next field, physically removing it.

The three methods ignore the locks while transversing the list, possibly passing over both logically and physically deleted nodes. The *add* and *remove* methods still lock *pred* and *curr* nodes as with *OptimisticList* solution, but the validation reduces to check that *curr* node has not been marked; as well as validating the same for *pred* node, and that it still points to *curr* (validating a couple of nodes is much better than transversing whole list

though). Finally, the introduction of logical removals implies a new contract for detecting that an item still belongs to set: it does so, if still referred by an unmarked reachable node.

The most relevant methods, *remove* and *contains*, are listed below:

```

1  public boolean remove(T item) {
2      int key = item.hashCode();
3      while (true) {
4          Node pred = this.head;
5          Node curr = head.next;
6          while (curr.key < key) {
7              pred = curr; curr = curr.next;
8          }
9          pred.lock();
10         try {
11             curr.lock();
12             try {
13                 if (validate(pred, curr)) {
14                     if (curr.key != key) { // present
15                         return false;
16                     } else { // absent
17                         curr.marked = true; // logically remove
18                         pred.next = curr.next; // physically remove
19                         return true;
20                     }
21                 }
22             } finally { // always unlock curr
23                 curr.unlock();
24             }
25         } finally { // always unlock pred
26             pred.unlock();
27         }
28     }
29 }
30
31 public boolean contains(T item) {
32     int key = item.hashCode();
33     Node curr = this.head;
34     while (curr.key < key)
35         curr = curr.next;
36     return curr.key == key && !curr.marked;
37 }

```


5.4.3 Experiment Description

The test is pretty much the same described for *LockFreeQueueTest*, with a few differences.

- The data structure here is a set, rather than a queue.
- The exception that it uses 8 threads instead of two for each operation (*add* / *remove*).
- The threads that remove elements do not care only in successfully removing certain number of times (like with the queue); here they expect to remove a particular subset of the values.

5.4.4 Observations and Interpretations

The test works as expected on a two cores machine, sample output below:

```
.parallel deq
.parallel both
.sequential push and pop
.parallel enq
```

Time: 0.03

OK (4 tests)

Interestingly, on a 24 cores machine, sometimes the test case *testParallelBoth* fails with exceptions like the one below:

```
junit.framework.AssertionFailedError: RemoveThread: duplicate remove
at junit.framework.Assert.fail(Assert.java:57)
at junit.framework.TestCase.fail(TestCase.java:227)
at lists.LazyListTest$RemoveThread.run(LazyListTest.java:142)
```

While debugging the error above, we found that the message of “duplicate remove” is a bit misleading; is not really that someone else tried to delete that value (as each *RemoveThread* cares about a unique set of values). The

real problem is that the removing threads just try once to remove each value, and fail if they did not find any of them. Since both the adder and remover threads are started concurrently, there is no guarantee that the adders will come first than the removers; so it could be that the removers try to pull out something that has not been inserted yet (leaving to the exception shown above).

Since the *LazyList* solution does not include an error-and-retry approach (as with our second rewrite of the *LockFreeQueue*, which used Java condition's await methods), the only way to fix this would be to rewrite the test program itself. Each remover thread will need to indefinitely try to remove all its items until completion, rather than expecting that all of them are available by the time they are to be removed. We tried that approach and made the proper adjustments to the test program, after which the problem got solved as expected.

5.5 Lock-Free List

5.5.1 Particular Case

We have seen some algorithms for implementing a concurrent list. Now it is time to see how a Lock-Free implementation can be done.

5.5.2 Solution

5.5.3 Experiment Description

Again, we have the same four test cases that we have been dealing with:

1. *testSequential()*. The main thread inserts multiple elements to the list, then checks if all elements are there and finally removes all elements one by one.
2. *testParallelAdd()*. The main thread spawns multiple threads that add elements to the list concurrently. When they are done, the main thread checks if all elements are in there and finally removes all elements.
3. *testParallelRemove()*. The main thread inserts multiple elements to the list, then checks if all elements are in there and finally it spawns multiple threads that remove the elements from the list concurrently.
4. *testParallelBoth()*. The main thread creates multiple threads. Some of them insert elements to the list and some of them remove elements.

5.5.4 Sample Results

Similar to the other implementations, the test cases turns out to have a problem. Here is the output:

```
[oraadm@gdlaa008 Lists]$ junit lists.LockFreeListTest
.sequential add, contains, and remove
.parallel both
Exception in thread "Thread-13" junit.framework.AssertionFailedError:
RemoveThread: duplicate remove: 109
    at junit.framework.Assert.fail(Assert.java:57)
    at junit.framework.TestCase.fail(TestCase.java:227)
    at lists.LockFreeListTest$RemoveThread.run(LockFreeListTest.java:145)
.parallel add
.parallel remove
```

```
Time: 0.022
```

```
OK (4 tests)
```

The fix for this is the same that has been discussed before. Here is the test output after the fix:

```
[oraadm@gdlaa008 Lists]$ junit lists.LockFreeListTest  
.sequential add, contains, and remove  
.parallel both  
.parallel add  
.parallel remove
```

```
Time: 0.016
```

```
OK (4 tests)
```

5.5.5 Interpretation

We observed a way in which a list could be implemented allowing it to be lock free. When using the same set of unit tests as before, we see that the implementation works correctly.

Chapter 6

Concurrent Queues and the ABA Problem

6.1 LockFreeQueueTest Test

6.1.1 Particular Case

In this exercise we deal with the problem of implementing a Queue data structure with the particular characteristic of being lock free.

6.1.2 Solution

Here we present the implementation of the Lock Free Queue that is based on the paper by Michael and Scott.

Here is the *enqueue()* method:

```
1  /**
2   * Append item to end of queue.
3   * @param item
4   */
5  public void enq(T item) {
6      if (item == null) throw new NullPointerException();
7      Node node = new Node(item); // allocate & initialize new node
8      while (true) {              // keep trying
9          Node last = tail.get();  // read tail
10         Node next = last.next.get(); // read next
11         if (last == tail.get()) { // are they consistent?
12             AtomicReference[] target = {last.next, tail};
13             Object[] expect = {next, last};
14             Object[] update = {node, node};
15             if (multiCompareAndSet(
16                 (AtomicReference<T>[]) target,
17                 (T[]) expect,
18                 (T[]) update)){
19                 return;
20             }
21         }
22     }
23 }
```

Here is the dequeue method.

```
1  /**
2   * Remove and return head of queue.
```

```

3      * @return remove first item in queue
4      * @throws queue.EmptyException
5      */
6      public T deq() throws EmptyException {
7          while (true) {
8              Node first = head.get();
9              Node last = tail.get();
10             Node next = first.next.get();
11             if (first == head.get()) {// are they consistent?
12                 if (first == last) {// is queue empty or tail falling behind?
13                     if (next == null) {// is queue empty?
14                         throw new EmptyException();
15                     }
16                     // tail is behind, try to advance
17                     tail.compareAndSet(last, next);
18                 } else {
19                     T value = next.value; // read value before dequeuing
20                     if (head.compareAndSet(first, next))
21                         return value;
22                 }
23             }
24         }

```

6.1.3 Experiment Description

The test consists of 4 test cases. Let us briefly describe each one of them:

- **testSequential.** In this test case 512 elements are enqueued in the data structure. After that, each of the elements are dequeued. At each dequeue, we verify that the dequeued elements are in sequential order (just as they were inserted).
- **testParallelEnqueue.** This test cases spawns 8 enqueueer threads. Each of them will add 512 elements to the queue. When they are done, the main thread dequeues each of the elements performing the same check described above.
- **testParallelDequeue.** In this test case, the master thread enqueues 512 elements. When done, 8 dequeuer threads are spawned. The job of each of these dequeuer threads is to remove the elements. Additionally, an array that keeps track of which elements have been already dequeued

is maintained. The test checks that no element is removed more than once.

- `testParallelBoth`. This is the most complete test in the suite. It spawns 8 enqueueer threads and 8 dequeueer threads. Each enqueueer adds 512 elements to the queue. The dequeueers must remove each of these elements. A check similar to the one described in the *testParallelDequeue* test is performed.

6.1.4 Sample Results

6.1.5 Interpretation

6.1.6 Proposed solution

6.2 HW Queue

6.2.1 Particular Case

The HW Queue uses the Atomic values defined by the Java language in order to implement a locking free queue.

6.2.2 Solution

The unbounded queue implementation used by the *HWQueue* class is blocking, meaning that the *deq()* method does not return until it has found an item to dequeue.

```
1  public T deq() {  
2      while (true) {  
3          int range = tail.get();  
4          for (int i = 0; i < range; i++) {  
5              T value = items[i].getAndSet(null);  
6              if (value != null) {  
7                  return value;  
8              }  
9          }  
10     }  
11 }
```

The *enq()* method uses the *AtomicInteger getAndIncrement()* method that atomically incremented the counter.

```
1  public void enq(T x) {  
2      int i = tail.getAndIncrement();  
3      items[i].set(x);  
4  }
```

6.2.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to add and remove values from a list. All threads have to cooperate in order to parallel add and remove values from the list while avoiding duplicate remove of the values. If that is not the case, a fail will be raised.

6.2.4 Observations and Interpretations

The tests executed as expected and no errors where found.

```
1 Tests run: 4, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   0,121 sec
2
3 ----- Standard Output -----
4 sequential push and pop
5 parallel enq
6 parallel deq
7 parallel both
8 -----
9 test-single:
10 BUILD SUCCESSFUL (total time: 1 second)
```

6.3 Bounded Partial Queue

6.3.1 Particular Case

In this experiment we implement a bounded queue as a linked list and using the Java *ReentrantLock* implementation for the enqueue-er and dequeue-er explicit call of the lock.

6.3.2 Solution

According to the theory there is a limited concurrency we can expect a bounded queue implementation with multiple concurrent enqueue-ers and dequeue-ers to provide, but can extend the `enq()` and `deq()` methods to operate on opposite ends of the queue, so as long as the queue is neither full nor empty, an `enq()` call and a `deq()` call should, in principle, be able to proceed without interference.

The `deq()` method reads the size field to check whether the queue is empty. If so, the dequeue-er must wait until an item is enqueued. The dequeue-er waits on *notEmptyCondition*, which temporarily releases *deqLock*, and blocks until the condition is signaled.

```
1  public T deq() {
2      T result;
3      boolean mustWakeEnqueuers = true;
4      deqLock.lock();
5      try {
6          while (size.get() == capacity) {
7              try {
8                  notEmptyCondition.await();
9              } catch (InterruptedException ex) {}
10         }
11         result = head.next.value;
12         head = head.next;
13         if (size.getAndIncrement() == 0) {
14             mustWakeEnqueuers = true;
15         }
16     } finally {
17         deqLock.unlock();
18     }
19     if (mustWakeEnqueuers) {
20         enqLock.lock();
21         try {
```

```

22         notFullCondition.signalAll();
23     } finally {
24         enqLock.unlock();
25     }
26 }
27 return result;
28 }

```

The *enq()* method thread acquires the *enqLock*, and reads the *size* field. While that field is equal to the capacity, the queue is full, and the enqueue-er must wait until a dequeue-er makes room. The enqueue-er waits on the *notFullCondition* field, releasing the enqueue lock temporarily, and blocking until that condition is signalled.

```

1  public void enq(T x) {
2      if (x == null) throw new NullPointerException();
3      boolean mustWakeDequeuers = false;
4      enqLock.lock();
5      try {
6          while (size.get() == 0) {
7              try {
8                  notFullCondition.await();
9              } catch (InterruptedException e) {}
10         }
11         Entry e = new Entry(x);
12         tail.next = e;
13         tail = e;
14         if (size.getAndDecrement() == capacity) {
15             mustWakeDequeuers = true;
16         }
17     } finally {
18         enqLock.unlock();
19     }
20     if (mustWakeDequeuers) {
21         deqLock.lock();
22         try {
23             notEmptyCondition.signalAll();
24         } finally {
25             deqLock.unlock();
26         }
27     }
28 }

```

6.3.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to *enq()* and *deq()* a range of numbers. All threads have to cooperate to add and remove elements from the queue. Each of the threads will enqueue and dequeue values into the queue, if everything works according to the test there will be mutual exclusion and the mapping of elements will corresponds to the queue elements. If that is not the case, a duplicate fail will be raised.

6.3.4 Observations and Interpretations

The tests executed as expected and no errors where found. Since the ReentrantLock is used to acquire an explicit lock, the mutual exclusion is guarantee by the Java `java.util.concurrent.locks` package.

```
1 Tests run: 4, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   0,076 sec
2
3 ----- Standard Output -----
4 sequential push and pop
5 parallel enq
6 parallel deq
7 parallel both
8 -----
9 test-single:
10 BUILD SUCCESSFUL (total time: 1 second)
```

6.4 Unbounded Queue

6.4.1 Particular Case

In this experiment we implement a unbounded queue, the representation is the same as the bounded queue, except there is no need to count the number of items in the queue, or to provide conditions on which to wait.

6.4.2 Solution

According to the theory an item is actually enqueued when the *enq()* call sets the last nodes next field to the new node, even before *enq()* resets tail to refer to the new node. After that instant, the new item is reachable along a chain of the next references. The actual head is the successor of the node referenced by head, and the actual tail is the last item reachable from the head. Both the *enq()* and *deq()* methods are total as they do not wait for the queue to become empty or full.

```
1  public T deq() throws EmptyException {
2      T result;
3      deqLock.lock();
4      try {
5          if (head.next == null) {
6              throw new EmptyException();
7          }
8          result = head.next.value;
9          head = head.next;
10     } finally {
11         deqLock.unlock();
12     }
13     return result;
14 }
```

```
1  public void enq(T x) {
2      if (x == null) throw new NullPointerException();
3      enqLock.lock();
4      try {
5          Node e = new Node(x);
6          tail.next = e;
7          tail = e;
8      } finally {
9          enqLock.unlock();
10     }
}
```

11 | }

This queue cannot deadlock, because each method acquires only one lock, either *enqLock* or *deqLock* and is much simpler than the bounded queue.

6.4.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to *enq()* and *deq()* a range of numbers. All threads have to cooperate to add and remove elements from the queue. Each of the threads will enqueue and dequeue values into the queue, if everything works according to the test there will be mutual exclusion and the mapping of elements will corresponds to the queue elements. If that is not the case, a duplicate fail will be raised.

6.4.4 Observations and Interpretations

The tests executed as expected and no errors where found. Since the *ReentrantLock* is used to acquire an explicit lock, the mutual exclusion is guarantee by the Java `java.util.concurrent.locks` package. As a interesting note, once each 100 executions the running time increases exponentially as shown below.

```
1 Tests run: 4, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   10,102 sec
2
3 ----- Standard Output -----
4 sequential push and pop
5 parallel enq
6 parallel deq
7 parallel both
8 -----
9 test-single:
10 BUILD SUCCESSFUL (total time: 11 seconds)
```

6.5 SynchronousDualQueueTest

6.5.1 Particular Case (problem)

The problem is to reduce the synchronization overhead of a synchronous queue.

6.5.2 Solution

The solution is given by using what is called a dual data structure; which splits the *enq* and *deq* operations in two parts. When a dequeuer tries to remove an item from the queue, it inserts a reservation object indicating that it is waiting for an enqueueer. Later when an enqueueer thread realizes about the reservation, it fulfills the same by depositing an item and setting the reservation's flag. On the same way, an enqueueer thread can make another reservation when it wants to add an item and spin on its reservation's flag.

This solution has some nice properties:

- Waiting threads can spin on a locally cached flag (scalability).
- Ensures fairness in a natural way. Reservations are queued in the order they arrive.
- This data structure is linearizable, since each partial method call can be ordered when it is fulfilled.

6.5.3 Experiment Description

The test creates 16 threads, grouping them on enqueueers and dequeuers; for each group it divides the workload (512) evenly. Then each thread proceeds to either enqueue or dequeue, as many times as its share of the workload indicated (512/8). There are some assertions to ensure that we do not repeat values (each dequeuer marks an array at the index of the value it got). among them.

6.5.4 Observations and Interpretations

The test runs normally most of the times, but in some occasions it hangs (in the test machine with two cores is a rare error, but on the one with 24 cores it almost always occurs) . During those cases we can see a null pointer exception, and some never ending enqueueers:

```
.parallel both
Exception in thread "Thread-9" java.lang.NullPointerException
    at queue.SynchronousDualQueueTest$DeqThread.run(SynchronousDualQueueTest.java:67)
2015-10-18 23:20:41
Full thread dump Java HotSpot(TM) 64-Bit Server VM (25.60-b23 mixed mode):

"Thread-12" #21 prio=5 os_prio=0 tid=0x00007f4314112800 nid=0x75cf runnable [0x00007f42fc980000]
    java.lang.Thread.State: RUNNABLE
        at queue.SynchronousDualQueue.enq(SynchronousDualQueue.java:49)
        at queue.SynchronousDualQueueTest$EnqThread.run(SynchronousDualQueueTest.java:60)

"Thread-8" #17 prio=5 os_prio=0 tid=0x00007f431410e800 nid=0x75cd runnable [0x00007f42fcb82000]
    java.lang.Thread.State: RUNNABLE
        at queue.SynchronousDualQueue.enq(SynchronousDualQueue.java:49)
        at queue.SynchronousDualQueueTest$EnqThread.run(SynchronousDualQueueTest.java:60)
```

The *NullPointerException* and the hanging enqueueers are related; the former appears cause the *deq* method returns null hence it fails to cast into integer for a dequeuer thread (leaving the overall situation unbalanced, as now nobody will consume some of the items being pushed into the queue).

The queue is populated with integers, so the *null* value returned must be either a defect of the tested class *SynchronousDualQueue*. or that the test case is badly designed. On either case, is shall be related with concurrency, as error occurs much more often on the machine with 24 cores. The easiest way to handle this is to modify the test to handle *null* values on dequeuers thread class:

```

class DeqThread extends Thread {
    public void run() {
        for (int i = 0; i < PER_THREAD; i++) {
            Object v = null;
            while ( (v = instance.deq()) == null );
            int value = (Integer) v;
            if (map[value]) {
                fail("DeqThread: _duplicate _pop");
            }
            map[value] = true;
        }
    }
}

```

With above modification, the test does not hang anymore (tried in both test machines, with 2 and 24 cores). Given more time, we would have preferred to dig further into root cause of this problem; instead of just putting a patch on the test program.

Now, if we dig deeper into the explanation, we can remember that the solution dos pairing between enqueueers and dequeuers; but if the dequeuers come first and there is no data, the *deq* method will return null, and the test is just not prepared to handle such data (it will need an adaptation like the one explained above). Since a dequeuer dies, there will be enqueueers that remain waiting forever for their pairing mates. This should be the reason behind the hanging.

6.6 SynchronousQueueTest Test

6.6.1 Particular Case

In this exercise we are dealing again with a first-in-first-out (FIFO) data structure, ie a Queue. The difference with respect of the previous examples is that in this case we have multiple producers and multiple consumers *rendezvousing* with one another. In other words, a producer that puts an item in the queue blocks until that item is removed by a consumer.

6.6.2 Solution

The following is a naive implementation of a Synchronous queue. Notice that it contains the following fields:

```
1  public class SynchronousQueue<T> {
2      T item = null;
3      boolean enqueueing;
4      Lock lock;
5      Condition condition;
6      ...
7  }
```

Let us talk about them:

- item. The first item in the queue (the one to be dequeued)
- enqueueing. Used to sync the enqueueers
- lock. Used for mutual exclusion
- condition. Used to block partial methods

Now, let us talk about how the enqueue and the dequeue actually work:

```
1  public T deq() throws InterruptedException {
2      lock.lock();
3      try {
4          while (item == null) {
5              condition.await();
6          }
7      }
```

```

7         T t = item;
8         item = null;
9         condition.signalAll();
10        return t;
11    } finally {
12        lock.unlock();
13    }
14 }
15
16 public void enq(T value) throws InterruptedException {
17     if (value == null) throw new NullPointerException();
18     try {
19         while (enqueueing) { // another enqueueer's turn
20             condition.await();
21         }
22         enqueueing = true; // my turn starts
23         item = value;
24         condition.signalAll();
25         while (item != null) {
26             condition.await();
27         }
28         enqueueing = false; // my turn ends
29         condition.signalAll();
30     } finally {
31         lock.unlock();
32     }
33 }
34 }

```

The enqueue method first has to check whether another enqueueer is enqueueing an element. If that is the case, then the enqueueer has to wait till the enqueueer is paired with a dequeuer. After that, it will set the flag *enqueueing* to true, sets the item and signals any other thread that is waiting for this one to finish its enqueue.

Once the enqueueer has put its value in the queue, it waits till someone consumes its value (*item* becomes *null*). Then the enqueueer sets the *enqueueing* flag back to false and finishes.

On the other hand, the dequeuer threads after acquiring the lock, they have to wait till an element is put in the queue. When it finds this is true, then they take the item and signal other threads to indicate they have finished dequeuing.

6.6.3 Experiment Description

This exercise consists of only one test and it works as follow:

It creates $2 * 8$ threads. Each of the threads will enqueue 512/8 elements to the queue. At the same time, 8 threads will try to dequeue the elements. When dequeuing, the test will put a value of *true* in the slot *arr[value]*. The test expects that each slot should not be set to *true* twice.

6.6.4 Sample Results

To start with, we notice that they consistently failes with the following error:

```
[oraadm@gdlaa008 Queue]$ junit queue.SynchronousQueueTest
.parallel both
Exception in thread "Thread-10" Exception in thread "Thread-6" Exception in thread "Thread-2" Exception in th
read "Thread-0" Exception in thread "Thread-8" Exception in thread "Thread-4" Exception in thread "Thread-14"
Exception in thread "Thread-12" java.lang.IllegalMonitorStateException
    at java.util.concurrent.locks.ReentrantLock$Sync.tryRelease(ReentrantLock.java:155)
    at java.util.concurrent.locks.AbstractQueuedSynchronizer.release(AbstractQueuedSynchronizer.java:1260)
    at java.util.concurrent.locks.ReentrantLock.unlock(ReentrantLock.java:460)
    at queue.SynchronousQueue.enq(SynchronousQueue.java:63)
    at queue.SynchronousQueueTest$EnqThread.run(SynchronousQueueTest.java:61)
java.lang.IllegalMonitorStateException
    at java.util.concurrent.locks.ReentrantLock$Sync.tryRelease(ReentrantLock.java:155)
    at java.util.concurrent.locks.AbstractQueuedSynchronizer.release(AbstractQueuedSynchronizer.java:1260)
    at java.util.concurrent.locks.ReentrantLock.unlock(ReentrantLock.java:460)
    at queue.SynchronousQueue.enq(SynchronousQueue.java:63)
    at queue.SynchronousQueueTest$EnqThread.run(SynchronousQueueTest.java:61)
java.lang.IllegalMonitorStateException
    at java.util.concurrent.locks.ReentrantLock$Sync.tryRelease(ReentrantLock.java:155)
    at java.util.concurrent.locks.AbstractQueuedSynchronizer.release(AbstractQueuedSynchronizer.java:1260)
    at java.util.concurrent.locks.ReentrantLock.unlock(ReentrantLock.java:460)
    at queue.SynchronousQueue.enq(SynchronousQueue.java:63)
    at queue.SynchronousQueueTest$EnqThread.run(SynchronousQueueTest.java:61)
java.lang.IllegalMonitorStateException
    at java.util.concurrent.locks.ReentrantLock$Sync.tryRelease(ReentrantLock.java:155)
    at java.util.concurrent.locks.AbstractQueuedSynchronizer.release(AbstractQueuedSynchronizer.java:1260)
    at java.util.concurrent.locks.ReentrantLock.unlock(ReentrantLock.java:460)
    at queue.SynchronousQueue.enq(SynchronousQueue.java:63)
    at queue.SynchronousQueueTest$EnqThread.run(SynchronousQueueTest.java:61)
java.lang.IllegalMonitorStateException
    at java.util.concurrent.locks.ReentrantLock$Sync.tryRelease(ReentrantLock.java:155)
    at java.util.concurrent.locks.AbstractQueuedSynchronizer.release(AbstractQueuedSynchronizer.java:1260)
    at java.util.concurrent.locks.ReentrantLock.unlock(ReentrantLock.java:460)
    at queue.SynchronousQueue.enq(SynchronousQueue.java:63)
    at queue.SynchronousQueueTest$EnqThread.run(SynchronousQueueTest.java:61)
java.lang.IllegalMonitorStateException
    at java.util.concurrent.locks.ReentrantLock$Sync.tryRelease(ReentrantLock.java:155)
    at java.util.concurrent.locks.AbstractQueuedSynchronizer.release(AbstractQueuedSynchronizer.java:1260)
    at java.util.concurrent.locks.ReentrantLock.unlock(ReentrantLock.java:460)
    at queue.SynchronousQueue.enq(SynchronousQueue.java:63)
    at queue.SynchronousQueueTest$EnqThread.run(SynchronousQueueTest.java:61)
```

```

java.lang.IllegalMonitorStateException
    at java.util.concurrent.locks.ReentrantLock$Sync.tryRelease(ReentrantLock.java:155)
    at java.util.concurrent.locks.AbstractQueuedSynchronizer.release(AbstractQueuedSynchronizer.java:1260)
    at java.util.concurrent.locks.ReentrantLock.unlock(ReentrantLock.java:460)
    at queue.SynchronousQueue.enq(SynchronousQueue.java:63)
    at queue.SynchronousQueueTest$EnqThread.run(SynchronousQueueTest.java:61)
java.lang.IllegalMonitorStateException
    at java.util.concurrent.locks.ReentrantLock$Sync.tryRelease(ReentrantLock.java:155)
    at java.util.concurrent.locks.AbstractQueuedSynchronizer.release(AbstractQueuedSynchronizer.java:1260)
    at java.util.concurrent.locks.ReentrantLock.unlock(ReentrantLock.java:460)
    at queue.SynchronousQueue.enq(SynchronousQueue.java:63)
    at queue.SynchronousQueueTest$EnqThread.run(SynchronousQueueTest.java:61)

```

6.6.5 Interpretation

6.6.6 Proposed solution

6.7 LockFreeQueueRecycleTest

6.7.1 Particular Case (problem)

From the generic problem of improving coarse grained lock approaches, the particular approach followed on this exercise corresponds to the extreme case: lock-free data structure. The particular case is for a queue, and it has the additional bonus of recycling memory.

6.7.2 Solution

The solution is based on an 1996 ACM article from Maged M. Michel and Michael L. Scott, who based on previous publications, suggest a new way of implementing a lock-free queue with recycles. They implement the queue with a single-linked list using *tail* and *head* pointers; where *head* always points to a dummy or sentinel node which is the first in the list, and *tail* points to either the last or second to last node in the list. The algorithm uses “compare and swap” (*CAS*) with modification counters to avoid the ABA problem.

Dequeuers are allowed to free dequeued nodes by ensuring that *tail* does not point to the dequeued node nor to any of its predecessors (that is, dequeued nodes may be safely reused). To obtain consistent values of various pointers the authors relied on sequences of reads that re-check earlier values to ensure they have not changed (these reads are claimed to be simpler than snapshots).

6.7.3 Experiment Description

The test is pretty much the same described for *LockFreeQueueTest* (with the exception that it uses 8 threads instead of two).

6.7.4 Observations and Interpretations

The test does not exhibit any pitfall, which suggests that the theory works just fine on the tested machines (we tried the one with 2 and with 24 cores). Sample output below:

```
.parallel enq  
.parallel deq  
.parallel both  
.sequential push and pop
```

Time: 0.023

OK (4 tests)

6.8 LockFreeQueueTest Test

6.8.1 Particular Case

In this exercise we deal with the problem of implementing a Queue data structure with the particular characteristic of being lock free.

6.8.2 Solution

Here we present the implementation of the Lock Free Queue that is based on the paper by Michael and Scott.

Here is the *enqueue()* method:

```
1  /**
2   * Append item to end of queue.
3   * @param item
4   */
5  public void enq(T item) {
6      if (item == null) throw new NullPointerException();
7      Node node = new Node(item); // allocate & initialize new node
8      while (true) {              // keep trying
9          Node last = tail.get(); // read tail
10         Node next = last.next.get(); // read next
11         if (last == tail.get()) { // are they consistent?
12             AtomicReference[] target = {last.next, tail};
13             Object[] expect = {next, last};
14             Object[] update = {node, node};
15             if (multiCompareAndSet(
16                 (AtomicReference<T>[]) target,
17                 (T[]) expect,
18                 (T[]) update)) {
19                 return;
20             }
21         }
22     }
23 }
```

Here is the dequeue method.

```
1  /**
2   * Remove and return head of queue.
```

```

3      * @return remove first item in queue
4      * @throws queue.EmptyException
5      */
6      public T deq() throws EmptyException {
7          while (true) {
8              Node first = head.get();
9              Node last = tail.get();
10             Node next = first.next.get();
11             if (first == head.get()) {// are they consistent?
12                 if (first == last) {// is queue empty or tail falling behind?
13                     if (next == null) {// is queue empty?
14                         throw new EmptyException();
15                     }
16                     // tail is behind, try to advance
17                     tail.compareAndSet(last, next);
18                 } else {
19                     T value = next.value; // read value before dequeuing
20                     if (head.compareAndSet(first, next))
21                         return value;
22                 }
23             }
24         }

```

6.8.3 Experiment Description

The test consists of 4 test cases. Let us briefly describe each one of them:

- **testSequential.** In this test case 512 elements are enqueued in the data structure. After that, each of the elements are dequeued. At each dequeue, we verify that the dequeued elements are in sequential order (just as they were inserted).
- **testParallelEnqueue.** This test cases spawns 8 enqueueer threads. Each of them will add 512 elements to the queue. When they are done, the main thread dequeues each of the elements performing the same check described above.
- **testParallelDequeue.** In this test case, the master thread enqueues 512 elements. When done, 8 dequeuer threads are spawned. The job of each of these dequeuer threads is to remove the elements. Additionally, an array that keeps track of which elements have been already dequeued

is maintained. The test checks that no element is removed more than once.

- `testParallelBoth`. This is the most complete test in the suite. It spawns 8 enqueueer threads and 8 dequeueer threads. Each enqueueer adds 512 elements to the queue. The dequeueers must remove each of these elements. A check similar to the one described in the *testParallelDequeue* test is performed.

6.8.4 Sample Results

6.8.5 Interpretation

6.8.6 Proposed solution

6.9 Bounded Partial Queue

6.9.1 Particular Case

In this experiment we implement a bounded queue as a linked list and using the Java *ReentrantLock* implementation for the enqueue-er and dequeue-er explicit call of the lock.

6.9.2 Solution

According to the theory there is a limited concurrency we can expect a bounded queue implementation with multiple concurrent enqueue-ers and dequeue-ers to provide, but can extend the `enq()` and `deq()` methods to operate on opposite ends of the queue, so as long as the queue is neither full nor empty, an `enq()` call and a `deq()` call should, in principle, be able to proceed without interference.

The `deq()` method reads the size field to check whether the queue is empty. If so, the dequeue-er must wait until an item is enqueued. The dequeue-er waits on *notEmptyCondition*, which temporarily releases *deqLock*, and blocks until the condition is signaled.

```
1  public T deq() {
2      T result;
3      boolean mustWakeEnqueuers = true;
4      deqLock.lock();
5      try {
6          while (size.get() == capacity) {
7              try {
8                  notEmptyCondition.await();
9              } catch (InterruptedException ex) {}
10         }
11         result = head.next.value;
12         head = head.next;
13         if (size.getAndIncrement() == 0) {
14             mustWakeEnqueuers = true;
15         }
16     } finally {
17         deqLock.unlock();
18     }
19     if (mustWakeEnqueuers) {
20         enqLock.lock();
21         try {
```

```

22         notFullCondition.signalAll();
23     } finally {
24         enqLock.unlock();
25     }
26 }
27 return result;
28 }

```

The *enq()* method thread acquires the *enqLock*, and reads the *size* field. While that field is equal to the capacity, the queue is full, and the enqueue-er must wait until a dequeue-er makes room. The enqueue-er waits on the *notFullCondition* field, releasing the enqueue lock temporarily, and blocking until that condition is signalled.

```

1  public void enq(T x) {
2      if (x == null) throw new NullPointerException();
3      boolean mustWakeDequeuers = false;
4      enqLock.lock();
5      try {
6          while (size.get() == 0) {
7              try {
8                  notFullCondition.await();
9              } catch (InterruptedException e) {}
10         }
11         Entry e = new Entry(x);
12         tail.next = e;
13         tail = e;
14         if (size.getAndDecrement() == capacity) {
15             mustWakeDequeuers = true;
16         }
17     } finally {
18         enqLock.unlock();
19     }
20     if (mustWakeDequeuers) {
21         deqLock.lock();
22         try {
23             notEmptyCondition.signalAll();
24         } finally {
25             deqLock.unlock();
26         }
27     }
28 }

```

6.9.3 Experiment Description

The test creates 8 threads that need to be coordinate in order to *enq()* and *deq()* a range of numbers. All threads have to cooperate to add and remove elements from the queue. Each of the threads will enqueue and dequeue values into the queue, if everything works according to the test there will be mutual exclusion and the mapping of elements will corresponds to the queue elements. If that is not the case, a duplicate fail will be raised.

6.9.4 Observations and Interpretations

The tests executed as expected and no errors where found. Since the ReentrantLock is used to acquire an explicit lock, the mutual exclusion is guarantee by the Java `java.util.concurrent.locks` package.

```
1 Tests run: 4, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   0,076 sec
2
3 ----- Standard Output -----
4 sequential push and pop
5 parallel enq
6 parallel deq
7 parallel both
8 -----
9 test-single:
10 BUILD SUCCESSFUL (total time: 1 second)
```

Chapter 7

Concurrent Stacks and Elimination

7.1 Exchanger

7.1.1 Particular Case

The problem we are trying to solve is how to allow two threads exchange values they hold. The idea is that if a Thread A calls an *exchange()* method with a given argument and thread B calls the *exchange()* method with another arguments, then thread A will return B's argument and B will return A's argument.

Let us discuss a little bit why such method is required. The motivation of this comes from the problem of implementing a concurrent queue. Imagine that two threads are trying to access the queue and one of them wants to enqueue an element and the other one wants to dequeue an element. In this situation, we can avoid contention in the top of the stack by letting the threads simply exchange their values. That is, the enqueueer can pass its value to the dequeuer thread. This approach is called the *Elimination backoff stack*.

7.1.2 Solution

The solution in this exercise is based on the *A Scalable Elimination-based Exchange Channel* paper. Since, this algorithm is a bit more involved and it is not discussed in the book, let us describe it in more detail.

The algorithm uses an array of *AtomicReferences*. Note that the size of this array depends on the number of processors in the system. Actually, it has to be of half the number of processors. The reason for this is that if we have n threads, then we can only have up to $n/2$ concurrent exchanges and each slot in the array is used for an exchange.

```
1 public class Exchanger<V> {
2     private static final int SIZE =
3         (Runtime.getRuntime().availableProcessors() + 1) / 2;
4     private static final long BACKOFF_BASE = 128L;
5     static final Object FAIL = new Object();
6     private final AtomicReference[] arena;
7     private final Random random;
8     public Exchanger() {
9         arena = new AtomicReference[SIZE + 1];
10        random = new Random();
11        for (int i = 0; i < arena.length; ++i)
```



```

12         arena[i] = new AtomicReference();
13     }

```

Now, from our program we can call the *exchange()* methods in two ways. One in which we can specify a timeout for the exchange operation and another where there is no timeout.

```

1     public V exchange(V x) throws InterruptedException {
2         try {
3             return (V)doExchange(x, Integer.MAX_VALUE);
4         } catch (TimeoutException cannotHappen) {
5             throw new Error(cannotHappen);
6         }
7     }
8     public V exchange(V x, long timeout, TimeUnit unit)
9     throws InterruptedException, TimeoutException {
10        return (V)doExchange(x, unit.toNanos(timeout));
11    }

```

The interesting method is *doExchange()*. Imagine that a thread first tries to do an exchange. So it calls this method and gets into the while loop. It takes the first slot in the arena and finds it is not occupied and hence it will try to occupy the slot setting its id in the slot using a *compareAndSet* operation. Then it will sit there and wait for someone else to come and exchange its value in the *waitForHole()* call.

Now imagine that another thread comes and tries to exchange its value with the previous thread. It then enters the loop, looks into the first slot and sees that someone has put a request in that slot. In other words, some other thread is willing to exchange a value. To do this, it fills the hole with the requested value and the other thread is signaled to continue. Our current thread then returns with the value of the other thread.

The first thread, as we mentioned, gets signaled and sets the slot to null to indicate that the exchange has finished and returns with the value of the other thread.

```

1     private Object doExchange(Object item, long nanos)
2     throws InterruptedException, TimeoutException {
3         Node me = new Node(item);
4         long lastTime = System.nanoTime();

```

```

5      int idx = 0;
6      int backoff = 0;
7      while (true) {
8          AtomicReference<Node> slot = (AtomicReference<Node>)arena[idx];
9          // If this slot is occupied, an item is waiting...
10         Node you = slot.get();
11         if (you != null) {
12             Object v = you.fillHole(item);
13             slot.compareAndSet(you, null);
14             if (v != FAIL) // ... unless it's cancelled
15                 return v;
16         }
17         // Try to occupy this slot
18         if (slot.compareAndSet(null, me)) {
19             Object v = ((idx == 0)?
20                 me.waitForHole(nanos) :
21                 me.waitForHole(randomDelay(backoff)));
22             slot.compareAndSet(me, null);
23             if (v != FAIL)
24                 return v;
25             if (Thread.interrupted())
26                 throw new InterruptedException();
27             long now = System.nanoTime();
28             nanos -= now - lastTime;
29             lastTime = now;
30             if (nanos <= 0)
31                 throw new TimeoutException();
32             me = new Node(item);
33             if (backoff < SIZE - 1)
34                 ++backoff;
35             idx = 0; // Restart at top
36         } else // Retry with a random non-top slot <= backoff
37             idx = 1 + random.nextInt(backoff + 1);
38     }
39 }
40 ...

```

The algorithm adds some other improvements to the base case already described. The main improvement is the backoff logic which comes to play when two threads are already exchanging their values in a slot. In that case, the thread will try using another slot.

7.1.3 Experiment Description

The experiment here creates an array of integers of size *THREADS*. It then spawns *THREADS* threads. Each thread will exchange its thread id with another one and will record this value in the array. For example, suppose that thread 1 and thread 2 are the participants. In this case thread 2 will store $array[2] = 1$ and thread 1 will store $array[1] = 2$. The condition that the test must satisfy is $i \neq array[array[i]]$. In this case $1 \neq array[array[1]]$.

7.1.4 Sample Results

For this test, we observed that the test always passes:

```
.exchange
OK

Time: 0.007

OK (1 test)
```

7.1.5 Interpretation

We see that the proposal of Scherer et. al works correctly. As we mentioned previously, this algorithm is very important if we want to implement a concurrent queue because this exchanger mechanism allows for what is called *Elimination*.

Chapter 8

Counting, Sorting and Distributed Coordination

8.1 TreeTest

8.1.1 Particular Case (problem)

This problem belongs to chapter 12, where the idea is to present problems which look inherently serial but that surprisingly accept interesting concurrent solutions (though not always easy to explain). The particular problem this test is exercising is that of shared counting, where we have N threads to increase a shared numeric variable up to certain value; but with the goal of producing less contention than a serial solution would generate.

8.1.2 Solution

The java program is called *Tree.java*, although it corresponds to the *CombiningTree* from the book, which is a binary tree structure where each thread is located at a leaf and at most two threads can share a leaf. The shared counter is at the root of the tree, and the rest of the nodes serve as intermediate result points. When a couple of threads collide in their attempt to increment, only one of them will serve as the representative and go up in the tree with the mission of propagating the combined increment (2) up to the shared counter at the root of the tree. In its way up, it may encounter further threads that collide again with it, making it wait or continue its journey to the root.

When a thread reaches the root it adds its accumulated result to the shared counter, and propagates down the news that the job is done to the rest of the threads that waited along the way. This solution to the counting problem has worse latency than lock-based solutions ($O(1)$ vs $O(\log(p))$), where p is the number of threads or processors; but it offers a better throughput.

8.1.3 Experiment Description

The program consists of a single unit test *testGetAndIncrement*, which creates 8 threads to perform each 2^{20} increments. The program uses an auxiliary array called *test* to record the individual increments done by each thread; assertions are made to ensure that none of the attempts results in a duplicated value, as well as to ensure that all threads completed their task.

8.1.4 Observations and Interpretations

The test performs without much controversy in an elapsed time between 3 and 4 seconds (laptop computer with i5 processor). Below a sample output:

```
.Parallel, 8 threads, 1048576 tries
```

```
Time: 3.641
```

To make the test a little more interesting, we created another couple of tests reusing sample template of existing one; difference was the Thread class used on each case.

The first additional test uses a thread class based on Java provider *java.util.concurrent.atomic.AtomicInteger*:

```
cnt2 = new AtomicInteger();

class MyThread2 extends Thread {
    public void run() {
        for (int j = 0; j < TRIES; j++) {
            int i = cnt2.getAndIncrement();
            if (test[i]) {
                System.out.printf("ERROR_duplicate_value_%d\n", i);
            } else {
                test[i] = true;
            }
        }
    }
}
```

while the second additional test was based on a custom class that used synchronized *getAndIncrement* method:

```
class MyThread3 extends Thread {
    public void run() {
        for (int j = 0; j < TRIES; j++) {
            int i = cnt3.getAndIncrement();
            if (test[i]) {
                System.out.printf("ERROR_duplicate_value_%d\n", i);
            } else {
                test[i] = true;
            }
        }
    }
}
```

```

    }
  }
}

class MyCounter {
  private int cnt = 0;

  public synchronized int getAndIncrement()
  {
    return cnt++;
  }
}

```

The expectation was that the test based on a synchronized method would be the slowest (*Parallel3* label on output), the one based on the *CombineTree* would become second (*Parallel1* label on output) and the fastest would be the one based on *AtomicInteger* (*Parallel2*); as the latest is likely to take advantage of hardware atomic instruction of 32bits. Surprisingly, the test based on *CombineTree* was the worst of all, by far:

```

.Parallel11, 8 threads, 1048576 tries took 3878
.Parallel12, 8 threads, 1048576 tries took 162
.Parallel13, 8 threads, 1048576 tries took 723

Time: 4.827

OK (3 tests)

```

Most likely we are not comparing apples with apples, as the theory does not match the experiment; perhaps the *CombineTest* class is not meant for real usage, so this comparison is not fair. Anyway, the test does not hang nor fails on neither the 2 cores nor the 24 machines.

8.2 Balancer Test

8.2.1 Particular Case

The problem we are dealing with in this experiment is that of creating a balancer class.

As stated in the text book, a *Balancer* is a simple switch with two input wires and two output wires (top and bottom wires). The token arrives to the input lines and the balancer must output the input one at a time.

A balancer has two states: *up* and *down*. If the state is *up*, the next token exits on the top of the wire. Otherwise, it exits on the bottom wire.

8.2.2 Solution

The following is a simple implementation of a Balancer class. As described above, the balancer consists of a toggle.

```
1  public synchronized int traverse(int input) {
2      try {
3          if (toggle) {
4              return 0;
5          } else {
6              return 1;
7          }
8      } finally {
9          toggle = !toggle;
10     }
11 }
```

As we can observe, the traverse method simply switches the state of the toggle and outputs either 0 or 1.

8.2.3 Experiment Description

The balancer class described above is used in one test case. The test case consists of a 2-slot array. Each slot keeps track of how many times the balancer has output a signal on the *up* wire and how many on the *down* wire.

So, the test activates the balancer exactly 256 times. The final result must be that each slot has the number 128. If that is the case, we declare the test as passed, otherwise it failed.

8.2.4 Sample Results

8.2.5 Interpretation

8.2.6 Proposed solution

8.3 Bitonic Counting Network

8.3.1 Particular Case

In this experiment we implement a *Bitonic* counting network, a *Bitonic* network is a comparison-based sorting network that can be run in parallel. It focuses on converting a random sequence of numbers into a bitonic sequence, one that monotonically increases, then decreases. We represent this network as switches in a counting network, but as we use shared-memory multiprocessor however, we represent this network as objects in memory mimicking a real network.

8.3.2 Solution

According to the theory each balancer is an object, whose wires are references from one balancer to another. Each thread repeatedly traverses the object, starting on some input wire and emerging at some output wire, effectively shepherding a token through the network. The *Balancer* class has a single Boolean field: *toggle*. The synchronized *traverse()* method complements the *toggle* field and returns as output wire, either 0 or 1.

```
1  public synchronized int traverse(int input) {
2      try {
3          if (toggle) {
4              return 0;
5          } else {
6              return 1;
7          }
8      } finally {
9          toggle = !toggle;
10     }
11 }
```

The *Merger* class has three fields; the *width* field must be a power of 2, *half[]* is a two-element array of half-width *Merger* objects, and *layer[]* is an array of width/2 balancers implementing the final network layer. The class provides a *traverse(i)* method, where *i* is the wire on which the token enters. If the token entered on the lower width/2 wires, then it passes through *half[0]*, otherwise *half[1]*.

```
1  public Merger(int _size) {
2      size = _size;
```

```

3     layer = new Balancer[size / 2];
4     for (int i = 0; i < size / 2; i++) {
5         layer[i] = new Balancer();
6     }
7     if (size > 2) {
8         half = new Merger[] { new Merger(size / 2), new Merger(size / 2)
9         };
10    }
11
12    public int traverse(int input) {
13        int output = 0;
14        if (size > 2) {
15            output = half[input % 2].traverse(input / 2);
16        }
17        return output + layer[output].traverse(0);
18    }

```

The *Bitonic* class provides a *traverse(i)* method. If the token entered on the lower width/2 wires, then it passes through *half[0]*, otherwise *half[1]*. A token that emerges from the half-merger subnetwork on wire *i* then traverses the final merger network from input wire *i*.

8.3.3 Experiment Description

The test creates a 256 array that get values mapped to it and later checks that the steps where properly executes, by validating the result in the *position-1* of each elements the bitonic sorted should have accommodated the array into a proper sequence. If the result doesn't match the mapping a new exception is thrown.

8.3.4 Observations and Interpretations

The tests executed as expected and no errors where found. The bitonic counting network as represented through software is replicating hardware that is well understood and easy to follow.

```

1 Tests run: 2, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   0,052 sec
2
3 test:
4 Deleting: /var/folders/hn/vlmcc1gj1kl7bb9rpqb8tdqr0000gn/T/TEST-
   counting.BitonicTest.xml

```

5 BUILD SUCCESSFUL (total time: 0 seconds)

Chapter 9

Concurrent Hashing and Natural Parallelism

9.1 CoarseHashSetTest Test

9.1.1 Particular Case

The problem we are trying to solve now is that of building a concurrent Hash data structure. In the particular case described in the book, the hash data structure is used to construct a representation of a set ADT. Hence, we ought to provide three operations:

1. add. Aggregate a new element to the set
2. remove. Remove a given element from the set
3. contains. Determine whether a given element is already in the set

9.1.2 Solution

Here we present a first approach to solving the problem stated above. We do so by providing a coarse version. As in other coarse data structures, the problem with this implementation is its poor performance in parallel environments.

Let us briefly describe how this implementation works. First we should mention that we have a base abstract class that implements the three operations mentioned previously. We extend this base class by providing four methods:

- resize. Resizes the hash based on a policy. In this implementation, it simply doubles the size of the hash.
- acquire. Acquires a lock on the hash.
- release. Releases the lock.
- policy. Returns true if each bucket contains more than 4 elements.

```
1  /**
2   * double the set size
3   */
4  public void resize() {
5      int oldCapacity = table.length;
```

```

6      lock.lock();
7      try {
8          if (oldCapacity != table.length) {
9              return; // someone beat us to it
10         }
11         int newCapacity = 2 * oldCapacity;
12         List<T>[] oldTable = table;
13         table = (List<T>[]) new List[newCapacity];
14         for (int i = 0; i < newCapacity; i++)
15             table[i] = new ArrayList<T>();
16         for (List<T> bucket : oldTable) {
17             for (T x : bucket) {
18                 int myBucket = Math.abs(x.hashCode() % table.length);
19                 table[myBucket].add(x);
20             }
21         }
22     } finally {
23         lock.unlock();
24     }
25 }
26 /**
27  * Synchronize before adding, removing, or testing for item
28  * @param x item involved
29  */
30 public final void acquire(T x) {
31     lock.lock();
32 }
33 /**
34  * synchronize after adding, removing, or testing for item
35  * @param x item involved
36  */
37 public void release(T x) {
38     lock.unlock();
39 }
40 public boolean policy() {
41     return size / table.length > 4;
42 }

```

This implementation is said to be coarse because any operation on the hash locks the entire data structure.

9.1.3 Experiment Description

As in other experiments, here we have four test cases:

- `testSequential`. The master thread adds 512 elements to the set, then it verifies that these elements were actually added and finally removes them.
- `testParallelEnq`. It creates 8 *adders* threads. Each of them will add 512 elements to the set. When all of them are done, the master thread checks that all the elements were actually added and finally removes them.
- `testParallelRemove`. In this test case, the master thread adds 512 elements to the set and verifies that the elements were actually added. After that, 8 *remover* threads are spawned. Each of them will remove 512/8 elements.
- `testParallelBoth`. This test is a mix of the two above. 8 *removers* and 8 *adders* are spawned. Removers should remove elements added by the adders.

9.1.4 Sample Results

9.1.5 Interpretation

9.1.6 Proposed solution

9.2 StripedHashSetTest Test

9.2.1 Particular Case

In this exercise we deal with the problem of building a hash data structure that can be accessed by multiple threads without losing performance.

In the coarse version of our hash, we saw that any access to the data structure ended up being serial. So, in other words, in this experiment we want to understand how to create a hash that allows more concurrent accesses.

9.2.2 Solution

A solution for this problem is the *Striped Hash*. The main characteristic of this implementation is that the data structure contains an array of locks. Each lock protect specific sections of the hash table. Here is the implementation of the *acquire()* and *release()* methods:

```
1  /**
2   * Synchronize before adding, removing, or testing for item
3   * @param x item involved
4   */
5  public final void acquire(T x) {
6      int myBucket = Math.abs(x.hashCode() % locks.length);
7      locks[myBucket].lock();
8  }
9
10 /**
11  * synchronize after adding, removing, or testing for item
12  * @param x item involved
13  */
14 public void release(T x) {
15     int myBucket = Math.abs(x.hashCode() % locks.length);
16     locks[myBucket].unlock();
17 }
```

Note that when we acquire or release a lock, first we need to choose which lock we should work with.

9.2.3 Experiment Description

The proposed experiment is exactly the same as the one described in the previous experiment. As mentioned before, there are four variants: *testSequential*, *testParallelEnq*, *testParallelRemove* and *testParallelBoth*.

9.2.4 Sample Results

9.2.5 Interpretation

9.2.6 Proposed solution

9.3 CoarseCuckooHashSetTest

9.3.1 Particular Case (problem)

The problem we try to solve is that of concurrent hashing, when we use the hash to implement a set data structure.

9.3.2 Solution

The solution uses three techniques combined:

- cuckoo-hashing: which consists in shifting to the right elements, by performing consecutive swaps on elements having a collision due hash function.
- double hashing where we use a couple of hash tables; the first one is based on a simple hash function (modulus of the element's hash code), and the second table is indexed with random numbers generated out of a sequence feed from the hash code of elements.
- coarse access: as the whole structure is being protected by mutual exclusion with a lock.

Below the most representative function, *add*, along with the two hash functions used:

```
1  private final int hash0(Object x) {
2      return x.hashCode() % size;
3  }
4
5  private final int hash1(Object x) {
6      random.setSeed(x.hashCode());
7      return random.nextInt(size);
8  }
9
10 public boolean add(T x) {
11     lock.lock();
12     try {
```

```

13         if (contains(x)) {
14             return false;
15         }
16         for (int i = 0; i < LIMIT; i++) {
17             if ((x = swap(0, hash0(x), x)) == null) {
18                 return true;
19             } else if ((x = swap(1, hash1(x), x)) == null) {
20                 return true;
21             }
22         }
23         System.out.println("uh-oh");
24         throw new CuckooException();
25     } finally {
26         lock.unlock();
27     }
28 }

```

Something additional to mention about this solution, is that when adding a new element we constraint the maximum number of swaps to perform (while shifting the elements to the right). In this implementation is *LIMIT* is 32; if after those attempts we could not arrange all elements on its slot we raise an exception.

9.3.3 Experiment Description

The test program consists of three test cases:

- *testSequential*: add 512 elements to the hash set (sequentially), perform a *contains* test for each one (sequentially) and finally remove the 512 elements (sequentially).
- *testParallelAdd*: add 512 elements to the hash set (in parallel with 8 threads), perform a *contains* test for each one (sequentially) and finally remove the 512 elements (sequentially).
- *testParallelRemove*: add 512 elements to the hash set (sequentially), perform a *contains* test for each one (sequentially) and finally remove the 512 elements (in parallel with 8 threads).

- *testParallelBoth*: adds and removes 512 elements to the hash set in parallel with 8 adder threads, and 8 remover threads; the removers complain if someone else removed its designated range of elements.

9.3.4 Observations and Interpretations

The tests works fine on both machines (2 and 24 cores respectively), with the exception that the parallel test remover threads complain when they try to delete an element which does not exist yet (which is an expected consequence of the test). A sample output is shown below:

```
[oraadm@gdlaa008 orig]$ junit hash.CoarseCuckooHashSetTest

.sequential add, contains, and remove
.parallel both
Exception in thread "Thread-5" Exception in thread "Thread-9"
        Exception in thread "Thread-11" Exception in thread
"Thread-1" junit.framework.AssertionFailedError: DeqThread: duplicate

remove
at junit.framework.Assert.fail(Assert.java:57)
at junit.framework.TestCase.fail(TestCase.java:227)
at
...

.parallel add
.parallel remove

Time: 0.053

OK (4 tests)
```

9.4 LockFreeHashSetTest

9.4.1 Particular Case (problem)

The problem we try to solve is that of implementing a lock-free hash set.

9.4.2 Solution

The solution uses a technique called “Recursive Split-Ordering”, where we keep all items on a single lock-free linked list, and we keep separately an increasing array of buckets (which are just references to different positions of the list). While any element could be found in the set by transversing the list, idea is to use hash functions to take advantage of the shortcuts that the buckets represent. The algorithm takes its name from the fact that the split-ordered-keys on the list, are the bitwise reverse representation of their items keys (hash codes).

9.4.3 Experiment Description

The test is exactly the same as that described for *CoarseCuckooHashSetTest*.

9.4.4 Observations and Interpretations

The test runs fine in 3 and 24 core machines, and similarly to *CoarseCuckooHashSetTest*, suffers from a design flag on the test program itself (remover threads may attempt to delete an element which has not been added yet). Below a sample output of its execution:

```
[loraadm@gdlaa008 orig]$ junit hash.LockFreeHashSetTest
.sequential add, contains, and remove
.parallel both
Exception in thread "Thread-1" junit.framework.AssertionFailedError:
    DeqThread: duplicate pop
at junit.framework.Assert.fail(Assert.java:57)
at junit.framework.TestCase.fail(TestCase.java:227)
at
hash.LockFreeHashSetTest$RemoveThread.run(LockFreeHashSetTest.java:143)
.parallel add
.parallel remove

Time: 0.037

OK (4 tests)
```


9.5 Cuckoo Hashing

9.5.1 Particular Case

Cuckoo hashing is a (sequential) hashing algorithm in which a newly added item displaces any earlier item occupying the same slot. In this test we modify the sequential Cuckoo hashing in order to change the sequential hashing algorithm to concurrent hashing.

9.5.2 Solution

We break up each method call into a sequence of phases, where each phase adds, removes, or displaces a single item x . We use a two-entry array *table[]* of tables, and two independent hash functions, (denoted as *hash0()* and *hash1()* in the code) mapping the set of possible keys to entries in the array.

```
1  public TCuckooHashSet(int capacity) {
2      locks = new Lock[2][LOCKS];
3      table = (T[][]) new Object[2][capacity];
4      size = capacity;
5      for (int i = 0; i < 2; i++) {
6          for (int j = 0; j < LOCKS; j++) {
7              locks[i][j] = new ReentrantLock();
8          }
9      }
10 }
11 private final int hash0(Object x) {
12     return Math.abs(x.hashCode() % size);
13 }
14 private final int hash1(Object x) {
15     random.setSeed(x.hashCode());
16     return random.nextInt(size);
17 }
```

To test whether a value x is in the set, *contains(x)* tests whether either *table[0][h0(x)]* or *table[1][h1(x)]* is equal to x . Similarly, *remove(x)* checks whether x is in either *table[0][h0(x)]* or *table[1][h1(x)]*, and removes it if found.

```
1  public boolean remove(T x) {
2      if (x == null) {
3          throw new IllegalArgumentException();
4      }
5  }
```



```

5      int h0 = hash0(x);
6      Lock lock0 = locks[0][h0 % LOCKS];
7      try {
8          lock0.lock();
9          if (x.equals(table[0][h0])) {
10             table[0][h0] = null;
11             return true;
12         } else {
13             int h1 = hash1(x);
14             Lock lock1 = locks[1][h1 % LOCKS];
15             try {
16                 lock1.lock();
17                 if (x.equals(table[1][h1])) {
18                     table[1][h1] = null;
19                     return true;
20                 }
21                 return false;
22             } finally {
23                 lock1.unlock();
24             }
25         }
26     } finally {
27         lock0.unlock();
28     }
29 }

```

The *add(x)* method is the most interesting. It successively removes conflicting items until every key has a slot. To add *x*, the method swaps *x* with *y*, the current occupant of *table[0][h0(x)]*. If the prior value *y* was null, it is done, otherwise, it swaps the newly nestles value *y* for the current occupant of *table[1][h1(y)]* in the same way and continues swapping entries (alternating tables) until it finds an empty slot.

9.5.3 Experiment Description

The test creates 8 threads that perform *enq()* and *deq()* operations in sequential and parallel. The dequeue values are revised and if a value is missing a fail assertion is thrown.

9.5.4 Observations and Interpretations

When executing the test the test fails on the parallel part due to a missing value when executing the parallel both test.

```
1 Exception in thread "Thread-5" Exception in thread "Thread-7"  
  junit.framework.AssertionFailedError: Th 0          DeqThread:  
    missing value: 3
```

9.5.5 Proposed changes to fix the problem

By adding terminal output through System.out in the *remove()* method or Debugging the file the execution works, thus becoming a problem of observation, since the fail assertion only happens when its not observed.

```
1 compile-test:  
2 .sequential add, contains, and remove  
3 .parallel add  
4 .parallel remove  
5 .parallel both  
6  
7 Time: 0.071  
8  
9 OK (4 tests)
```

9.5.6 Proposed solution

It is not clear why the observation would affect the execution, thus it is possible that the issue is a propagation of the values in the memory which are forced to be executed constantly through the debug implementation.

9.6 Refinable Hash Set

9.6.1 Particular Case

In this test we want to refine the granularity of locking as the table size grows, so that the number of locations in a stripe does not continuously grow by using a globally shared owner field that combines a *Boolean* value with a reference to a thread.

9.6.2 Solution

While a resizing is in progress, the *owner* Boolean value is true, and the associated reference indicates the thread that is in charge of resizing. We use the owner as a mutual exclusion flag between the *resize()* method and any of the *add()* methods, so that while resizing, there will be no successful updates, and while updating, there will be no successful resizes.

```
1  public void resize() {
2      int oldCapacity = table.length;
3      int newCapacity = 2 * oldCapacity;
4      Thread me = Thread.currentThread();
5      if (owner.compareAndSet(null, me, false, true)) {
6          try {
7              if (table.length != oldCapacity) { // someone else
                  resized first
8                  System.out.println("Someone_else_resized_first");
9                  return;
10             }
11             quiesce();
12             List<T>[] oldTable = table;
13             table = (List<T>[]) new List[newCapacity];
14             for (int i = 0; i < newCapacity; i++)
15                 table[i] = new ArrayList<T>();
16             locks = new ReentrantLock[newCapacity];
17             for (int j = 0; j < locks.length; j++) {
18                 locks[j] = new ReentrantLock();
19             }
20             initializeFrom(oldTable);
21         } finally {
22             owner.set(null, false); // restore prior state
23         }
24     }
25 }
```

The *acquire()* and the *resize()* methods guarantee mutually exclusive access via the flag principle using the mark field of the *owner* flag and the tables locks array, *acquire()* first acquires its locks and then reads the mark field, while *resize()* first sets mark and then reads the locks during the *quiesce()* call. This ordering ensures that any thread that acquires the locks after *quiesce()* has completed will see that the set is in the processes of being resized, and will back off until the resizing is complete.

```

1  protected void quiesce() {
2      for (ReentrantLock lock : locks) {
3          while (lock.isLocked()) {} // spin
4      }
5  }

1  public void acquire(T x) {
2      boolean[] mark = {true};
3      Thread me = Thread.currentThread();
4      Thread who;
5      while (true) {
6          do { // wait until not resizing
7              who = owner.get(mark);
8          } while (mark[0] && who != me);
9          ReentrantLock[] oldLocks = this.locks;
10         int myBucket = Math.abs(x.hashCode() % oldLocks.length);
11         ReentrantLock oldLock = oldLocks[myBucket];
12         oldLock.lock(); // acquire lock
13         who = owner.get(mark);
14         if ((!mark[0] || who == me) && this.locks == oldLocks) {
15             // recheck
16             return;
17         } else { // unlock & try again
18             oldLock.unlock();
19         }
20     }

```

9.6.3 Experiment Description

The test creates 8 threads that perform *enq()* and *deq()* operations in sequential and parallel. The dequeue values are revised and if a value is missing a fail assertion is thrown.

9.6.4 Observations and Interpretations

When executing the test the test fails on the parallel part due to a duplicate value when executing the parallel both test.

```
1 Assertion error: duplicate pop Exception in thread "Thread-13"  
  junit.framework.AssertionFailedError: DeqThread: duplicate  
  pop
```

9.6.5 Proposed changes to fix the problem

By adding terminal output through `System.out` in the `remove()` method or Debugging the file the execution works, thus becoming a problem of observation, since the fail assertion only happens when its not observed, this happens with the TCuckooHash and CuckooHas tests.

```
1 compile-test :  
2 .sequential add, contains, and remove  
3 .parallel add  
4 .parallel remove  
5 .parallel both  
6  
7 Time: 0.081  
8  
9 OK (4 tests)
```

9.6.6 Proposed solution

It is not clear why the observation would affect the execution, thus it is possible that the issue is a propagation of the values in the memory which are forced to be executed constantly through the debug implementation. Also worth mentioning that several hashing test suffer from this same situation and adding volatile didn't help as in other tests.

9.7 RefinableCuckooHashSetTest Test

9.7.1 Particular Case

The particular problem we are trying to solve here is how to implement a concurrent open-addressed hash (specifically a hash set).

We ought to remember that an open-addressed hash table is one in which each table entry holds one and only one item. In previous examples we had seen that an entry in the table can handle multiple elements which are stored in a linked list within the bucket.

The other thing we have to consider is that in this exercise we are willing to have a concurrent hash. Previously we saw a open-addressed hash that allowed coarse granularity. Here we want to go beyond that.

9.7.2 Solution

The solution to this problem is given by what is called the *Refinable Concurrent Cuckoo Hash*. The difference between this implementation and previously presented *Cuckoo Hashes* is that this one is a fine-grained approach where a lock is always responsible for protecting a fixed number of elements in the hash.

The interesting bits of this implementation are the *acquire()*, *release()* and *resize()* methods. Let us take a quick look at those:

```
1  /**
2   * Synchronize before adding, removing, or testing for item
3   * @param x item involved
4   */
5  public void acquire(T x) {
6      boolean[] mark = {true};
7      Thread me = Thread.currentThread();
8      Thread who;
9      while (true) {
10         do { // wait until not resizing
11             who = owner.get(mark);
12         } while (mark[0] && who != me);
13         ReentrantLock[][] oldLocks = this.locks;
14         ReentrantLock oldLock0 = oldLocks[0][hash0(x) % oldLocks[0].length];
15         ReentrantLock oldLock1 = oldLocks[1][hash1(x) % oldLocks[1].length];
16         oldLock0.lock(); // acquire locks
17         oldLock1.lock();
```

```

18         who = owner.get(mark);
19         if ((!mark[0] || who == me) && this.locks == oldLocks) { // recheck
20             return;
21         } else { // unlock & try again
22             oldLock0.unlock();
23             oldLock1.unlock();
24         }
25     }
26 }

```

The *acquire()* method first spins till it finds that no one is doing a resizing. This means that resizing and adding are exclusive operations. After this step, the thread acquires the two locks that protect the intended slot in the table. If after this point, the lock array is still the same we saw before taking the lock, then we are free to continue. Otherwise we attempt to take the locks again.

The *release()* method simply unlocks the slot:

```

1  /**
2   * synchronize after adding, removing, or testing for item
3   * @param x item involved
4   */
5  public final void release(T x) {
6      Lock lock0 = locks[0][hash0(x) % locks[0].length];
7      Lock lock1 = locks[1][hash1(x) % locks[1].length];
8      lock0.unlock();
9      lock1.unlock();
10 }

```

The *resize()* method is essentially the same as the one showed in the previous exercise.

9.7.3 Experiment Description

As for the tests provided for this experiment. They are the same discussed in previous exercises in this chapter. Nothing new has been added, so let us go directly to the results.

9.7.4 Sample Results

9.7.5 Interpretation

9.7.6 Proposed solution

9.8 StripedCuckooHashSetTest

9.8.1 Particular Case (problem)

The problem we try to solve is that of concurrent hashing, when we use the hash to implement a set data structure.

9.8.2 Solution

This solution is a refinement of the cuckoo technique we mentioned for other solutions like *CoarseCuckooHashSetTest*; providing a matrix of reentrant locks, where each cell protects a cell on the hash table. The following constructor shows how the locks are initialized:

```
1  public StripedCuckooHashSet(int capacity) {
2      super(capacity);
3      lock = new ReentrantLock[2][capacity];
4      for (int i = 0; i < 2; i++) {
5          for (int j = 0; j < capacity; j++) {
6              lock[i][j] = new ReentrantLock();
7          }
8      }
9  }
```

This solution extends the solution named *PhasedCuckooHashSet*. Part of its main differences are the *acquire* and *release* methods; which follow certain order on the locks associated to items, in order to avoid deadlocks:

```
1  public final void acquire(T x) {
2      Lock lock0 = lock[0][hash0(x) % lock[0].length];
3      Lock lock1 = lock[1][hash1(x) % lock[1].length];
4      lock0.lock();
5      lock1.lock();
6  }
7
8  public final void release(T x) {
9      Lock lock0 = lock[0][hash0(x) % lock[0].length];
10     Lock lock1 = lock[1][hash1(x) % lock[1].length];
11     lock0.unlock();
12     lock1.unlock();
13 }
```

9.8.3 Experiment Description

The test is in essence the same as that described for *CoarseCuckooHashSetTest*, with some small differences: among them, the most relevant is the remover thread class used in *testParallelBoth*, which does not fail anymore if the element it tried to remove was not present; it only increments a counter of missed deletions. This minor change prevents the test from failing like the others.

9.8.4 Observations and Interpretations

The test runs fine in 2 and 24 cores, without any assertion error; below a sample execution:

```
[oraadm@gdlaa008 orig]$ junit hash.StripedCuckooHashSetTest
.sequential add, contains, and remove
.parallel both
.parallel add
.parallel remove

Time: 0.017

OK (4 tests)
```

Chapter 10

Skiplists and Balanced Search

10.1 Double Ended Queue

10.1.1 Particular Case

In this experiment each thread keeps a pool of tasks waiting to be executed in the form of a double-ended queue (DEQueue), providing *pushBottom()*, *popBottom()*, and *popTop()* methods.

10.1.2 Solution

According to the theory when a thread creates a new task, it calls *pushBottom()* to push that task onto its *DEQueue*. When a thread needs a task to work on, it calls *popBottom()* to remove a task from its own *DEQueue*. If the thread discovers its queue is empty, then it becomes a thief, it chooses a victim thread at random, and calls that threads DEQueues *popTop()* method to steal a task for itself.

```
1  public Runnable popTop() {
2      int [] stamp = new int [1];
3      int oldTop = top.get(stamp), newTop = oldTop + 1;
4      int oldStamp = stamp[0], newStamp = oldStamp + 1;
5      if (bottom <= oldTop) // empty
6          return null;
7      Runnable r= tasks[oldTop];
8      if (top.compareAndSet(oldTop, newTop, oldStamp, newStamp))
9          return r;
10     return null;
11 }
```

10.1.3 Experiment Description

The test creates 16 threads, eight that need to be coordinate in order to *pushBottom()* and *popBottom()* an array of values. All threads have to cooperate to add and remove elements from the queue and eight *Dummy()* threads will be stealing execution from the running ones. After each thread execution, if everything works according to the test the map will have all values set to true. If that is not the case, a duplicate or missing fail will be raised.

10.1.4 Observations and Interpretations

The tests executed as expected and no errors where found.

```
1 Tests run: 2, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   1.162 sec
2
3 ----- Standard Output -----
4 sequential pushBottom and popBottom
5 concurrent pushBottom and popBottom
6 -----
7 test-single:
8 BUILD SUCCESSFUL (total time: 2 seconds)
```

10.2 Parallel Matrix Multiplication

10.2.1 Particular Case

In this experiment, the particular case we are trying to study is matrix parallel multiplication by devising methods to add and multiply matrices.

10.2.2 Solution

To implement a parallel multiplication we create a Matrix class that provides *set()* and *get()* methods to access matrix elements, along with a constant-time *split()* method that splits an n -by- n matrix into four $(n/2)$ -by- $(n/2)$ submatrices.

```
1  double get(int row, int col) {
2      return data[row+rowDisplace][col+colDisplace];
3  }
4
5  void set(int row, int col, double value) {
6      data[row+rowDisplace][col+colDisplace] = value;
7  }
8
9  Matrix [][] split() {
10     Matrix [][] result = new Matrix[2][2];
11     int newDim = dim / 2;
12     result[0][0] = new Matrix(data, rowDisplace, colDisplace,
13                               newDim);
14     result[0][1] = new Matrix(data, rowDisplace, colDisplace +
15                               newDim, newDim);
16     result[1][0] = new Matrix(data, rowDisplace + newDim,
17                               colDisplace, newDim);
18     result[1][1] = new Matrix(data, rowDisplace + newDim,
19                               colDisplace + newDim, newDim);
20     return result;
21 }
```

For simplicity, we consider matrices whose dimension n is a power of 2. Any such matrix can be decomposed into four submatrices, thus their sums can be done in parallel.

10.2.3 Experiment Description

The test creates two matrices with a dimension of four which will be multiplied in parallel, the resulting matrix will be verified and a fail assertion will be thrown if the result is incorrect.

10.2.4 Observations and Interpretations

When compiling the code the Matrix class had to be changed from private to public and the offset and dimension of the Matrix must be added to the test class as shown below.

```
1 Matrix aa = new Matrix(a, 0, 0, 4);
2 Matrix bb = new Matrix(b, 0, 0, 4);
```

After making this changes the codes compiles and everything works fine.

```
1 Tests run: 1, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
  0.369 sec
2
3 ----- Standard Output -----
4 run
5 -----
6 test-single:
7 BUILD SUCCESSFUL (total time: 0 seconds)
```

10.2.5 Proposed changes to fix the problem

The MatrixTask class needs to have the visibility of the inner Matrix class change to public in order for the test to view it.

```
1 public static class Matrix {
2     int dim;
3     double[][] data;
4     int rowDisplace;
5     int colDisplace;
6     -----
7 }
```

10.3 Unbounded Double Ended Queue

10.3.1 Particular Case

In this experiment we implement an unbounded double queue based on the previous double queue implemented.

10.3.2 Solution

According to the theory the *UnboundedDeque* class dynamically resizes itself as needed. We implement the *UnboundedDeque* in a cyclic array, with top and bottom fields as in the *BoundedDeque* (except indexed modulo the array's capacity). As before, if bottom is less than or equal to top, the *UnboundedDeque* is empty. The difference resides in the *CircularArray*() class which provides *get()* and *put()* methods that add and remove tasks, and a *resize()* method that allocates a new circular array and copies the old array's contents into the new array.

```
1  class CircularArray {
2      private int logCapacity;
3      private Runnable[] currentTasks;
4      CircularArray(int logCapacity) {
5          this.logCapacity = logCapacity;
6          currentTasks = new Runnable[1 << logCapacity];
7      }
8      int capacity() {
9          return 1 << this.logCapacity;
10     }
11     Runnable get(int i) {
12         return this.currentTasks[i % capacity()];
13     }
14     void put(int i, Runnable task) {
15         this.currentTasks[i % capacity()] = task;
16     }
17     CircularArray resize(int bottom, int top) {
18         CircularArray newTasks =
19             new CircularArray(this.logCapacity+1);
20         for (int i = top; i < bottom; i++) {
21             newTasks.put(i, this.get(i));
22         }
23         return newTasks;
24     }
25 }
```


10.3.3 Experiment Description

The test creates 16 threads, eight that need to be coordinate in order to *pushBottom()* and *popBottom()* an array of values. All threads have to cooperate to add and remove elements from the queue and eight *Dummy()* threads will be stealing execution from the running ones. After each thread execution, if everything works according to the test the map will have all values set to true. If that is not the case, a duplicate or missing fail will be raised.

10.3.4 Observations and Interpretations

The test needs some changes in order to compile, once this changes are implemented the test executes as expected.

```
1 Tests run: 2, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   1.4 sec
2
3 ----- Standard Output -----
4 sequential pushBottom and popBottom
5 concurrent pushBottom and popBottom
6 -----
7 test-single:
8 BUILD SUCCESSFUL (total time: 1 second)
```

10.3.5 Proposed changes to fix the problem

The UDEQueueTest class needs to have replaced the calls to non defined reset by new instance in order to compile.

```
1 //instance.reset();
2 instance = new UDEQueue(THREADS);
```

10.4 ArraySumTest

10.4.1 Particular Case (problem)

The problem we want to solve is that of summing up all the elements of an array; but of course, we want to do it in parallel.

10.4.2 Solution

The solution uses the divide and conquer approach, creating a tree of tasks (*Future* from java concurrency library); which represent the processing we do while splitting the arrays in two halves and summing up each half recursively. Base case are arrays of a single element. Below the heart of the solution, the recursive function which create new tasks for subarrays:

```
1      public Integer call() throws InterruptedException , ExecutionException {
2          if (size == 1) {
3              return a[start];
4          } else {
5              int lhsSize = size / 2;
6              int rhsStart = lhsSize + 1;
7              int rhsSize = size - lhsSize;
8              Future<Integer> lhs = exec.submit(new SumTask(a, start , lhsSize));
9              Future<Integer> rhs = exec.submit(new SumTask(a, rhsStart , rhsSize));
10             return rhs.get() + lhs.get();
11         }
12     }
```

10.4.3 Experiment Description

The test merely consists in trying the algorithm with a couple of arrays of consecutive positive integers, of respective sizes 3 and 4. The test validates that the expected sum is 6 and 10 respectively.

10.4.4 Observations and Interpretations

The test presents failures like the one below, where the array of 3 elements gave as sum 7 instead of 6:

```

[oraadm@gdlaa008 orig]$ junit steal.ArraySumTest
.F
Time: 0.009
There was 1 failure:
1) testRun(steal.ArraySumTest)junit.framework.AssertionFailedError:
    expected:<6> but was:<7>
at steal.ArraySumTest.testRun(ArraySumTest.java:31)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at
sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)

FAILURES!!!
Tests run: 1, Failures: 1, Errors: 0

```

The problem lied on the program *ArraySum.java*, on the assignment of variable *rhsStart*, which represents the starting point of the right subarray. While the original program had *lhsSize + 1* as assignment, the proper expression was *lhsSize + start* (so we take into account both the relative start of the initial subarray we got, as well as the length of the left sub-subarray). After this fix, the test worked fine on both 2 and 24 cores machines.

10.5 FibTaskTest

10.5.1 Particular Case (problem)

The problem is calculating the sum of the first n Fibonacci numbers, in parallel.

10.5.2 Solution

The solution merely follows the recursive definition of the Fibonacci sequence, whose base cases are 1 (for first and second elements of the sequence). Each recursive call of the function creates a new task (*Future* from java concurrency library). Below the code of the main function:

```
1  public Integer call() {
2      try {
3          if (arg > 2) {
4              Future<Integer> left = exec.submit(new FibTask(arg - 1));
5              Future<Integer> right = exec.submit(new FibTask(arg - 2));
6              return left.get() + right.get();
7          } else {
8              return 1;
9          }
10     } catch (Exception ex) {
11         ex.printStackTrace();
12         return 1;
13     }
14 }
```

Note that *Future* class represents asynchronous calculations, and that the *get* method is blocking (waits for calculation to be complete); hence we would effectively create a tree of tasks until we reach the base cases and start evaluating bottom up.

10.5.3 Experiment Description

The experiment merely request the calculation of the sum up to a known number (16), and asserts the expected result (987).

10.5.4 Observations and Interpretations

The test works as expected in both 2 and 24 cores machines. Below a sample output:

```
[oraadm@gdlaa008 orig]$ junit steal.FibTaskTest  
.run
```

```
Time: 0.128
```

```
OK (1 test)
```

10.6 MMThreadTest

10.6.1 Particular Case (problem)

The problem is that of matrix multiplication, but done in parallel.

10.6.2 Solution

The solution creates one thread per cell on the result matrix, and for each one it creates a new worker thread; which in turn calculates the cell (i, j) value by computing the dot product of row i of left matrix by column j of right matrix (according to the definition of the matrix product). Below the most relevant sections of the code:

```
1  public MMThread(double [][] a, double [][] b) {
2      n = a.length;
3      this.a = a;
4      this.b = b;
5      this.c = new double[n][n];
6  }
7
8  void multiply() {
9      Worker [][] worker = new Worker[n][n];
10     // create one thread per matrix entry
11     for (int row = 0; row < n; row++) {
12         for (int col = 0; col < n; col++) {
13             worker[row][col] = new Worker(row, col);
14         }
15     }
16     ...
17 }
18
19 class Worker extends Thread {
20     int row, col;
21     Worker(int row, int col) {
22         this.row = row; this.col = col;
23     }
24     public void run() {
25         double dotProduct = 0.0;
26         for (int i = 0; i < n; i++) {
27             dotProduct += a[row][i] * b[i][col];
28         }
29         c[row][col] = dotProduct;
```

```
30     }  
31 }
```

10.6.3 Experiment Description

The test consists of multiplying two identity matrices of dimensions 3×3 , and compare the result with one of the operands (it should be identity matrix again).

10.6.4 Observations and Interpretations

The test runs fine on 2 and 24 core machines, below a sample output:

```
[oraadm@gdlaa008 orig]$ junit steal.MMThreadTest  
.run
```

```
Time: 0.004
```

```
OK (1 test)
```

Chapter 11

Priority Queues

11.1 BitReversedCounterTest

11.1.1 Particular Case (problem)

The problem is actually a subproblem of the *FineGrainedHeap* program, which has a *next* variable indicating next slot to use on insert. The problem is to ensure that two consecutive paths from root to leaves, where insertion takes place, share no common nodes other than the root (this is useful to ensure that each insertion thread locks a disjoint set of nodes, reducing contention on multi-threaded scenarios).

11.1.2 Solution

The solution is a bit cryptic, but small enough to fit here. It can be proved that the *reverseIncrement* method below, produces sequence of numbers with the above property. Given time constraints, we did not prove such claim (specially cause this test itself, did not exhibit concurrency):

```
1  public int reverseIncrement() {
2      if (counter++ == 0) {
3          reverse = highBit = 1;
4          return reverse;
5      }
6      int bit = highBit >> 1;
7      while (bit != 0) {
8          reverse ^= bit;
9          if ((reverse & bit) != 0) break;
10         bit >>= 1;
11     }
12     if (bit == 0)
13         reverse = highBit <<= 1;
14     return reverse;
15 }
```

11.1.3 Experiment Description

The test does not really make any assertions, and simply call the *reverseIncrement* and *reverseDecrement* methods printing their results.

11.1.4 Observations and Interpretations

The test did not show any failures or hanging, given that is totally sequential and does not really validate anything. Below a sample execution output:

```
[oraadm@gdlaa008 orig]$ junit priority.BitReversedCounterTest
.increment
inc:      0      1
inc:      1      2
inc:      2      3
inc:      3      4
inc:      4      6
inc:      5      5
inc:      6      7
inc:      7      8
inc:      8     12
inc:      9     10
inc:     10     14
inc:     11      9
inc:     12     13
inc:     13     11
inc:     14     15
inc:     15     16
inc:     16     24
inc:     17     20
inc:     18     28
inc:     19     18
inc:     20     26
inc:     21     22
inc:     22     30
inc:     23     17
inc:     24     25
inc:     25     21
inc:     26     29
inc:     27     19
inc:     28     27
inc:     29     23
inc:     30     31
inc:     31     32
inc:     32     48
inc:     33     40
inc:     34     56
inc:     35     36
inc:     36     52
inc:     37     44
inc:     38     60
inc:     39     34
inc:     40     50
inc:     41     42
inc:     42     58
inc:     43     38
inc:     44     54
inc:     45     46
inc:     46     62
inc:     47     33
```

inc:	48	49
inc:	49	41
inc:	50	57
inc:	51	37
inc:	52	53
inc:	53	45
inc:	54	61
inc:	55	35
inc:	56	51
inc:	57	43
inc:	58	59
inc:	59	39
inc:	60	55
inc:	61	47
inc:	62	63
inc:	63	64
decrement		
dec:	63	
dec:	47	
dec:	55	
dec:	39	
dec:	59	
dec:	43	
dec:	51	
dec:	35	
dec:	61	
dec:	45	
dec:	53	
dec:	37	
dec:	57	
dec:	41	
dec:	49	
dec:	33	
dec:	62	
dec:	46	
dec:	54	
dec:	38	
dec:	58	
dec:	42	
dec:	50	
dec:	34	
dec:	60	
dec:	44	
dec:	52	
dec:	36	
dec:	56	
dec:	40	
dec:	48	
dec:	32	

Time: 0.022

OK (1 test)

11.2 SimpleTreeTest Test

11.2.1 Particular Case

The problem we want to solve now is how we can implement a bounded priority queue that can be accessed concurrently by multiple threads.

Remember that priority queue is said to be *bounded* when the range of the priorities are taken from the range $0, \dots, m - 1$.

11.2.2 Solution

One possible solution to this problem is given by the *Tree-Based Bounded Priority Queue* data structure. This structure is essentially a binary tree where the leaf nodes have a bin holding items of priority i . Internal nodes maintain a shared counter that indicates the number of elements stored to the left of a given node.

The interesting methods in this data structure are the *add()* and *removeMin()*, so let us talk about them.

The *add()* method first takes the bin associated with the given priority. It puts the item into that bin and then starts increasing the counter of the parent nodes all the way to the root. Note that it only increases the counter if the current node happens to be a left child of the parent.

```
1  /**
2   *  add item to priority queue
3   *  @param item new item
4   *  @param priority item 's priority
5   */
6  public void add(T item, int priority) {
7      TreeNode node = leaves.get(priority);
8      node.bin.put(item);
9      while(node != root) {
10         TreeNode parent = node.parent;
11         if (node == parent.left) { // increment if ascending from left
12             parent.counter.getAndIncrement();
13         }
14         node = parent;
15     }
16 }
```

The *removeMin()* method has to retrieve the item with the lowest priority. In order to do that, the method starts from the root and traverses the tree. If the current counter of the inner node is 0, then it knows that the min must be on the left. Otherwise, it goes through the right child.

```
1  public T removeMin() {
2      TreeNode node = root;
3      while (!node.isLeaf()) {
4          if (node.counter.getAndDecrement() > 0 ) {
5              node = node.left;
6          } else {
7              node = node.right;
8          }
9      }
10     return node.bin.get(); // if null pqueue is empty
11 }
```

11.2.3 Experiment Description

Three test cases were provided to exercise this code:

- testAdd. This test adds elements with random priority to the priority queue one by one. At the end, it extracts each element and it must be the case that the previously removed element had a lower priority than the current element.
- testParallelAdd. This test spawns 8 threads. Each of the test adds 8 elements to the queue concurrently. At the end, it performs the same checking as in the previous test case. Notice that each thread is in charge of adding elements from a fixed range in ascending order.
- testParallelBoth. Does the same as the previous test case but also removes elements in parallel. Each remover thread will perform the same checking. Since adders add elements in ascending order, it is guaranteed that any removal from any thread will retrieve an element with a lower priority than the next one.

- 11.2.4 Sample Results
- 11.2.5 Interpretation
- 11.2.6 Proposed solution

11.3 Sequential Heap

11.3.1 Particular Case

In this experiment we implement a linearisable priority queue that supports priorities from an unbounded range. It uses fine-grained locking for synchronization.

11.3.2 Solution

According to the theory a sequential heap implementation represent a binary heap by using an array of nodes, where the trees root is array entry 1, and the right and left children of array entry i are entries $2i$ and $(2i) + 1$, respectively. Each node has an item and a score field. To add an item, the *add()* method sets child to the index of the first empty array slot.

```
1  public void add(T item, int priority) {
2      int child = next++;
3      heap[child].init(item, priority);
4      while (child > ROOT) {
5          int parent = child / 2;
6          int oldChild = child;
7          if (heap[child].priority < heap[parent].priority) {
8              swap(child, parent);
9              child = parent;
10         } else {
11             return;
12         }
13     }
14 }
```

To remove and return the highest-priority item, the *removeMin()* method records the roots item, which is the highest-priority item in the tree.

```
1  public T removeMin() {
2      int bottom = --next;
3      T item = heap[ROOT].item;
4      swap(ROOT, bottom);
5      if (bottom == ROOT) {
6          return item;
7      }
8      int child = 0;
9      int parent = ROOT;
10     while (parent < heap.length / 2) {
```

```

11     int left = parent * 2; int right = (parent * 2) + 1;
12     if (left >= next) {
13         break;
14     } else if (right >= next || heap[left].priority < heap[
15         right].priority) {
16         child = left;
17     } else {
18         child = right;
19     }
20     // If child higher priority than parent swap then else
21     stop
22     if (heap[child].priority < heap[parent].priority) {
23         swap(parent, child); // Swap item, key, and tag of heap[
24         i] and heap[child]
25         parent = child;
26     } else {
27         break;
28     }
29 }
30 return item;
31 }

```

11.3.3 Experiment Description

The test adds 512 elements to the heap, does a check of the elements and removes the values with min priority. We verify the priorities of the removed elements and assert a fail if we have non ascending priorities.

11.3.4 Observations and Interpretations

The test execute as expected and no problems are thrown.

```

1 Tests run: 1, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
  0.413 sec
2
3 ----- Standard Output -----
4 testSequential
5 OK.
6 -----
7 test-single:
8 BUILD SUCCESSFUL (total time: 1 second)

```


11.4 Skiplist-Based Unbounded Priority Queue

11.4.1 Particular Case

In this experiment we implement a skip-list priority queue which in contrast to the *FineGrainedHeap* priority queue algorithm, that the underlying heap structure requires complex, coordinated rebalancing. We requires no rebalancing.

11.4.2 Solution

According to the theory the *PrioritySkipList* class sorts items by priority instead of by hash value, ensuring that high-priority items (the ones we want to remove first) appear at the front of the list. Removing the item with highest priority is done lazily. A node is logically removed by marking it as removed, and is later physically removed by unlinking it from the list.

```
1  public T removeMin() {  
2      Node<T> node = skiplist.findAndMarkMin();  
3  
4      if (node != null) {  
5          skiplist.remove(node);  
6          return node.item;  
7      } else {  
8          return null;  
9      }  
10 }
```

The *add()* and *remove()* calls take skiplist nodes instead of items as arguments and results. These methods are straightforward adaptations of the corresponding *LockFreeSkipList* methods. This class nodes differ from *LockFreeSkipList* nodes in two fields, an integer score field, and an *AtomicBoolean* marked field used for logical deletion from the priority queue (not from the skiplist).

```
1  boolean add(Node node) {  
2      int bottomLevel = 0;  
3      Node<T>[] preds = (Node<T>[]) new Node[MAXLEVEL + 1];  
4      Node<T>[] succs = (Node<T>[]) new Node[MAXLEVEL + 1];  
5      while (true) {  
6          boolean found = find(node, preds, succs);  
7          if (found) { // if found it's not marked  
8              return false;  
          }  
      }
```

```

9      } else {
10         for (int level = bottomLevel; level <= node.topLevel;
              level++) {
11             Node<T> succ = succs[level];
12             node.next[level].set(succ, false);
13         }
14         // try to splice in new node in bottomLevel going up
15         Node<T> pred = preds[bottomLevel];
16         Node<T> succ = succs[bottomLevel];
17         node.next[bottomLevel].set(succ, false);
18         if (!pred.next[bottomLevel].compareAndSet(succ, node,
              false, false)) { // lin point
19             continue; // retry from start
20         }
21         // splice in remaining levels going up
22         for (int level = bottomLevel+1; level <= node.topLevel;
              level++) {
23             while (true) {
24                 pred = preds[level];
25                 succ = succs[level];
26                 if (pred.next[level].compareAndSet(succ, node, false,
              false))
27                     break;
28                 find(node, preds, succs); // find new preds and
              succs
29             }
30         }
31         return true;
32     }
33 }
34 }
35
36 boolean remove(Node<T> node) {
37     int bottomLevel = 0;
38     Node<T>[] preds = (Node<T>[]) new Node[MAXLEVEL + 1];
39     Node<T>[] succs = (Node<T>[]) new Node[MAXLEVEL + 1];
40     Node<T> succ;
41     while (true) {
42         boolean found = find(node, preds, succs);
43         if (!found) {
44             return false;
45         } else {
46             // proceed to mark all levels
47             // some levels could stil be unthreaded by concurrent
              add() while being marked

```

```

48      // other find()s could be modifying node's pointers
        concurrently
49      for (int level = node.topLevel; level >= bottomLevel+1;
        level--) {
50          boolean[] marked = {false};
51          succ = node.next[level].get(marked);
52          while (!marked[0]) { // until I succeed in marking
53              node.next[level].attemptMark(succ, true);
54              succ = node.next[level].get(marked);
55          }
56      }
57      // proceed to remove from bottom level
58      boolean[] marked = {false};
59      succ = node.next[bottomLevel].get(marked);
60      while (true) { // until someone succeeded in marking
61          boolean iMarkedIt = node.next[bottomLevel].
            compareAndSet(succ, succ, false, true);
62          succ = succs[bottomLevel].next[bottomLevel].get(marked
            );
63          if (iMarkedIt) {
64              // run find to remove links of the logically removed
                node
65              find(node, preds, succs);
66              return true;
67          } else if (marked[0]) return false; // someone else
            removed node
68          // else only succ changed so repeat
69      }
70  }
71 }
72 }

```

11.4.3 Experiment Description

The test creates 8 threads, eight that need to be coordinate in order to *add()* and *removeMin()* an array of values. All threads have to cooperate to add and remove elements from the queue and a validation of the priorities is done to validate non-ascending priorities. If that is not the case, a fail will be raised.

11.4.4 Observations and Interpretations

The test needs some changes in order to execute and avoid starvation.

```
1  _____ Standard Output _____
2  add
3  OK.
4  testParallelAdd
5  OK.
6  testParallelBoth
```

11.4.5 Proposed changes to fix the problem

The *findAndMarkMin()* method must choose the corresponding element that will be removed in order to mark it and later on the *remove()* method we verify the marked node and remove it, but the atomicity of this operations are not guarantee, thus when a node is marked it can be marked by other thread causing the remove process to hand, since the internal logic tries to remove the node and has no escape sequence when the mark is affected.

```
1  public T removeMin() {
2      synchronized(this) {
3          Node<T> node = skiplist.findAndMarkMin();
4
5          if (node != null) {
6              skiplist.remove(node);
7              return node.item;
8          } else{
9              return null;
10         }
11     }
12 }
```

By doing so, the execution finishes as expected and no errors are found.

```
1  Tests run: 3, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   0.358 sec
2
3  _____ Standard Output _____
4  add
5  OK.
6  testParallelAdd
7  OK.
8  testParallelBoth
```

```
9 | OK.  
10 | _____  
11 | test-single:  
12 | BUILD SUCCESSFUL (total time: 0 seconds)
```

11.5 FineGrainedHeapTest Test

11.5.1 Particular Case

Now we look at the problem of implementing an Unbounded Priority Queue in a Fine-Grained fashion.

Again, remember that Unbounded in this context means that the range is unbounded, for example the set of 32-bits integers or floating point values.

11.5.2 Solution

The solution proposed in this exercise is the Fine-Grained heap. The particular characteristic of this implementation is that percolation happens in a sequence of discrete atomic steps that can be interleaved with other such steps. The same happens with the *removeMin()* method.

The *add()* method first locks the heap and then takes the lock to the next available child and puts the new value there. After that, it releases the heap lock. After this point, it must percolate this new node. To do this, it has to swap child and parent after locking them. This only happens if the parent is available and the child has a higher priority. Otherwise, the node is already where it belongs. At the end of this step, child and parent are unlocked. Then it goes up one node in the heap, and does the same.

```
1  /**
2   * Add item to heap.
3   * @param item Uninterpreted item.
4   * @param priority item priority
5   */
6  public void add(T item, int priority) {
7      heapLock.lock();
8      int child = next++;
9      heap[child].lock();
10     heapLock.unlock();
11     heap[child].init(item, priority);
12     heap[child].unlock();
13     while (child > ROOT) {
14         int parent = child / 2;
15         heap[parent].lock();
16         heap[child].lock();
17         int oldChild = child;
18         try {
```

```

19         if (heap[parent].tag == Status.AVAILABLE && heap[child].amOwner()) {
20             if (heap[child].score < heap[parent].score) {
21                 swap(child, parent);
22                 child = parent;
23             } else {
24                 heap[child].tag = Status.AVAILABLE;
25                 heap[child].owner = NO_ONE;
26                 return;
27             }
28         } else if (!heap[child].amOwner()) {
29             child = parent;
30         }
31     } finally {
32         heap[oldChild].unlock();
33         heap[parent].unlock();
34     }
35 }
36 if (child == ROOT) {
37     heap[ROOT].lock();
38     if (heap[ROOT].amOwner()) {
39         heap[ROOT].tag = Status.AVAILABLE;
40         heap[child].owner = NO_ONE;
41     }
42     heap[ROOT].unlock();
43 }
44 }

```

The *removeMin()* method first locks the heap, the bottom and the root nodes. If the root is empty, then we are done. Otherwise, it swaps root and bottom and unlocks bottom. After that, it has to percolate the new root element. It follows a similar approach as the one mentioned in the *add()* method. One difference is that this percolation goes from top to bottom.

```

1  /**
2   * Returns and removes lowest-priority item in heap.
3   * @return lowest-priority item.
4   */
5  public T removeMin() {
6      heapLock.lock();
7      int bottom = --next;
8      heap[bottom].lock();
9      heap[ROOT].lock();
10     heapLock.unlock();

```

```

11     if (heap[ROOT].tag == Status.EMPTY) {
12         heap[ROOT].unlock();
13         heap[bottom].lock();
14         return null;
15     }
16     T item = heap[ROOT].item;
17     heap[ROOT].tag = Status.EMPTY;
18     swap(bottom, ROOT);
19     heap[bottom].owner = NO_ONE;
20     heap[bottom].unlock();
21     if (heap[ROOT].tag == Status.EMPTY) {
22         heap[ROOT].unlock();
23         return item;
24     }
25     int child = 0;
26     int parent = ROOT;
27     while (parent < heap.length / 2) {
28         int left = parent * 2;
29         int right = (parent * 2) + 1;
30         heap[left].lock();
31         heap[right].lock();
32         if (heap[left].tag == Status.EMPTY) {
33             heap[right].unlock();
34             heap[left].unlock();
35             break;
36         } else if (heap[right].tag == Status.EMPTY || heap[left].score < heap[right].score) {
37             heap[right].unlock();
38             child = left;
39         } else {
40             heap[left].unlock();
41             child = right;
42         }
43         if (heap[child].score < heap[parent].score) {
44             swap(parent, child);
45             heap[parent].unlock();
46             parent = child;
47         } else {
48             heap[child].unlock();
49             break;
50         }
51     }
52     heap[parent].unlock();
53     return item;
54 }

```


11.5.3 Experiment Description

The same three test cases were provided in this experiment, so let us go directly to the results.

11.5.4 Sample Results

11.5.5 Interpretation

11.5.6 Proposed solution

11.6 SimpleLinearTest

11.6.1 Particular Case (problem)

The problem is to implement a bounded priority queue, which are a multi-set data structure; each set has an associated priority. The contract or interface for a priority queue is essentially two methods: *add(x,k)* which adds element *x* to the set of priority *k*, and *removeMin* which returns and removes the element with smaller (higher) priority from all the sets. This type of priority queues are called “bounded”, because the priority of elements is taken from a finite predefined set.

11.6.2 Solution

Given that the problem is to implement a bounded priority queue, an array shall suffice. So we create an array of sets as big as the range of values for the priority. The *add* and *removeMin* methods are straightforward, but they are built on the assumption that the auxiliary class *Bin* is thread safe (which is the case, as it is implemented with locks). Below the relevant snippets of code:

```
1 public class SimpleLinear<T> implements PQueue<T> {
2     int range;
3     Bin<T>[] pqueue;
4
5     ...
6
7     public void add(T item, int key) {
8         pqueue[key].put(item);
9     }
10
11     ...
12
13     public T removeMin() {
14         for (int i = 0; i < range; i++) {
15             T item = pqueue[i].get();
16             if (item != null) {
17                 return item;
18             }
19         }
20         return null;
21     }
22 }
```

```
21     }  
22 }
```

11.6.3 Experiment Description

The test program consists of the following test cases:

- *testAdd*: Serial test which adds 512 random numbers in a fixed range (these numbers become the priority), and later validates that they are removed in ascending order with *removeMin*.
- *testParallelAdd*: Same as *testAdd*, with the exception that the addition of the elements is done in parallel with 8 threads.
- *testParallelBoth*: Same as *testParallelAdd*, except that the removal of items is also performed in parallel with 8 threads.

11.6.4 Observations and Interpretations

The test runs fine in 2 and 24 core machines, below a sample output:

```
[oraadm@gdlaa008 orig]$ junit priority.SimpleLinearTest  
.add  
OK.  
...  
OK.  
OK.  
.testParallelBoth  
OK.  
.testParallelAdd  
OK.  
  
Time: 0.116  
  
OK (3 tests)
```

Chapter 12

Futures, Scheduling and Work Distribution

12.1 ArraySumTest

12.1.1 Particular Case (problem)

The problem we want to solve is that of summing up all the elements of an array; but of course, we want to do it in parallel (this is the same as the test program of same name from chapter 14).

12.1.2 Solution

The solution uses the divide and conquer approach, creating a tree of tasks (*Future* from java concurrency library); which represent the processing we do while splitting the arrays in two halves and summing up each half recursively. Base case are arrays of a single element. Below the heart of the solution, the recursive function which create new tasks for subarrays:

```
1      public Integer call() throws InterruptedException , ExecutionException {
2          if (size == 1) {
3              return a[start];
4          } else {
5              int lhsSize = size / 2;
6              int rhsStart = lhsSize + 1;
7              int rhsSize = size - lhsSize;
8              Future<Integer> lhs = exec.submit(new SumTask(a, start, lhsSize));
9              Future<Integer> rhs = exec.submit(new SumTask(a, rhsStart, rhsSize));
10             return rhs.get() + lhs.get();
11         }
12     }
```

Above code is exactly the same as the one from chapter 14; we are not sure whether this repetition was intentional, but putting anyway for the sake of completeness.

12.1.3 Experiment Description

The test merely consists in trying the algorithm with a couple of arrays of consecutive positive integers, of respective sizes 3 and 4. The test validates that the expected sum is 6 and 10 respectively (again, same as chapter 14).

12.1.4 Observations and Interpretations

The test presents failures like the one below, where the array of 3 elements gave as sum 7 instead of 6:

```
[oraadm@gdlaa008 orig]$ junit steal.ArraySumTest
.F
Time: 0.01
There was 1 failure:
1) testRun(steal.ArraySumTest)junit.framework.AssertionFailedError:
    expected:<6> but was:<7>
    at steal.ArraySumTest.testRun(ArraySumTest.java:31)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at
    sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
    at
    sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)

FAILURES!!!
Tests run: 1, Failures: 1, Errors: 0
```

The problem lied on the program *ArraySum.java*, on the assignment of variable *rhsStart*, which represents the starting point of the right subarray. While the original program had *lhsSize* + 1 as assignment, the proper expression was *lhsSize* + *start* (so we take into account both the relative start of the initial subarray we got, as well as the length of the left sub-subarray). After this fix, the test worked fine on both 2 and 24 cores machines.

Again, above fix and behavior were the same as homonym program from chapter 14.

12.2 FibTaskTest

12.2.1 Particular Case (problem)

The problem is calculating the sum of the first n Fibonacci numbers, in parallel.

Note: this exercise is exactly the same as the homonym from chapter 14; we believe the authors of the book made a mistake and ended up with this repetition. Still we are repeating here for the sake of completeness.

12.2.2 Solution

The solution merely follows the recursive definition of the Fibonacci sequence, whose base cases are 1 (for first and second elements of the sequence). Each recursive call of the function creates a new task (*Future* from java concurrency library). Below the code of the main function:

```
1  public Integer call() {
2      try {
3          if (arg > 2) {
4              Future<Integer> left = exec.submit(new FibTask(arg - 1));
5              Future<Integer> right = exec.submit(new FibTask(arg - 2));
6              return left.get() + right.get();
7          } else {
8              return 1;
9          }
10     } catch (Exception ex) {
11         ex.printStackTrace();
12         return 1;
13     }
14 }
```

Note that *Future* class represents asynchronous calculations, and that the *get* method is blocking (waits for calculation to be complete); hence we would effectively create a tree of tasks until we reach the base cases and start evaluating bottom up.

12.2.3 Experiment Description

The experiment merely request the calculation of the sum up to a known number (16), and asserts the expected result (987).

12.2.4 Observations and Interpretations

The test works as expected in both 2 and 24 cores machines. Below a sample output:

```
[oraadm@gdlaa008 orig]$ junit steal.FibTaskTest  
.run
```

```
Time: 0.176
```

```
OK (1 test)
```


12.3 MMThreadTest

12.3.1 Particular Case (problem)

The problem is that of matrix multiplication, but done in parallel.

Note: this exercise is exactly the same as the homonym from chapter 14; we believe the authors of the book made a mistake and ended up with this repetition. Still we are repeating here for the sake of completeness.

12.3.2 Solution

The solution creates one thread per cell on the result matrix, and for each one it creates a new worker thread; which in turn calculates the cell (i, j) value by computing the dot product of row i of left matrix by column j of right matrix (according to the definition of the matrix product). Below the most relevant sections of the code:

```
1  public MMThread(double [][] a, double [][] b) {
2      n = a.length;
3      this.a = a;
4      this.b = b;
5      this.c = new double[n][n];
6  }
7
8  void multiply() {
9      Worker [][] worker = new Worker[n][n];
10     // create one thread per matrix entry
11     for (int row = 0; row < n; row++) {
12         for (int col = 0; col < n; col++) {
13             worker[row][col] = new Worker(row, col);
14         }
15     }
16     ...
17 }
18
19 class Worker extends Thread {
20     int row, col;
21     Worker(int row, int col) {
22         this.row = row; this.col = col;
23     }
24     public void run() {
```

```

25         double dotProduct = 0.0;
26         for (int i = 0; i < n; i++) {
27             dotProduct += a[row][i] * b[i][col];
28         }
29         c[row][col] = dotProduct;
30     }
31 }

```

12.3.3 Experiment Description

The test consists of multiplying two identity matrices of dimensions 3×3 , and compare the result with one of the operands (it should be identity matrix again).

12.3.4 Observations and Interpretations

The test runs fine on 2 and 24 core machines, below a sample output:

```

[oraadm@gdlaa008 orig]$ junit steal.MMThreadTest
.run

Time: 0.004

OK (1 test)

```

Chapter 13

Barriers

13.1 DisBarrierTest

13.1.1 Particular Case (problem)

The problem we are trying to solve is that of synchronization on a multi-threaded environment; and the particular mechanism we want to use for that is a barrier.

13.1.2 Solution

The particular type of barrier proposed for this exercise is called a Dissemination Barrier, and it has a nice tree structure. Imagine that in order to implement the barrier, you need to perform “r” rounds; where each rounds means that all the threads involved communicated with other two threads (everyone sent a message to one predefined thread and waited to receive a message from another predefined thread). We use the communication word loosely here, as in a shared memory computer it will simply mean writing or reading from a shared variable.

Now, each round is essentially some sort of permutation; is like you align the n threads in one line, duplicate that line and do a shuffle on it, then you pair its elements with the original line. It turns out (can be proved), that if you do $\text{ceil}(\log_2(n))$ rounds and on each round k use the formula $(i + 2^k) \bmod n$ for thread communication (where i is the slot associated to each thread on an array), then all the threads can safely assumed that everybody has finished its call to *await* and they can move on. The implementation is compact enough to fit here, so here it goes:

```
1 public class DisBarrier implements Barrier {
2     int size;
3     int logSize;
4     Node[][] node;
5     ThreadLocal<Boolean> mySense;
6     ThreadLocal<Integer> myParity;
7
8     public DisBarrier(int capacity) {
9         size = capacity;
10        logSize = 0;
11        while (capacity != 1) {
12            this.logSize++;
```

```

13         capacity >>= 1;
14     }
15     node = new Node[logSize][size];
16     for (int r = 0; r < logSize; r++) {
17         for (int i = 0; i < size; i++) {
18             node[r][i] = new Node();
19         }
20     }
21     int distance = 1;
22     for (int r = 0; r < logSize; r++) {
23         for (int i = 0; i < size; i++) {
24             node[r][i].partner = node[r][(i + distance) % size];
25         }
26         distance *= 2;
27     }
28     this.mySense = new ThreadLocal<Boolean>() {
29         protected Boolean initialValue() { return true; };
30     };
31     this.myParity = new ThreadLocal<Integer>() {
32         protected Integer initialValue() { return 0; };
33     };
34 }
35
36 public void await() {
37     int parity = myParity.get();
38     boolean sense = mySense.get();
39     int i = ThreadID.get();
40     for (int r = 0; r < logSize; r++) {
41         node[r][i].partner.flag[parity] = sense;
42         while (node[r][i].flag[parity] != sense) {}
43     }
44     if (parity == 1) {
45         mySense.set(!sense);
46     }
47     myParity.set(1 - parity);
48 }
49
50 private static class Node {
51     boolean[] flag = {false, false}; // signal when done
52     Node partner; // partner node
53 }
54 }

```

13.1.3 Experiment Description

The test consists in having a shared array of 8 slots, one per thread to be created. Then the test has two phases, both using the barrier; the first phase is to wait all threads to write on their assigned slot on the array and the second stage is to allow all threads to validate that the rest actually finished their writing. These two phases are repeated 8 times.

13.1.4 Observations and Interpretations

The test hangs time to time, on the 24 cores machine (with 2 cores machine, issue was more difficult to reproduce); the hanging threads have an stack like this:

```
"Thread-5" prio=10 tid=0x00007f1bc0174800 nid=0x8414 runnable [0x00007f1b6bb59000]  
java.lang.Thread.State: RUNNABLE  
at barrier.DisBarrier.await(DisBarrier.java:59)  
at barrier.DisBarrierTest$TestThread.run(DisBarrierTest.java:79)
```

The stack above refers to the following wait cycle from the *await* method:

```
while (node[r][i].flag[parity] != sense) {}
```

The problem then implies that the thread in charge of writing to the flag of the hanging thread, never made the write ... well, that is the appearance but in reality what happened is that the flag variable was not declared as volatile, therefore there is no guarantee about when that new value will be propagated from caches to main memory, and from there to the hanging thread core's cache. Thus, even when the write was actually done the other thread could not see it:

As with other tests from chapter 2, we can solve this issue simply by declaring the flag as volatile:

```
volatile boolean[] flag = {false , false}
```

With above fix the test no longer hangs on both 2 and 24 core machines.

13.2 SenseBarrierTest

13.2.1 Particular Case (problem)

The problem we are trying to solve is that of synchronization on a multi-threaded environment; and the particular mechanism we want to use for that is a barrier (same as *DisBarrierTest*).

13.2.2 Solution

The solution called a “Sense Barrier” is well suited for reusing barrier structures (given that barriers construction, in general, is kind of expensive). The approach introduces the concept of a *sense*; which is basically a boolean flag used for both the whole barrier structure and for each thread. An implementation detail is that the barrier’s sense is a shared variable among threads (hence declared *volatile*) and the thread’s sense is implemented with a *ThreadLocal* variable.

The main idea behind this solution is that the barrier’s sense is initialized to the negation of the threads’ senses; and that a call to the *await* method will block all threads to wait for the barrier’s sense to be negated. Exception is the last thread to come, which is allowed to actually change the barrier’s sense hence unblocking the rest of the threads. Prior leaving the method each thread also negates its own local sense flag (so the contract between barrier and thread senses is kept for next usage). The code is small enough to fit here, and is pasted below:

```
1 public class SenseBarrier implements Barrier {
2     AtomicInteger count;        // how many threads have arrived
3     int size;                  // number of threads
4     volatile boolean sense;    // object's sense
5     ThreadLocal<Boolean> threadSense;
6
7     public SenseBarrier(int n) {
8         count = new AtomicInteger(n);
9         size = n;
10        sense = false;
11        threadSense = new ThreadLocal<Boolean>() {
12            protected Boolean initialValue() { return !sense; };
13        };
14    }
15 }
```

```

14     }
15
16     public void await() {
17         boolean mySense = threadSense.get();
18         int position = count.getAndDecrement();
19         if (position == 1) { // I'm last
20             count.set(this.size); // reset counter
21             sense = mySense;      // reverse sense
22         } else {
23             while (sense != mySense) {} // busy-wait
24         }
25         threadSense.set(!mySense);
26     }
27 }

```

One specially attractive feature of this approach, that we also saw while studying Spin locks in chapter 7, is that the threads are spinning on a local variable (which is expected to be cached on their respective core). One may think that the while expression is using both a local and a global variable, but the global one is likely to be cached for most of the time; and its entry in the cache gets invalidated only when is time to leave the barrier (after last thread changed it), so traffic on the cores' bus only occurs at the end.

13.2.3 Experiment Description

The test consists in having a shared array of 8 slots, one per thread to be created. Then the test has two phases, both using the barrier; the first phase is to wait all threads to write on their assigned slot on the array and the second stage is to allow all threads to validate that the rest actually finished their writing. These two phases are repeated 8 times.

The test above is basically the same as that described for *DisBarrierTest*.

13.2.4 Observations and Interpretations

Contrary to *DisBarrierTest*, this test does not hang on 2, nor on 24 core machines. This is because the shared variable *sense* is properly declared as *volatile*. Below a sample successful execution:


```
[oraadm@gdlaa008 orig]$ junit barrier.SenseBarrierTest  
.Testing 8 threads, 8 rounds
```

```
Time: 0.012
```

```
OK (1 test)
```

13.3 Reusable Barrier

13.3.1 Particular Case

In this experiment we implement an n-thread reusable barrier from an n-wire counting network and a single Boolean variable.

13.3.2 Solution

To implement a reusable barrier, we use a Combining Tree Barrier implementation and change the constructor in order to support the reuse of the barrier.

```
1  public RevBarrier(int _size, int _radix) {
2      radix = _radix;
3      leaves = 0;
4      this.leaf = new Node[_size / _radix];
5      int depth = 0;
6      threadSense = new ThreadLocal<Boolean>() {
7          protected Boolean initialValue() { return true; };
8      };
9      // compute tree depth
10     while (_size > 1) {
11         depth++;
12         _size = _size / _radix;
13         assert _size > 0;
14     }
15     Node root = new Node();
16     build(root, depth - 1);
17 }
18
19 // recursive tree constructor
20 void build(Node parent, int depth) {
21     // are we at a leaf node?
22     if (depth == 0) {
23         leaf[leaves++] = parent;
24     } else {
25         for (int i = 0; i < radix; i++) {
26             Node child = new Node(parent);
27             build(child, depth - 1);
28         }
29     }
30 }
```

13.3.3 Experiment Description

The test creates 8 threads that need to use the barrier *await()* in order to write values into the log and let the writer finish. If the round doesn't correspond to the value in the log then a System.out showing the expected value and found are displayed.

13.3.4 Observations and Interpretations

The tests executed as expected and no errors where found.

```
1 Tests run: 1, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
   0.511 sec
2
3 ----- Standard Output -----
4 Testing 8 threads , 8 rounds
5 -----
6 test-single:
7 BUILD SUCCESSFUL (total time: 0 seconds)
```

13.4 TreeBarrierTest

13.4.1 Particular Case (problem)

The problem we are trying to solve is that of synchronization on a multi-threaded environment; and the particular mechanism we want to use for that is a barrier (same as other cases from same chapter 17). The subproblem we care about here, is to reduce the contention generated on the last phase of solutions like *SenseBarrierTest*.

13.4.2 Solution

This solution is a natural extension of the “Sense Barrier” seen for exercise *SenseBarrierTest*; where the barrier as a whole is split into a several smaller sense barriers, and they are placed as the leaves of a tree (each leaf node is setup with same amount of elements, and this size is also called the *radix*). Then inside each leaf node, we will suffer smaller contention than with a single sense barrier (simply cause the number of threads to be unblocked is smaller). A key difference with the single sense barrier approach though, is that the children nodes need to call *await* method on the parent; this effectively means that all calls will eventually reach the root node (and that is the way to know all the threads have reached a consensus, and they can move on). The relevant parts of code are shown below:

```
1 public class TreeBarrier implements Barrier {
2     int radix;    // tree fan-in
3     Node[] leaf;  // array of leaf nodes
4     int leaves;   // used to build tree
5     ThreadLocal<Boolean> threadSense; // thread-local sense
6     public TreeBarrier(int n, int r) {
7         radix = r;
8         leaves = 0;
9         this.leaf = new Node[n / r];
10        int depth = 0;
11        threadSense = new ThreadLocal<Boolean>() {
12            protected Boolean initialValue() { return true; };
13        };
14        // compute tree depth
15        while (n > 1) {
16            depth++;
17            n = n / r;
```

```

18     }
19     Node root = new Node();
20     build(root, depth - 1);
21 }
22
23 void build(Node parent, int depth) {
24     // are we at a leaf node?
25     if (depth == 0) {
26         leaf[leaves++] = parent;
27     } else {
28         for (int i = 0; i < radix; i++) {
29             Node child = new Node(parent);
30             build(child, depth - 1);
31         }
32     }
33 }
34
35 public void await() {
36     int me = ThreadID.get();
37     Node myLeaf = leaf[me / radix];
38     myLeaf.await();
39 }
40
41 private class Node {
42     AtomicInteger count;
43     Node parent;
44     volatile boolean sense;
45     // construct root node
46     public Node() {
47         sense = false;
48         parent = null;
49         count = new AtomicInteger(radix);
50     }
51     public Node(Node parent) {
52         this();
53         this.parent = parent;
54     }
55     public void await() {
56         boolean mySense = threadSense.get();
57         int position = count.getAndDecrement();
58         if (position == 1) { // I'm last
59             if (parent != null) { // root?
60                 parent.await();
61             }
62             count.set(radix); // reset counter

```

```

63         sense = mySense;
64     } else {
65         while (sense != mySense) {};
66     }
67     threadSense.set(!mySense);
68 }
69 }
70 }

```

13.4.3 Experiment Description

The test consists in having a shared array of slots, as big as the number of threads to use. Then the test has two phases, both using the barrier; the first phase is to wait all threads to write on their assigned slot on the array and the second stage is to allow all threads to validate that the rest actually finished their writing. These two phases are repeated 8 times.

The test above is basically the same as that described for *DisBarrierTest* and *SenseBarrierTest*; except that we have a couple of tests instead of one (each one tries a different radix ¹ and number of threads; the combinations of these two parameters are *threads* = 8, *radix* = 2 and *threads* = 27, *radix* = 3 respectively.

13.4.4 Observations and Interpretations

The test runs fine in 2 and 24 core machines, a sample execution is shown below:

```

[oraadm@gdlaa008 orig]$ junit barrier.TreeBarrierTest
.Testing 27 threads, 8 rounds
finished radix 3
.Testing 8 threads, 8 rounds
finished radix 2

Time: 0.023

OK (2 tests)

```

¹The *radix* parameter controls the size of the leaf nodes, which represent smaller sense barriers

Chapter 14

Transactional Memory

14.1 Restricted Bloom Filter

14.1.1 Particular Case

In this experiment we implement a bloom filters to reduce expensive disk lookups for non-existent keys. If the element is not in the bloom filter, then we know for sure we don't need to perform the expensive lookup. On the other hand, if it is in the bloom filter, we perform the lookup.

14.1.2 Solution

The basic bloom filter supports two operations, *contains()* and *add()*. *Contains* is used to check whether a given element is in the set or not.

```
1 public boolean contains(Object x) {  
2     return filter.get(hash0(x)) && filter.get(hash1(x));  
3 }
```

Add simply adds an element to the set. Removal is impossible without introducing false negatives.

```
1 public void add(Object x) {  
2     filter.set(hash0(x));  
3     filter.set(hash1(x));  
4 }
```

14.1.3 Experiment Description

The test creates three values and adds two of them to the filter, then we check if the hash contains them and only fail if the assertion fails if a non-existing value returns true.

14.1.4 Observations and Interpretations

The tests executed as expected and no errors were found.

```
1 Testsuite: TinyTM.locking.BloomFilterTest  
2 Tests run: 1, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:  
   0.393 sec  
3  
4 test-single:  
5 BUILD SUCCESSFUL (total time: 0 seconds)
```


14.2 Lazy Skiplist

14.2.1 Particular Case

In this experiment we implement a Lazy Skiplist which uses an optimistic fine-grained locking, meaning that the method searches for its target node without locking, and acquires locks and validates only when it discovers the target.

14.2.2 Solution

According to the theory we use a skiplist to implement a set providing the methods, *add(x)* adds *x* to the set, *remove(x)* removes *x* from the set, and *contains(x)* returns true if, and only if *x* is in the set. Recall that a skiplist is a collection of linked lists. Each node in the list contains an item field (an element of the set), a key field (the items hash code), and a next field, which is an array of references to successor nodes in the list. Using *TinyTM*, threads communicate via shared atomic objects, which provide synchronization, ensuring that transactions cannot see one another's uncommitted effects, and recovery, undoing the effects of aborted transactions. Fields of atomic objects are not directly accessible. Instead, they are accessed indirectly through getter and setter methods. Accessing fields through getters and setters provides the ability to interpose transactional synchronization and recovery on each field access.

14.2.3 Experiment Description

The test creates 32 thread that parallel insert values, commits them and aborts then and compare the length and commits. If values are missing an assert fail is thrown.

14.2.4 Observations and Interpretations

The tests executed as expected and no errors where found.

1	Tests run: 3, Failures: 0, Errors: 0, Skipped: 0, Time elapsed:
2	12.445 sec
3	----- Standard Output -----
4	TestSequential

```
5 | inserts: 186147, removes: 185635, missed: 370334, delta: 512
6 | commits: 742116, aborts: 0
7 | SkipListSet OK
8 | length = 1536 (expected 1536)
9 | Commits: 742116 (expected 742116)
10 | TestParallel(32)
11 | inserts: 24, removes: 0, missed: 41, delta: 24
12 | commits: 65, aborts: 66
13 | SkipListSet OK
14 | length = 1048 (expected 1048)
15 | Commits: 65 (expected 65)
16 | _____
17 | test-single:
18 | BUILD SUCCESSFUL (total time: 12 seconds)
```

14.3 ListTest (ofree)

14.3.1 Particular Case (problem)

The problem we try to solve here is how to implement a concurrent list by using software transactional (*STM*) memory.

14.3.2 Solution

The solution uses the “obstruction-free” flavor of the *STM*, where the building blocks of the list use the so called *FreeObject*. Each one of these free objects has three logical fields: an *owner* field to keep track of what thread is using it, *oldVersion* which is the object state prior beginning of transaction, and *newVersion* version that reflects current transaction updates. The possible values for *oldVersion* and *newVersion* fields are *COMMITTED*, *ABORTED* and *ABORTED*.

Each time a transaction *A* accesses one of these objects for the first time, it “opens” that object (possibly resetting the logical fields mentioned above). If we call *B* the previous owner transaction, the following rules govern this solution (quoted from the book):

- If *B* was *COMMITTED*, then the new version is current. *A* installs itself as the object’s current owner, sets the old version to the prior new version, and the new version to a copy of the prior new version (if the call is a setter), or to the new version itself (if the call is a getter).
- Symmetrically, if *B* was *ABORTED*, then the old version is current. *A* installs itself as the objects current owner, sets the old version to the prior old version, and the new version to a copy of the prior old version (if the call is a setter), or to the old version itself (if the call is a getter).
- If *B* is still *ACTIVE*, then *A* and *B* conflict, so *A* consults the contention manager for advice whether to abort *B*, or to pause, giving *B* a chance to finish. One transaction aborts another by successfully calling *compareAndSet()* to change the victim’s status to *ABORTED*.

Contrary to other chapters, the transactional memory one implies much more code to consider; so we did not consider relevant to paste it all here. The explanation given above hopefully suffices for our purposes.

14.3.3 Experiment Description

The test program populates serially the list protected by STM (initial size of 1024), and then creates certain amount of threads to perform indefinitely, alternating random insertions and deletions. The main test thread interrupts them and performs sanity checks. This test is performed with 1 (sequential) and 32 threads (parallel).

14.3.4 Observations and Interpretations

The test using a single thread (sequential), gives the following null pointer exception on the 24 cores machine:

```
.TestSequential
TinyTM.exceptions.PanicException: java.lang.NullPointerException
at TinyTM.TThread.doIt(TThread.java:51)
at TinyTM.list.ofree.ListTest$TestThread.run(ListTest.java:117)
Caused by: java.lang.NullPointerException
at TinyTM.list.ofree.TNode.getNext(TNode.java:54)
at TinyTM.list.ofree.List.add(List.java:40)
at TinyTM.list.ofree.ListTest$TestThread$1.call(ListTest.java:119)
at TinyTM.list.ofree.ListTest$TestThread$1.call(ListTest.java:117)
at TinyTM.TThread.doIt(TThread.java:41)
... 1 more
inserts: 0, removes: 0, missed: 0, delta: 0
commits: 0, aborts: 0
```

While debugging the sequential exception above, we found that in method *FreeObject:openRead*, a new location was never getting a *newVersion* value after its construction:

```
Locator newLocator = new Locator();
... more code ...
return newLocator.newVersion <== NULL !!!
```

As possible solutions for this problem, we realized that either we could use constructor that takes *newVersion* as argument or we set it like with method *openWrite*. We opted for the first option, using existing *locator* object:

```
Locator newLocator = new Locator(locator.newVersion);
```

With above fix the sequential test no longer presented problems on 2 nor on 24 cores. But the parallel one was still failing with assertion errors like the ones below:

```
Testcase: testParallel(TinyTM.list.ofree.ListTest): FAILED
Bad size, expected:<1518> but was:<1502>
junit.framework.AssertionFailedError: Bad size, expected:<1518> but
was:<1502>
at TinyTM.list.ofree.ListTest.sanityCheck(ListTest.java:196)
at TinyTM.list.ofree.ListTest.testParallel(ListTest.java:84)
```

Due time constraints, we could not dig further into the reasons for the above error; but professor acknowledged it was ok (after all, we covered the entire set of programs from the book).